



CLOUDNATIVE **SECURITYCON**

NORTH AMERICA 2023





CLOUDNATIVE
SECURITYCON

NORTH AMERICA 2023

From the Cluster to the Cloud: Lateral Movements in Kubernetes

Yossi Weizman & Ram Pliskin, Microsoft



Agenda



- Identity types in Kubernetes
- Inner-cluster lateral movement
- Cluster-to-cloud lateral movement: Azure, AWS and GCP
- Detections & mitigations
- Key takeaways

Identity types in Kubernetes



Three main areas:

- How users (or applications) from outside the cluster authenticate with the cluster.
- How workloads in the cluster authenticate within the cluster.
- How workloads in the cluster authenticate with resources in the cloud outside the cluster.

Identity types in Kubernetes



Three main areas:

- How users (or applications) from outside the cluster authenticate with the cluster.
- **How workloads in the cluster authenticate within the cluster.**
(Inner-cluster lateral movement)
- **How workloads in the cluster authenticate with resources in the cloud outside the cluster.**
(Cluster-to-cloud lateral movement)

Inner-cluster lateral movement



Inner-cluster lateral movement




Let's assume a pod is compromised

K8s control plane



Node 1




Node 2




Node 3




K8s control plane



Node 1




Node 2




Node 3



K8s control plane



Node 1



Node 2




Node 3




K8s control plane



Node 1



Node 2




Node 3



K8s control plane



Node 1




Node 2




Node 3



K8s control plane




Node 1



Node 2

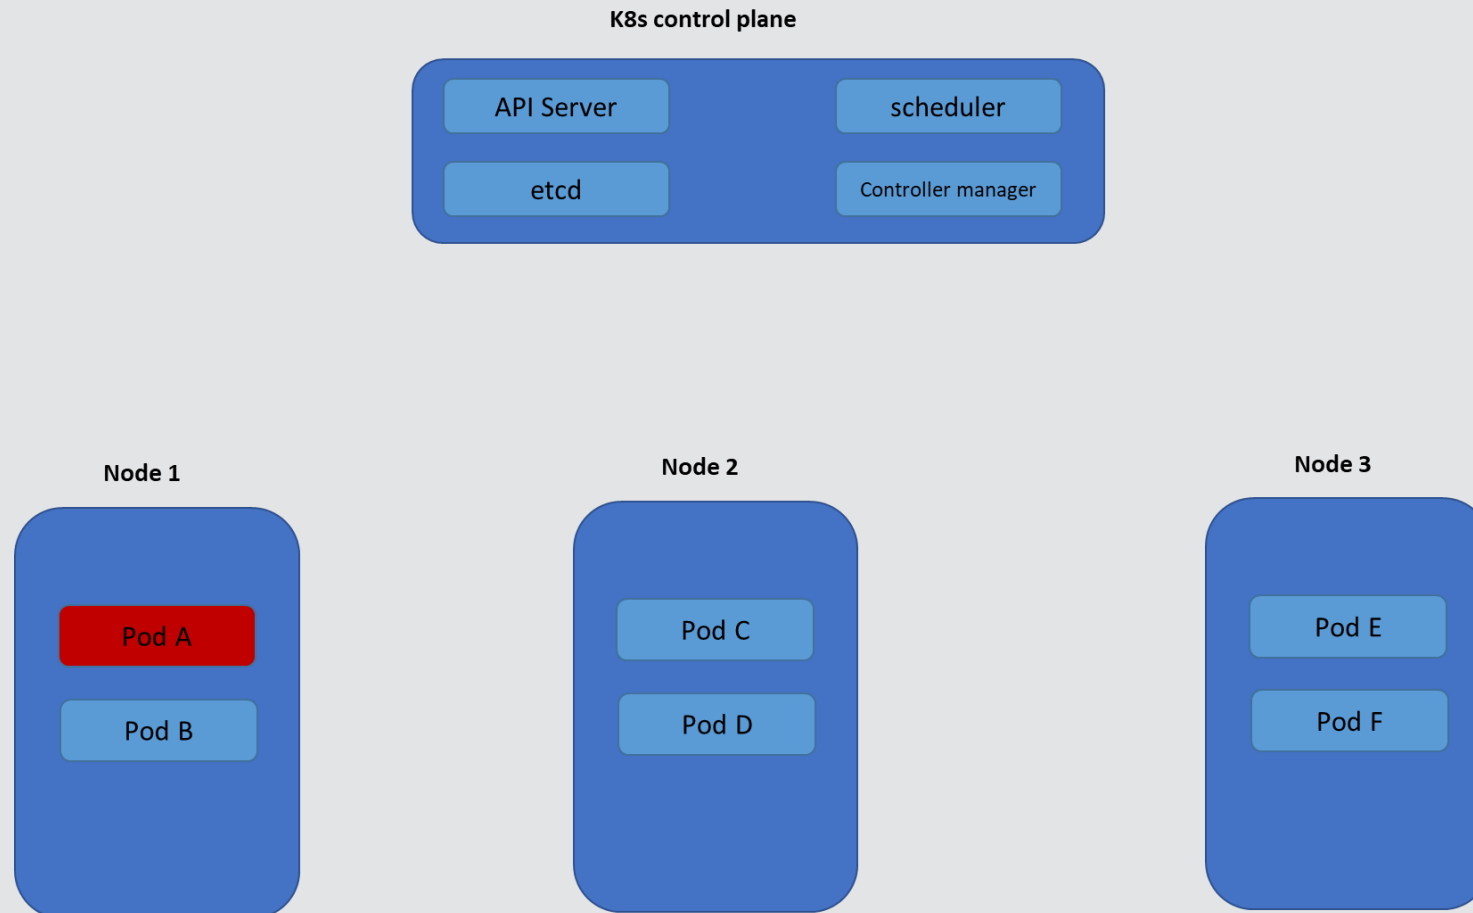


Node 3



Inner-cluster lateral movement


How can attackers leverage a compromised pod for cluster takeover?




K8s control plane



Node 1



Node 2



Node 3



Inner-cluster lateral movement

How can attackers leverage a compromised pod for cluster takeover?

Good news – it becomes more difficult:

- Read secrets permission isn't enough
 - Newer versions of K8s don't create long-lived SA tokens as secret objects
- Node takeover \neq cluster takeover
 - Node authorizer + NodeRestriction admission controller limit the permissions of the nodes' identities

Inner-cluster lateral movement

How can attackers leverage a compromised pod for cluster takeover?

Good news – it becomes more difficult:

- Read secrets permission isn't enough
 - Newer versions of K8s don't create long-lived SA tokens as secret objects
- Node takeover \neq cluster takeover
 - Node authorizer + NodeRestriction admission controller limit the permissions of the nodes' identities

However, common misconfigurations still allow it.

Inner-cluster lateral movement



How can attackers leverage a compromised pod for cluster takeover?

Good news – it becomes more difficult:

- Read secrets permission isn't enough
 - Newer versions of K8s don't create long-lived SA tokens as secret objects
- Node takeover \neq cluster takeover
 - Node authorizer + NodeRestriction admission controller limit the permissions of the nodes' identities

However, common misconfigurations still allow it.

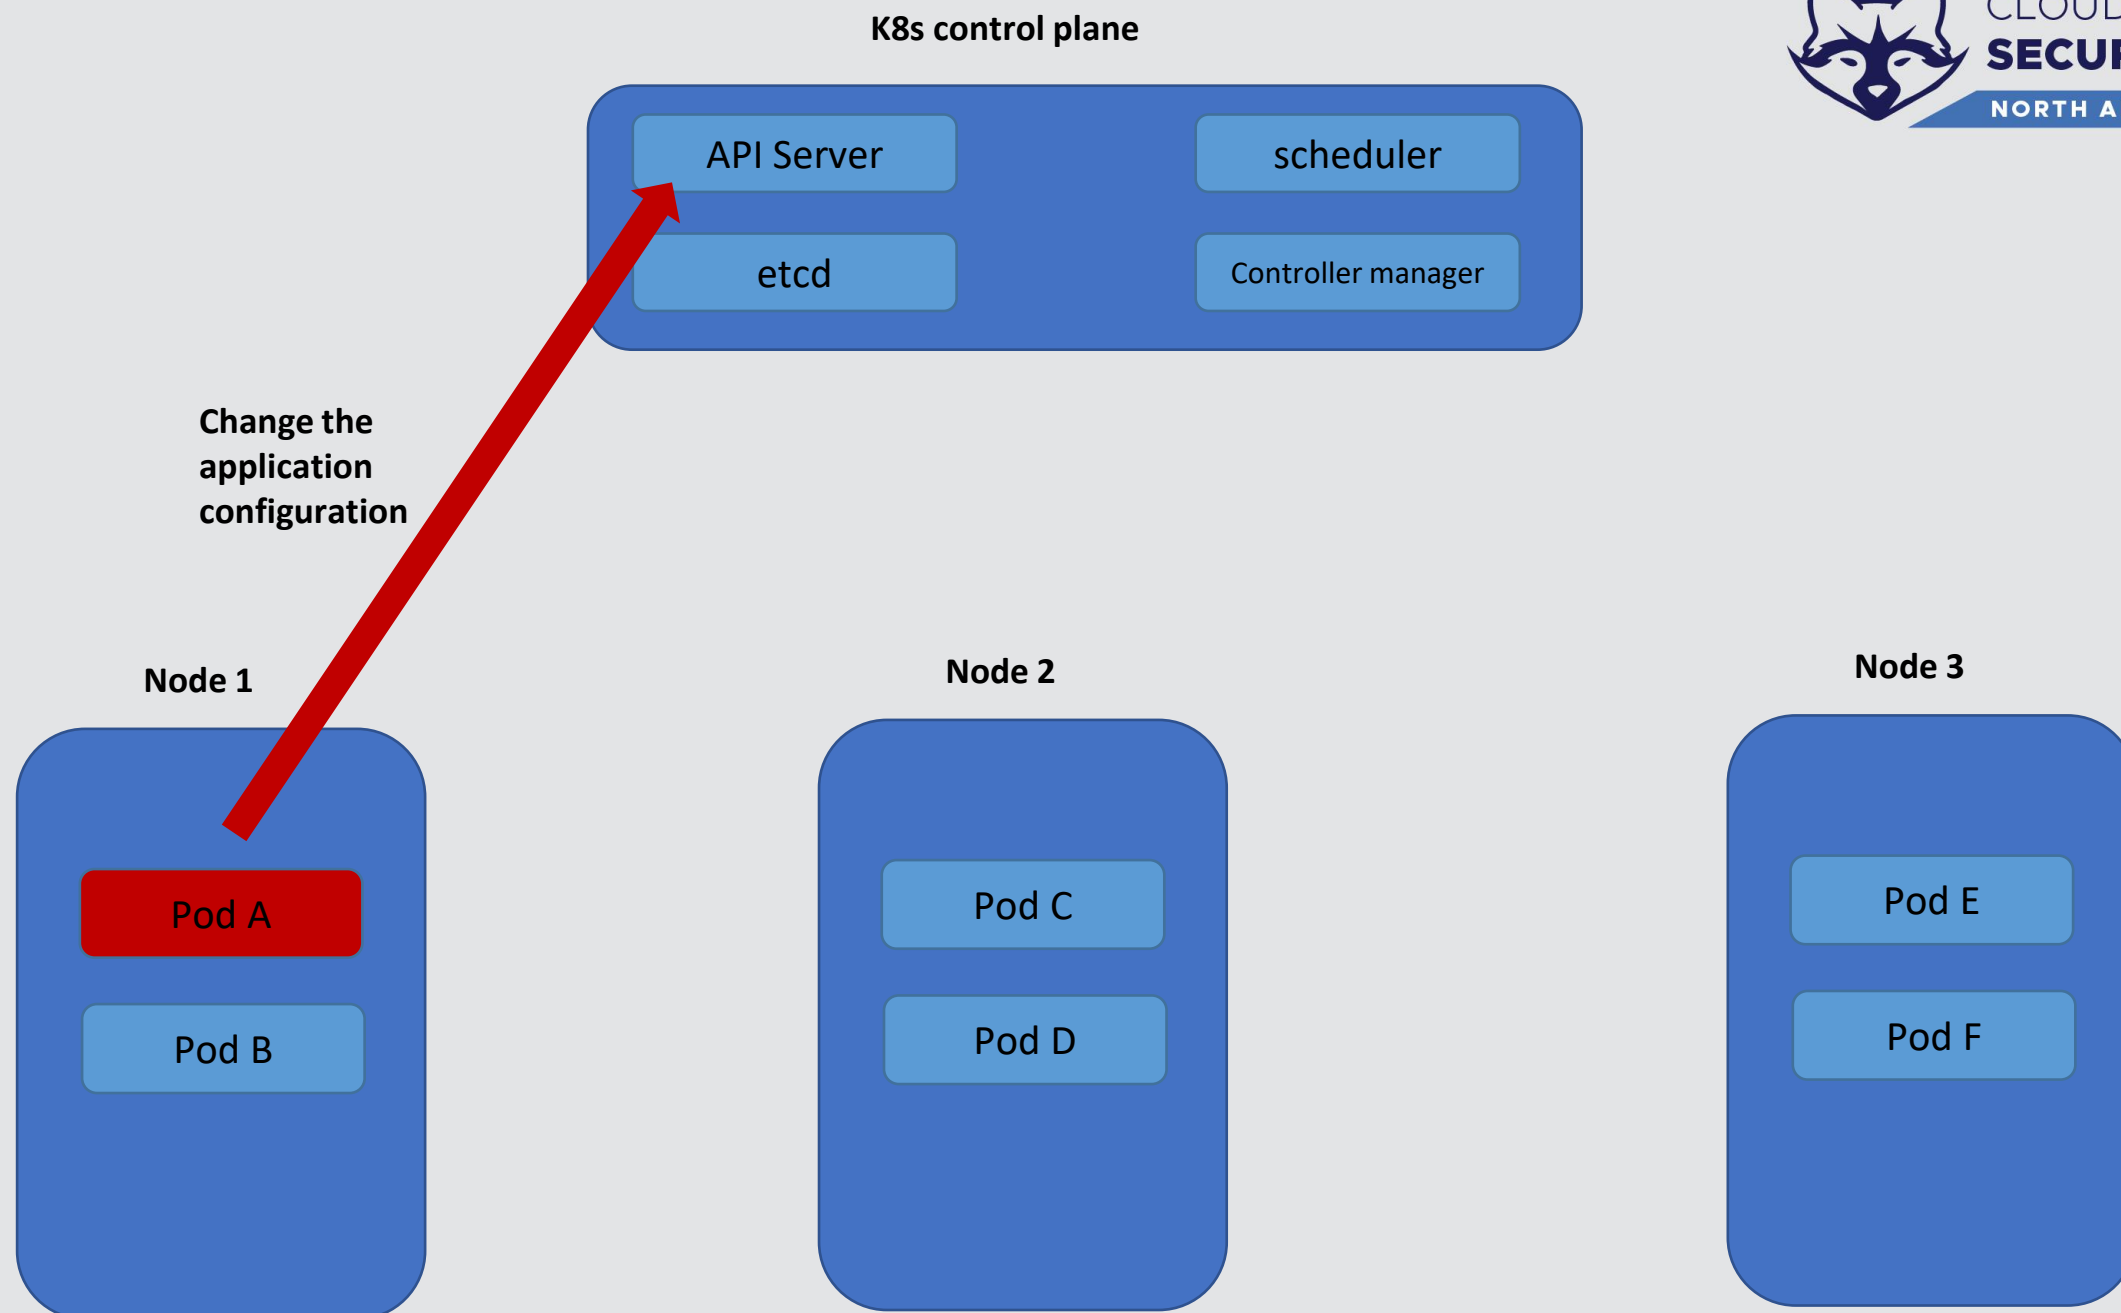
Let's see an example

Inner-cluster lateral movement

Example: Self-update permissions

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: app-update
rules:
- apiGroups:
  - 'apps'
  resources:
  - 'deployments'
  verbs:
  - 'update'
  resourceNames:
  - 'my-app'
```

- We've observed cases in which applications had permissions to update themselves.
- This allows the applications to change their own configuration. For example: update their configuration to be privileged, change their service account, schedule to specific node(s).
- Effectively, this can lead to cluster takeover




K8s control plane



Node 1




Node 2




Node 3



K8s control plane



Node 1




Node 2




Node 3



K8s control plane



Node 1




Node 2




Node 3



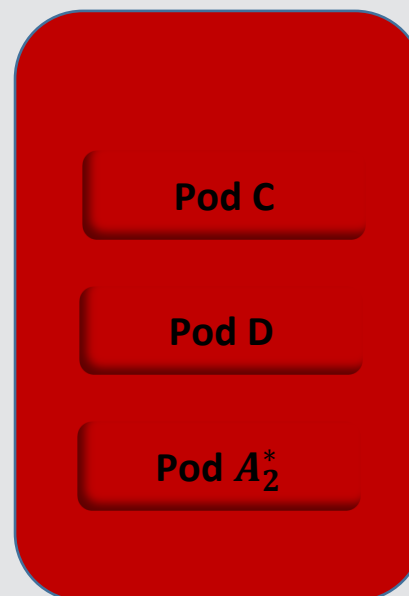
K8s control plane




Node 1



Node 2



Node 3



Inner-cluster lateral movement

Other sensitive permissions (partial list):

Permission	Allows to
Create pod\controllers	Use a privileged service account in a new pod
Update controller	Change the configuration to use a privileged service account
Create secrets & get\list secret	Create a new long-lived token for SA and read its value
Create serviceaccounts/token	Create a short-lived token for SA

Cluster-to-cloud lateral movement



Cluster-to-cloud lateral movement



How tightly coupled are Kubernetes clusters with their Cloud hosting environment?

1. Managed clusters use cloud services for their ongoing operation (e.g., VMs allocation)
2. Workloads in Kubernetes may need access to cloud resources (e.g., cloud storage, cloud secret store, etc.)

Cluster-to-cloud lateral movement



How do workloads in Kubernetes authenticate with the cloud?

We'll go over the following authentication methods:

1. Storing cloud identity credential on the node
2. Direct access to IMDS
3. Indirect access to IMDS
4. Using OIDC (identity federation)

Cluster-to-cloud lateral movement

Storing cloud identity credential on the node



- Used to be the default authentication method in AKS in the past.
- In this method, each Kubernetes node stores a file with service principal (SPN) credentials. SPNs are Azure application identities.
- By default, this SPN has Contributor role for the node resource group.
- Users can bring their own SPN or grant more permissions to the SPN if their applications need access to more cloud resources. For example: add permissions to a cloud storage.

Cluster-to-cloud lateral movement

Storing cloud identity credential on the node



- Used to be the default authentication method in AKS in the past.
 - In this method, each Kubernetes node stores a file with service principal (SPN) credentials. SPNs are Azure application identities.
 - By default, this SPN has Contributor role for the node resource group.
 - Users can bring their own SPN or grant more permissions to the SPN if their applications need access to more cloud resources. For example: add permissions to a cloud storage.
- Access to the node's File System meant elevation to Contributor Role on the K8s resource group



Kubernetes Service Account




Azure Service principal



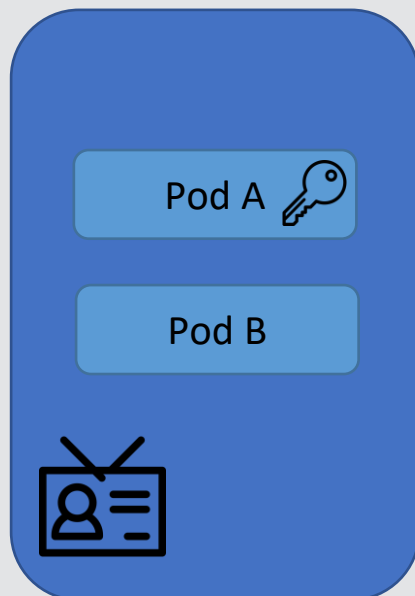
CLOUDNATIVE
SECURITYCON

NORTH AMERICA 2023


K8s control plane



Node 1



Node 2



Node 3





Kubernetes Service Account




Azure Service principal




CLOUDNATIVE
SECURITYCON
NORTH AMERICA 2023

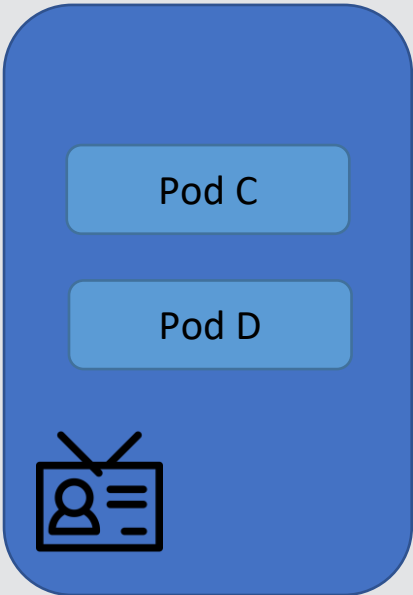
K8s control plane



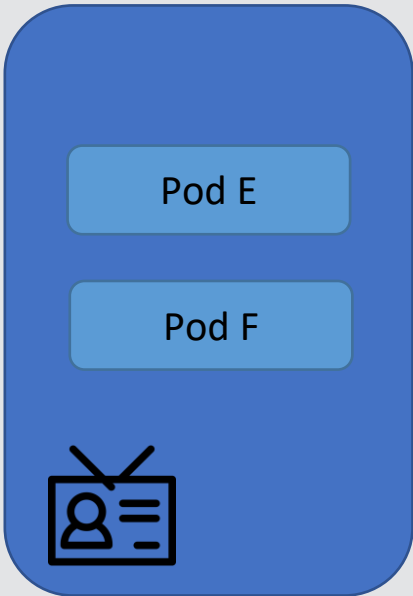
Node 1



Node 2



Node 3





Kubernetes Service Account




CLOUDNATIVE
SECURITYCON

NORTH AMERICA 2023



Azure Service principal


K8s control plane




“Pod create” operation

In the pod config: mount the SPN
into the container

Node 1



Node 2



Node 3





Kubernetes Service Account



Azure Service principal




CLOUDNATIVE
SECURITYCON
NORTH AMERICA 2023


K8s control plane



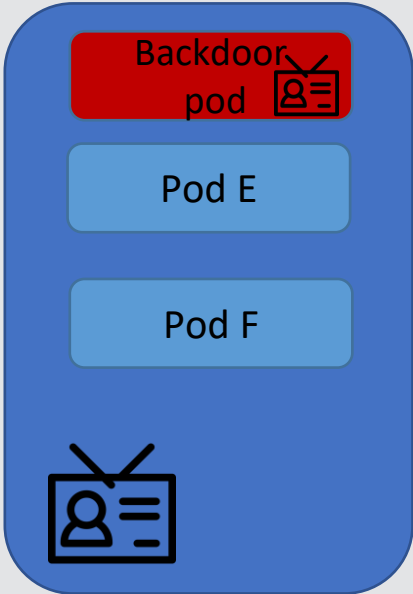
Node 1



Node 2



Node 3






Kubernetes Service Account




Azure Service principal

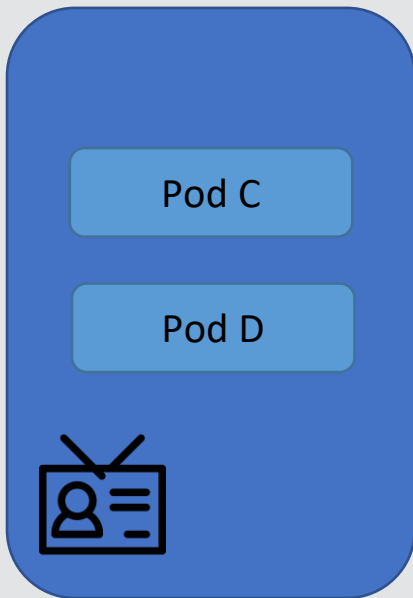
K8s control plane



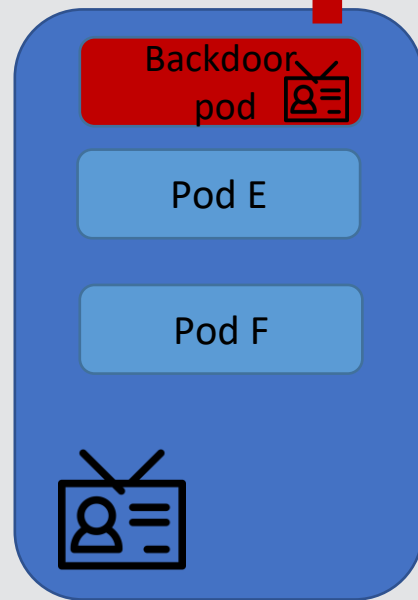
Node 1



Node 2



Node 3



Cluster-to-cloud lateral movement

Direct access to IMDS



Cluster-to-cloud lateral movement



Direct access to IMDS

- The metadata service is a special endpoint that is accessible to VMs, allowing them to retrieve information about the VM.
- Implemented by all major cloud providers.
- Metadata service allows retrieving tokens for the cloud identity that is attached to the VM:
 - Azure: 169.254.169.254/metadata/identity/oauth2
 - AWS: 169.254.169.254/latest/meta-data/iam/security-credentials
 - GCP: metadata.google.internal/computeMetadata/v1/instance/service-accounts
- Querying the metadata service doesn't require any authentication.

Cluster-to-cloud lateral movement

Direct access to IMDS



- In managed K8s clusters, the nodes are VMs which can access to their metadata service.
- By default, pods can access to the metadata service of their nodes.
- Thus, pods can acquire tokens of cloud identities attached to the nodes.
- The permissions of the identities depend on the cloud provider and the specific environment.

Cluster-to-cloud lateral movement

Direct access to IMDS - AKS



- AKS uses several managed identities to operate the cluster.
- Users can change the default permissions of those identities, or alternatively attach additional managed identities to the nodes.

Cluster-to-cloud lateral movement

Direct access to IMDS - AKS

Identity	Name	Use case	Default permissions	Bring your own identity
Control plane	AKS Cluster Name	Used by AKS control plane components to manage cluster resources including ingress load balancers and AKS managed public IPs, Cluster Autoscaler, Azure Disk & File CSI drivers	Contributor role for Node resource group	Supported
Kubelet	AKS Cluster Name-agentpool	Authentication with Azure Container Registry (ACR)	NA (for kubernetes v1.15+)	Supported
Add-on	AzureNPM	No identity required	NA	No
Add-on	AzureCNI network monitoring	No identity required	NA	No
Add-on	azure-policy (gatekeeper)	No identity required	NA	No
Add-on	azure-policy	No identity required	NA	No
Add-on	Calico	No identity required	NA	No
Add-on	Dashboard	No identity required	NA	No
Add-on	HTTPApplicationRouting	Manages required network resources	Reader role for node resource group, contributor role for DNS zone	No
Add-on	Ingress application gateway	Manages required network resources	Contributor role for node resource group	No
Add-on	omsagent	Used to send AKS metrics to Azure Monitor	Monitoring Metrics Publisher role	No
Add-on	Virtual-Node (ACIConnector)	Manages required network resources for Azure Container Instances (ACI)	Contributor role for node resource group	No
OSS project	aad-pod-identity	Enables applications to access cloud resources securely with Microsoft Azure Active Directory (Azure AD)	NA	Steps to grant permission at https://github.com/Azure/aad-pod-identity#role-assignment ↗.

Cluster-to-cloud lateral movement

Direct access to IMDS - EKS



- EKS uses EC2 Roles for the Kubernetes nodes.
- By default, the EC2 Role has the following policies:
 - AmazonEC2ContainerRegistryReadOnly - Pull permissions to the container registry
 - AmazonEKSWorkerNodePolicy - Read permissions of the compute environment (EC2, VPC, etc.)
 - AmazonEKS_CNI_Policy – Attach network interfaces and IPs to VMs
- Users can add more policies, if their containers require access to cloud resources.

Cluster-to-cloud lateral movement

Direct access to IMDS - GKE



- GKE uses IAM service accounts to authenticate with the cloud.
- By default, all the VMs in a project, including the Kubernetes nodes, share a default SA.
- This SA has Editor role for the project.
- While the access scope limits the permissions, they are still powerful by default:

Security		
Sandbox with gVisor		Disabled
Service account ?		default
Access scopes	Service Control	Enabled
	Service Management	Read Only
	Stackdriver Logging API	Write Only
	Stackdriver Monitoring API	Full
	Stackdriver Trace	Write Only
	Storage	Read Only
GKE Metadata Server		Disabled
Integrity monitoring		Enabled
Secure boot		Disabled

Cluster-to-cloud lateral movement

Direct access to IMDS




How does lateral movement from the cluster to the cloud would look like?

K8s control plane



Node 1




Node 2




Node 3



K8s control plane





Node 1



K8s control plane



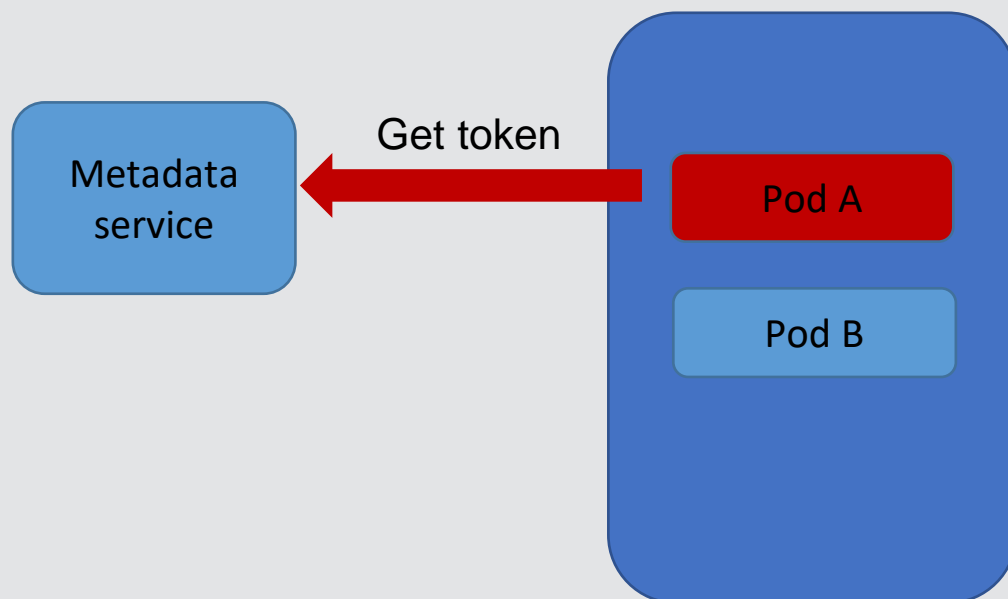
Node 1




K8s control plane




Node 1

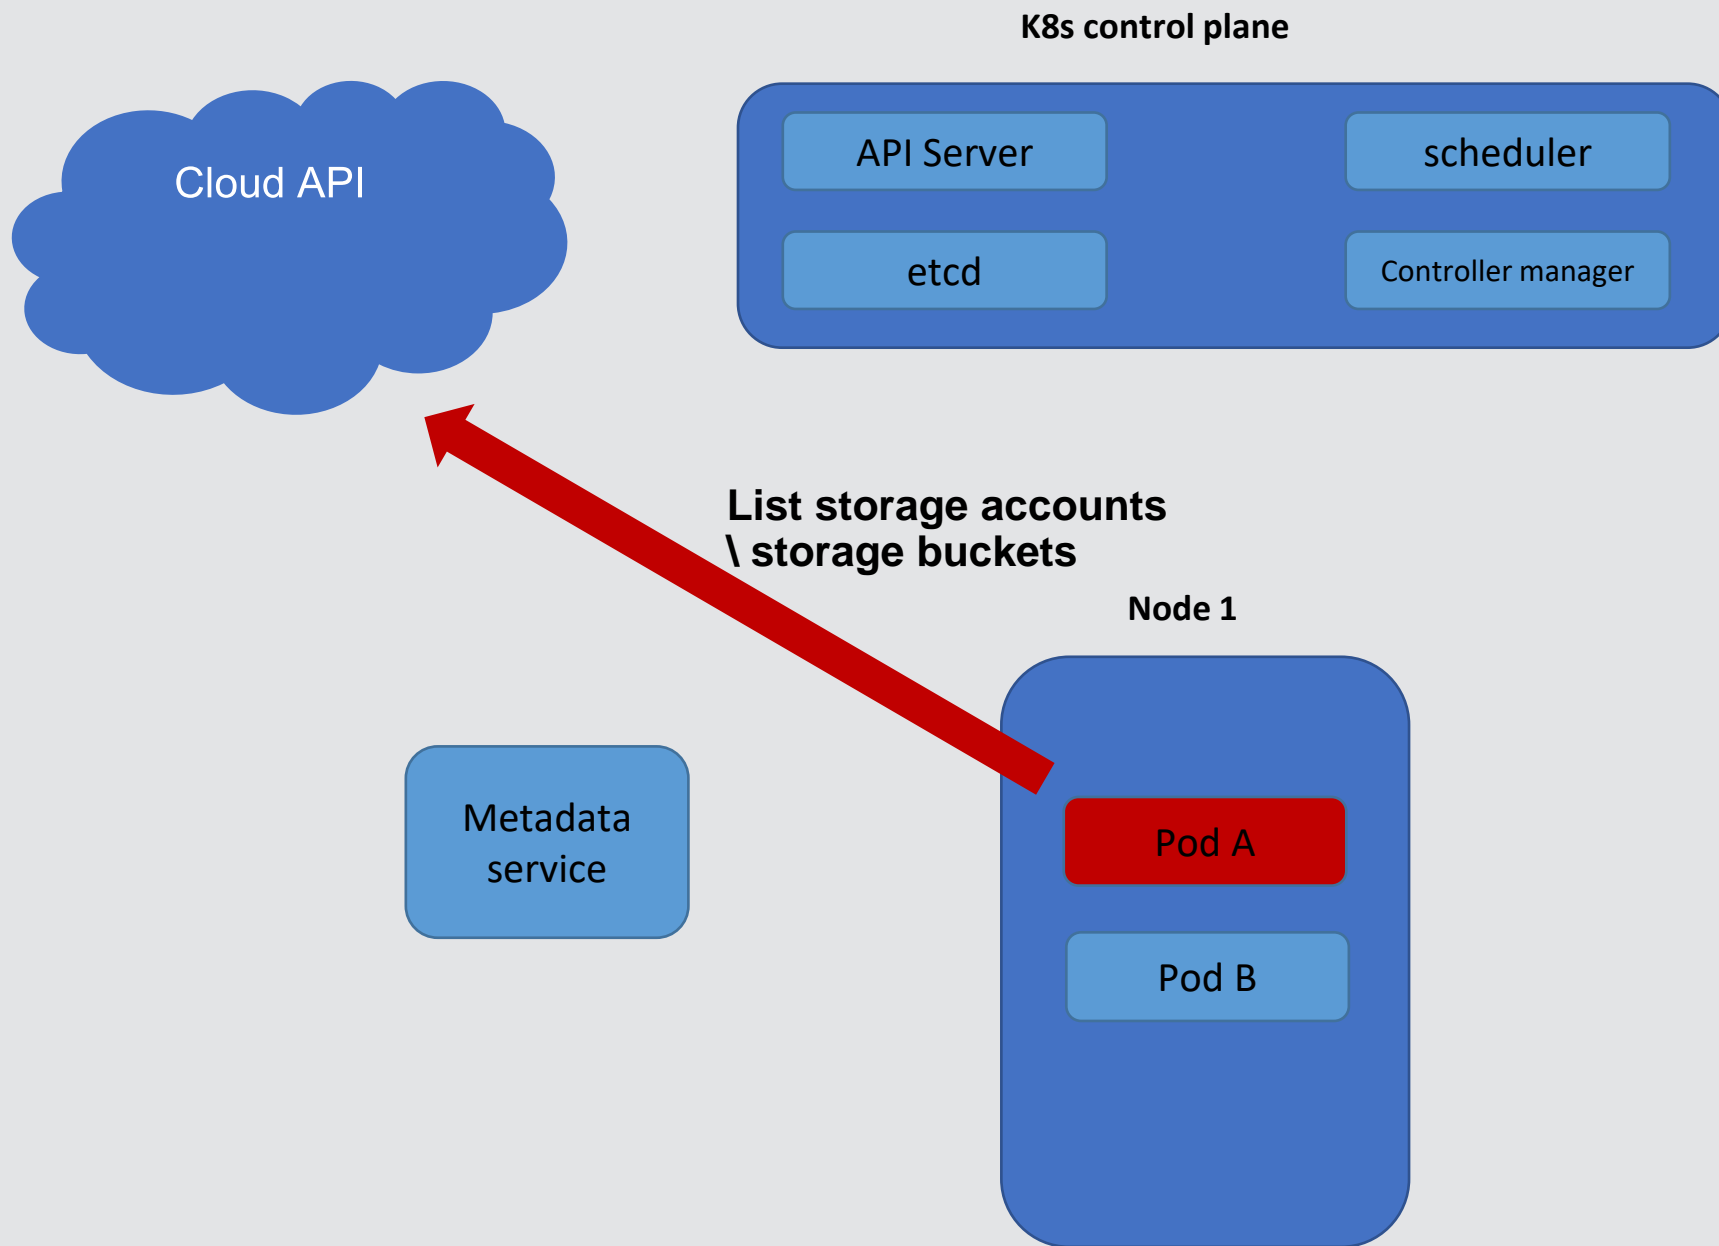


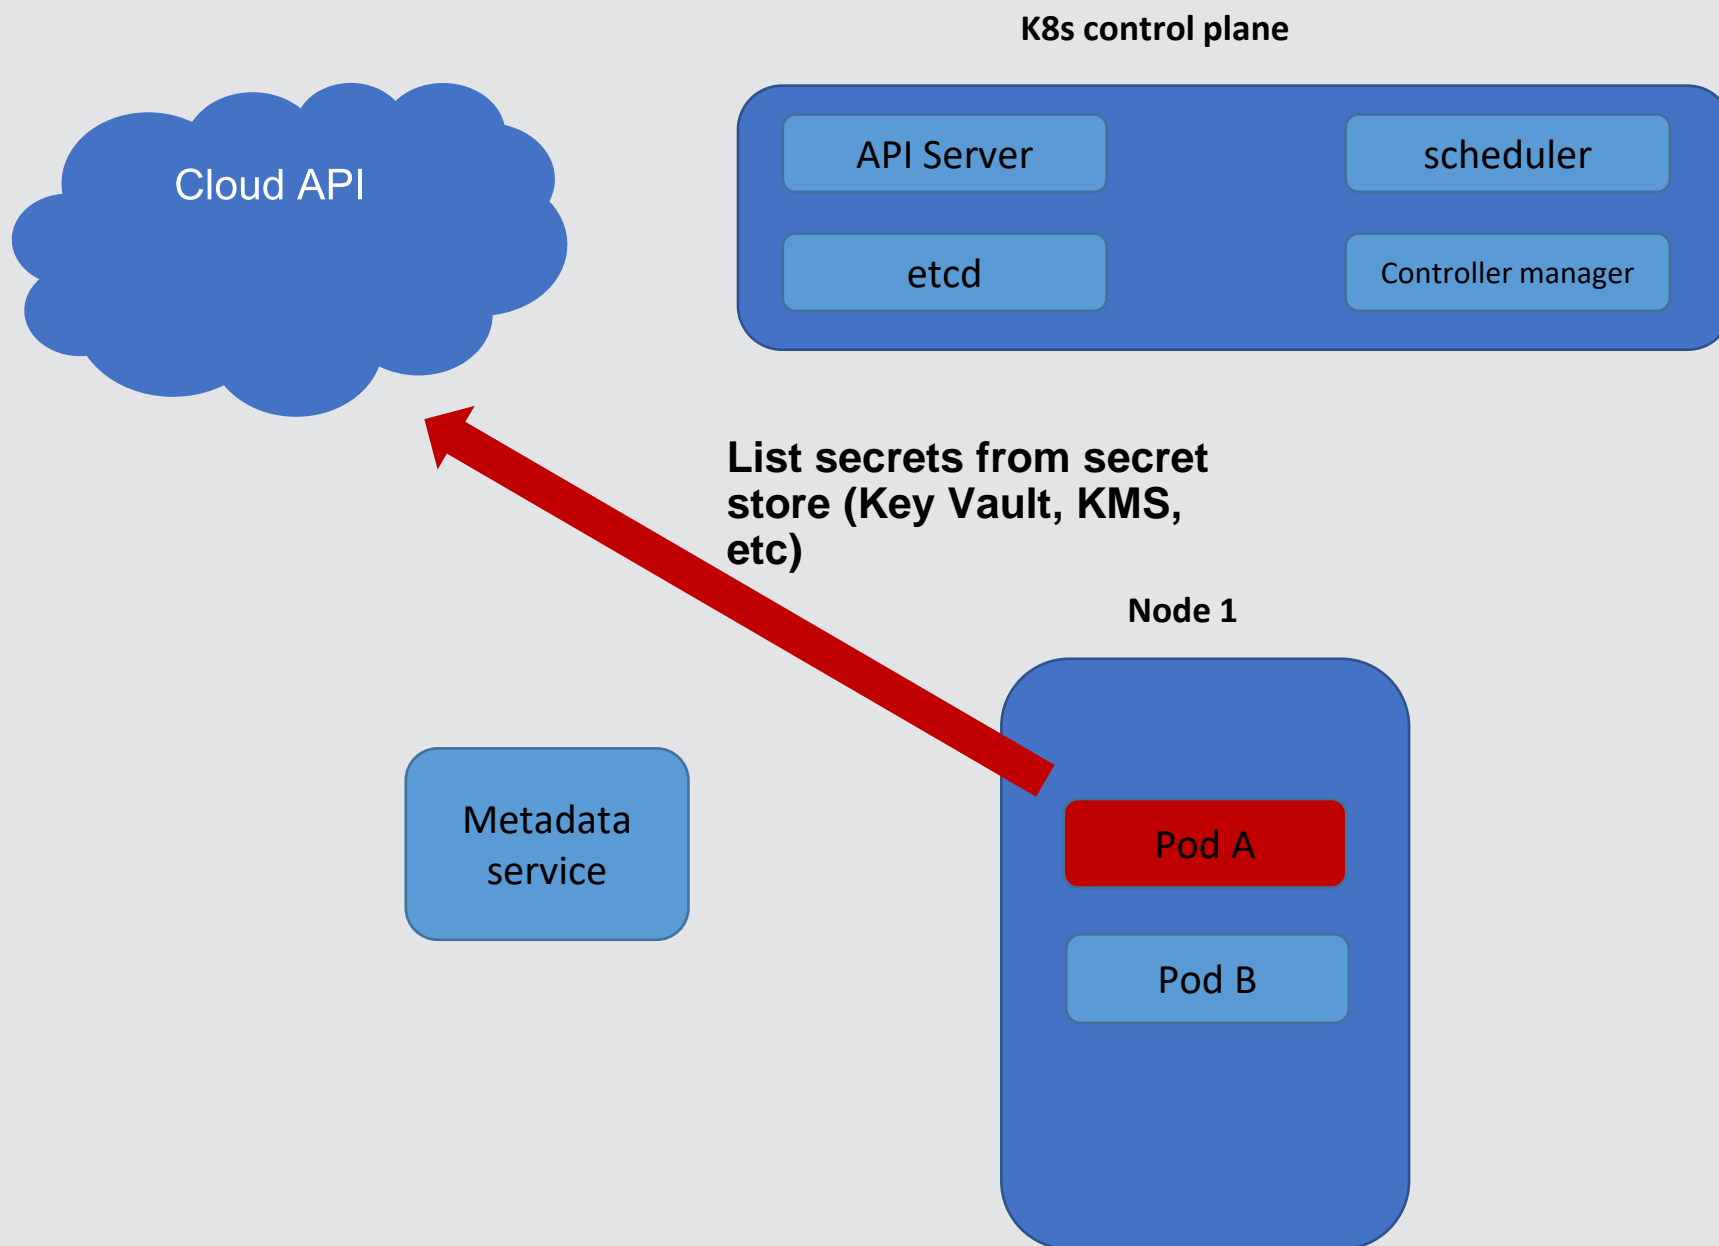
K8s control plane

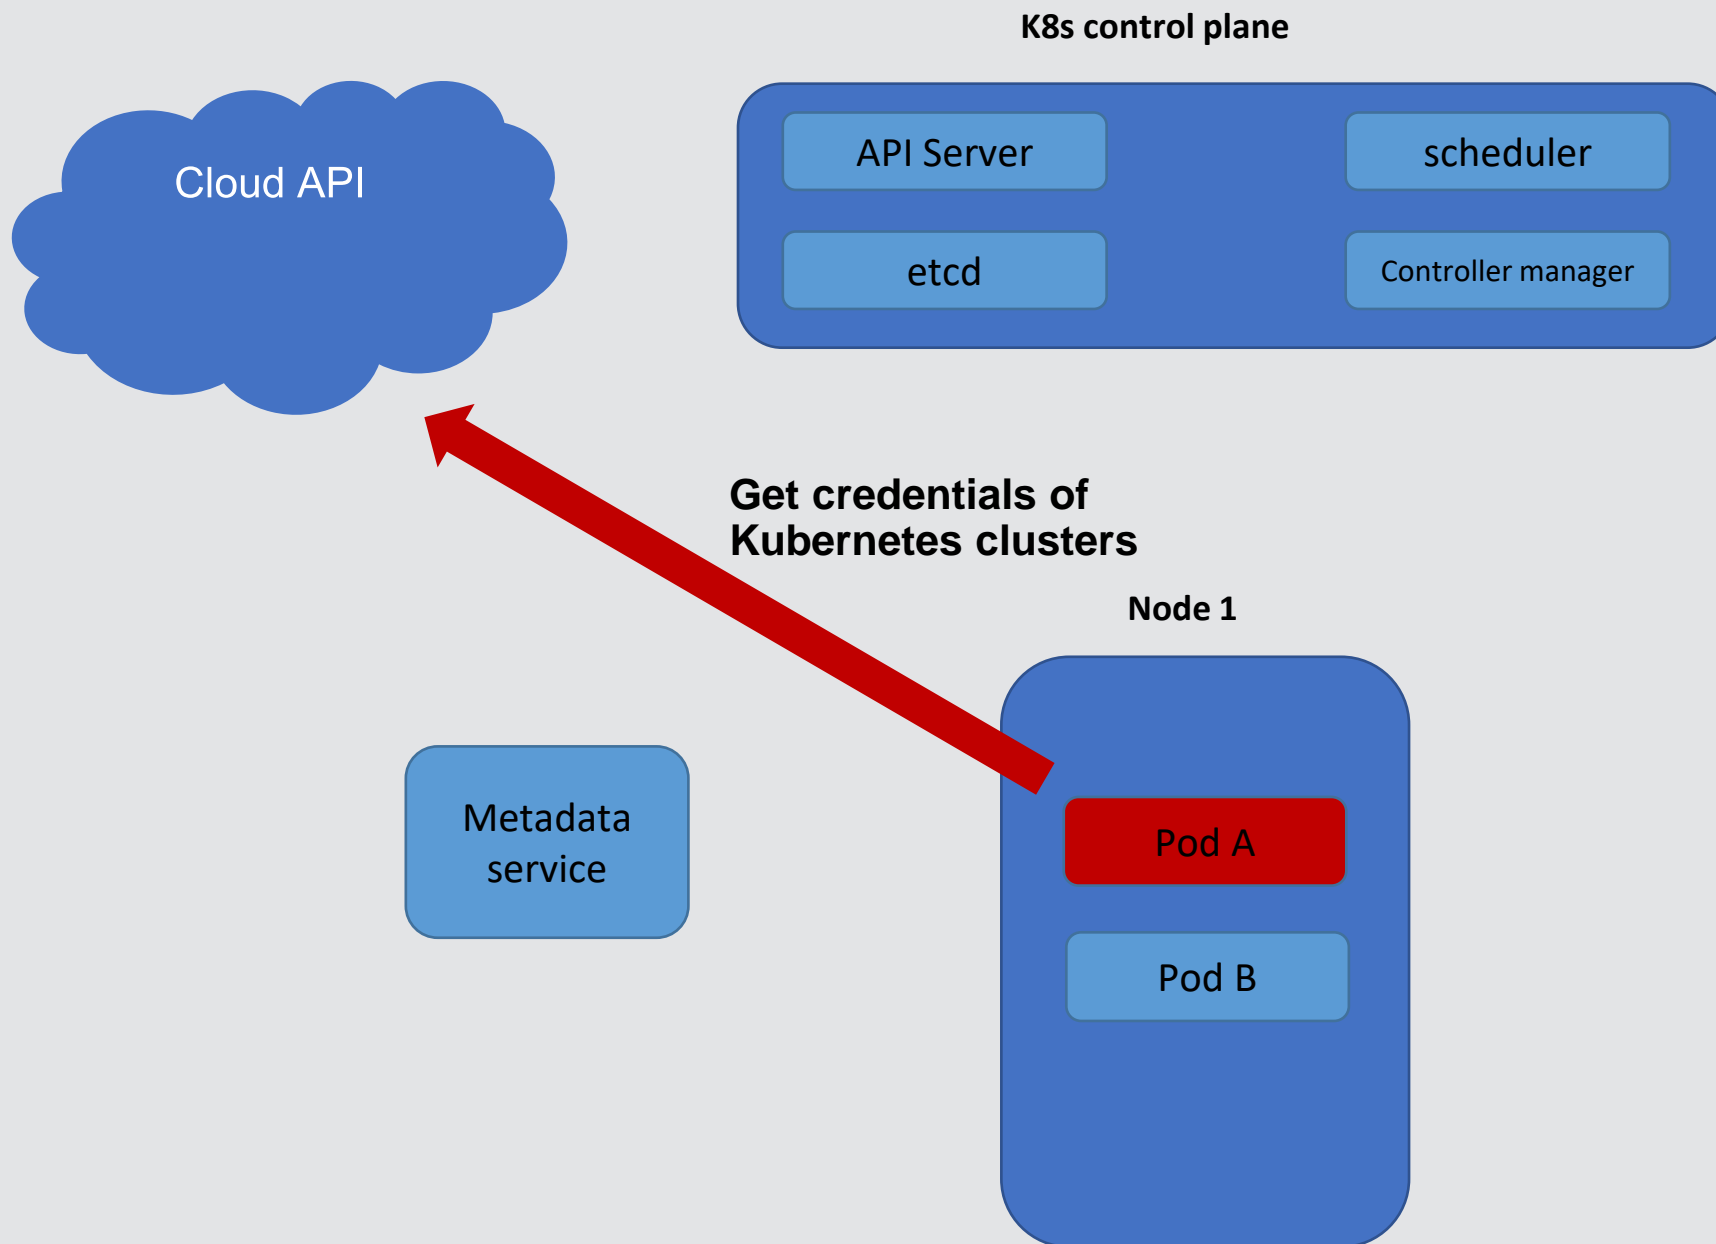


Node 1









Cluster-to-cloud lateral movement

Direct access to IMDS



The problem

- Pods can freely access to their node's cloud identities.
- All pods share the same cloud identities (the node's identities).

What we want:

- Allocate a specific identity to each pod (that needs access to cloud resources) with the minimal needed permissions.
- Make sure pods can only acquire tokens for their own identities.

Cluster-to-cloud lateral movement

Indirect access to IMDS



Cluster-to-cloud lateral movement

Indirect access to IMDS




- Users allocate different identities to the various applications in the cluster.
- When pods query IMDS, the traffic is intercepted and redirected to a local server in the cluster.
- The local server is K8s-aware, thus can identify the querying pod.
- The local server queries IMDS on behalf of the pod and request the pod-specific identity.
- This concept was implemented in Azure by AAD Pod Identity [in deprecation].


K8s control plane



1. Requests to the metadata server are intercepted and sent to the NMI pod (by modifying the IPTables of the node).
2. NMI recognize the querying pod and retrieves its token from the metadata server.
3. NMI send back the token to the pod




K8s control plane



Limitations

1. Works only for Linux containers (uses IPTables).
2. Not supported by all Kubernetes network configuration.

Node 1



Cluster-to-cloud lateral movement

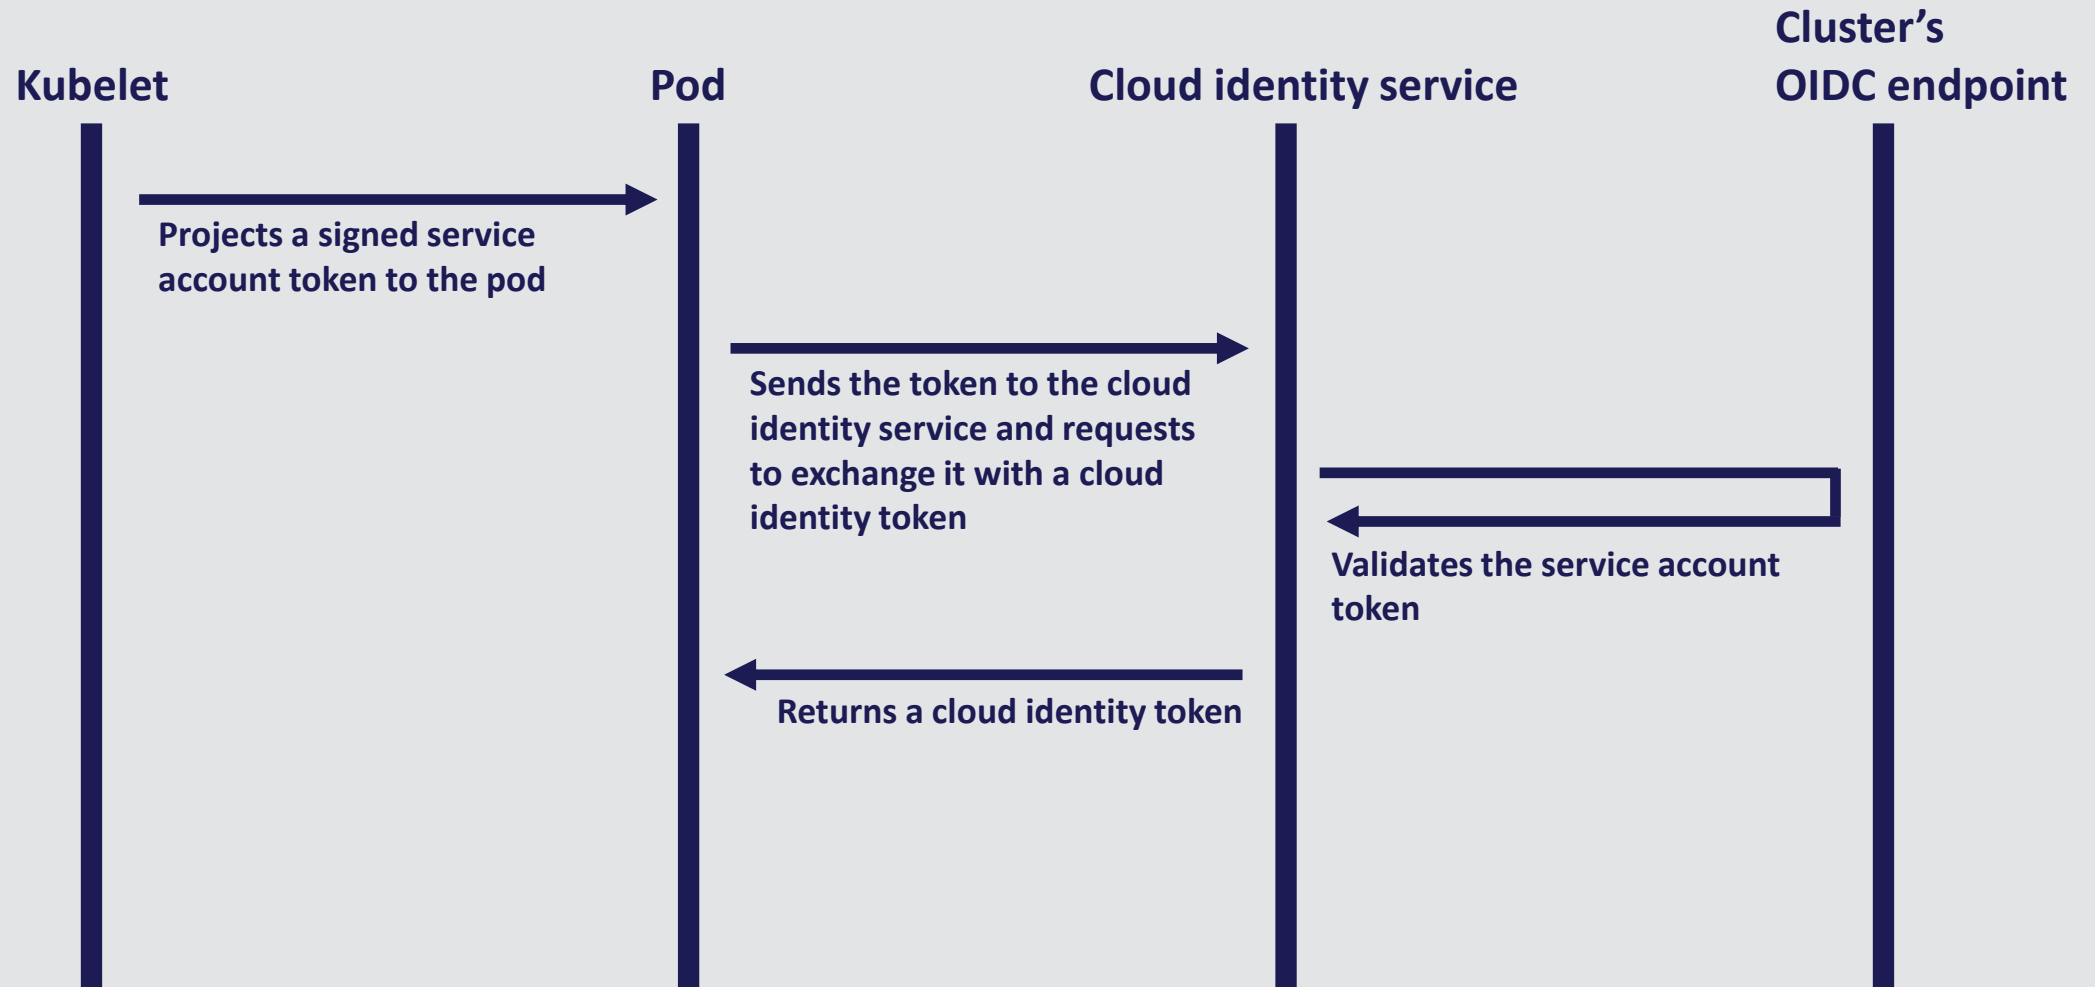
Using OIDC (identity federation)



- Implemented by all major cloud providers:
 - Azure: AAD workload identity
 - AWS: IAM roles for service accounts (IRSA)
 - GCP: Workload identity
- The Kubernetes cluster is used as an OIDC identity provider (IdP).
- Trust relation is created such as the cloud identity platform (AAD, AWS IAM, GCP IAM) trusts the service accounts issued by the K8s cluster.
- This trust relation allows applications in the cluster to exchange a K8s service account token with a cloud identity token.

Cluster-to-cloud lateral movement

Using OIDC (identity federation) – General flow



Cluster-to-cloud lateral movement

Using OIDC (identity federation)

Example: Azure



Binding of a Kubernetes service
account to an AAD application (federation)

Microsoft Azure (Preview) Search resources, services, and docs (G+/)

Dashboard > workload-identity-test-app | Certificates & secrets >

Edit a credential

Configure an Azure AD managed identity or an identity from an external OpenID Connect Provider to get tokens as this application and access Azure resources.

Federated credential scenario * Kubernetes accessing Azure resources

Connect your Kubernetes service account

Please enter the details of the Kubernetes cluster that you want to connect to Azure Active Directory. These values will be used by Azure AD to validate the connection and should match your Kubernetes OIDC configuration. Issuer has a limit of 600 characters. Subject Identifier is a calculated field with a 600 character limit.

Cluster issuer URL *	<input type="text" value="https://oidc.prod-aks.azure.com/240bc132-ba1c-9ba5-fe1a-ce8c3169c320/"/>
Namespace *	<input type="text" value="default"/>
Service account name *	<input type="text" value="workload-identity-sa"/>
Subject identifier	<input type="text" value="system:serviceaccount:default:workload-identity-sa"/>


This value is generated based on the Kubernetes account details provided. [Edit \(optional\)](#)

Cluster-to-cloud lateral movement

Using OIDC (identity federation)

GCP

- In GCP, there's a unified identity pool for the entire project.
- Meaning, there's a single binding of a K8s service account (namespace + SA name) to a cloud identity.

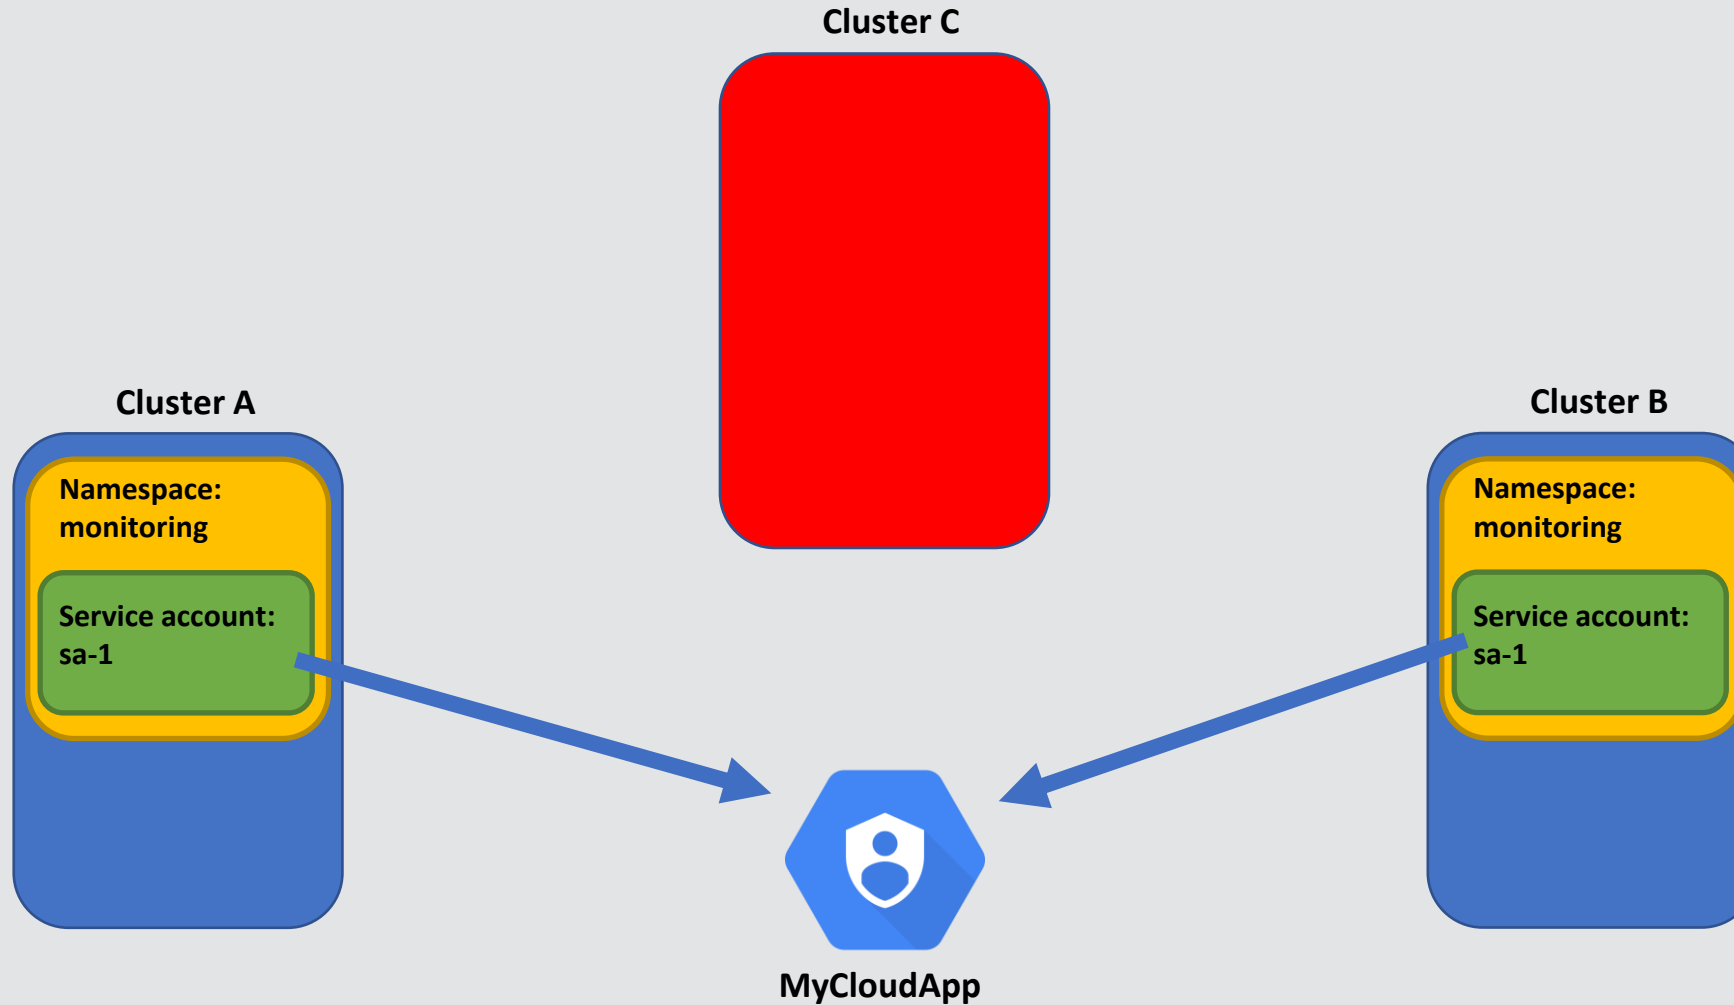


```
gcloud iam service-accounts add-iam-policy-binding MyCloudApp@My-GCP-Project.iam.gserviceaccount.com \
--role roles/iam.workloadIdentityUser \
--member "serviceAccount:My-GCP-Project.svc.id.goog[monitoring/sa-1]"
```

Cluster-to-cloud lateral movement

Using OIDC (identity federation)

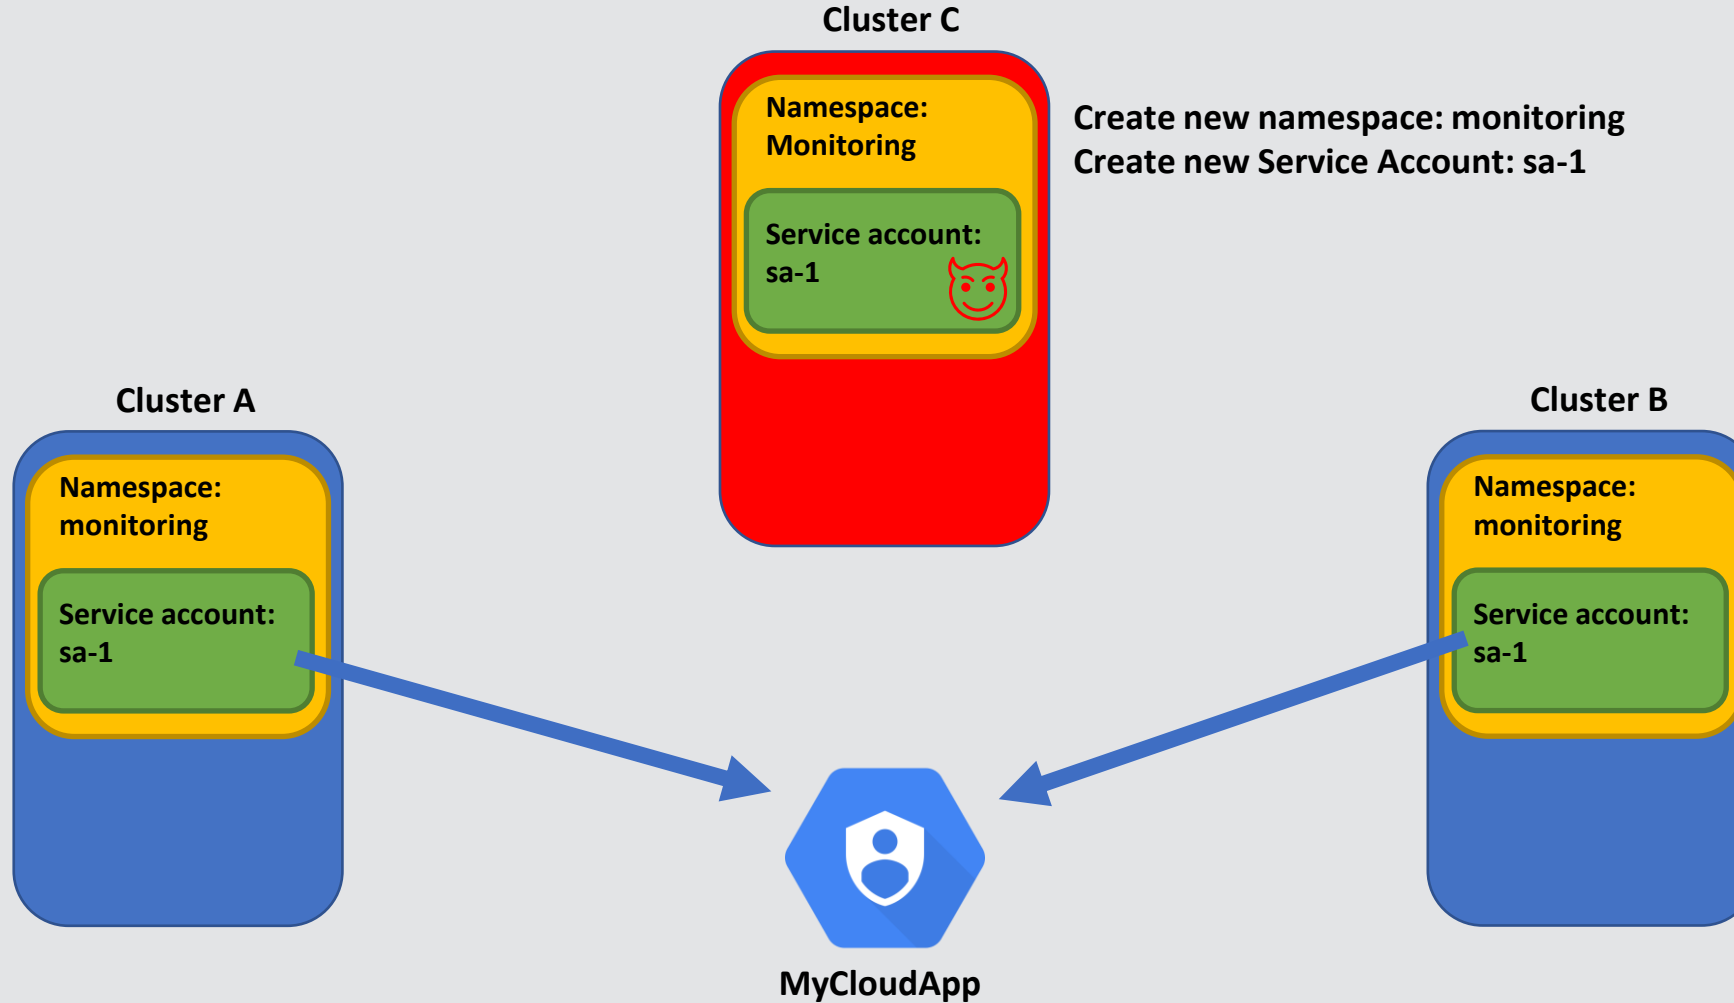
GCP



Cluster-to-cloud lateral movement

Using OIDC (identity federation)

GCP

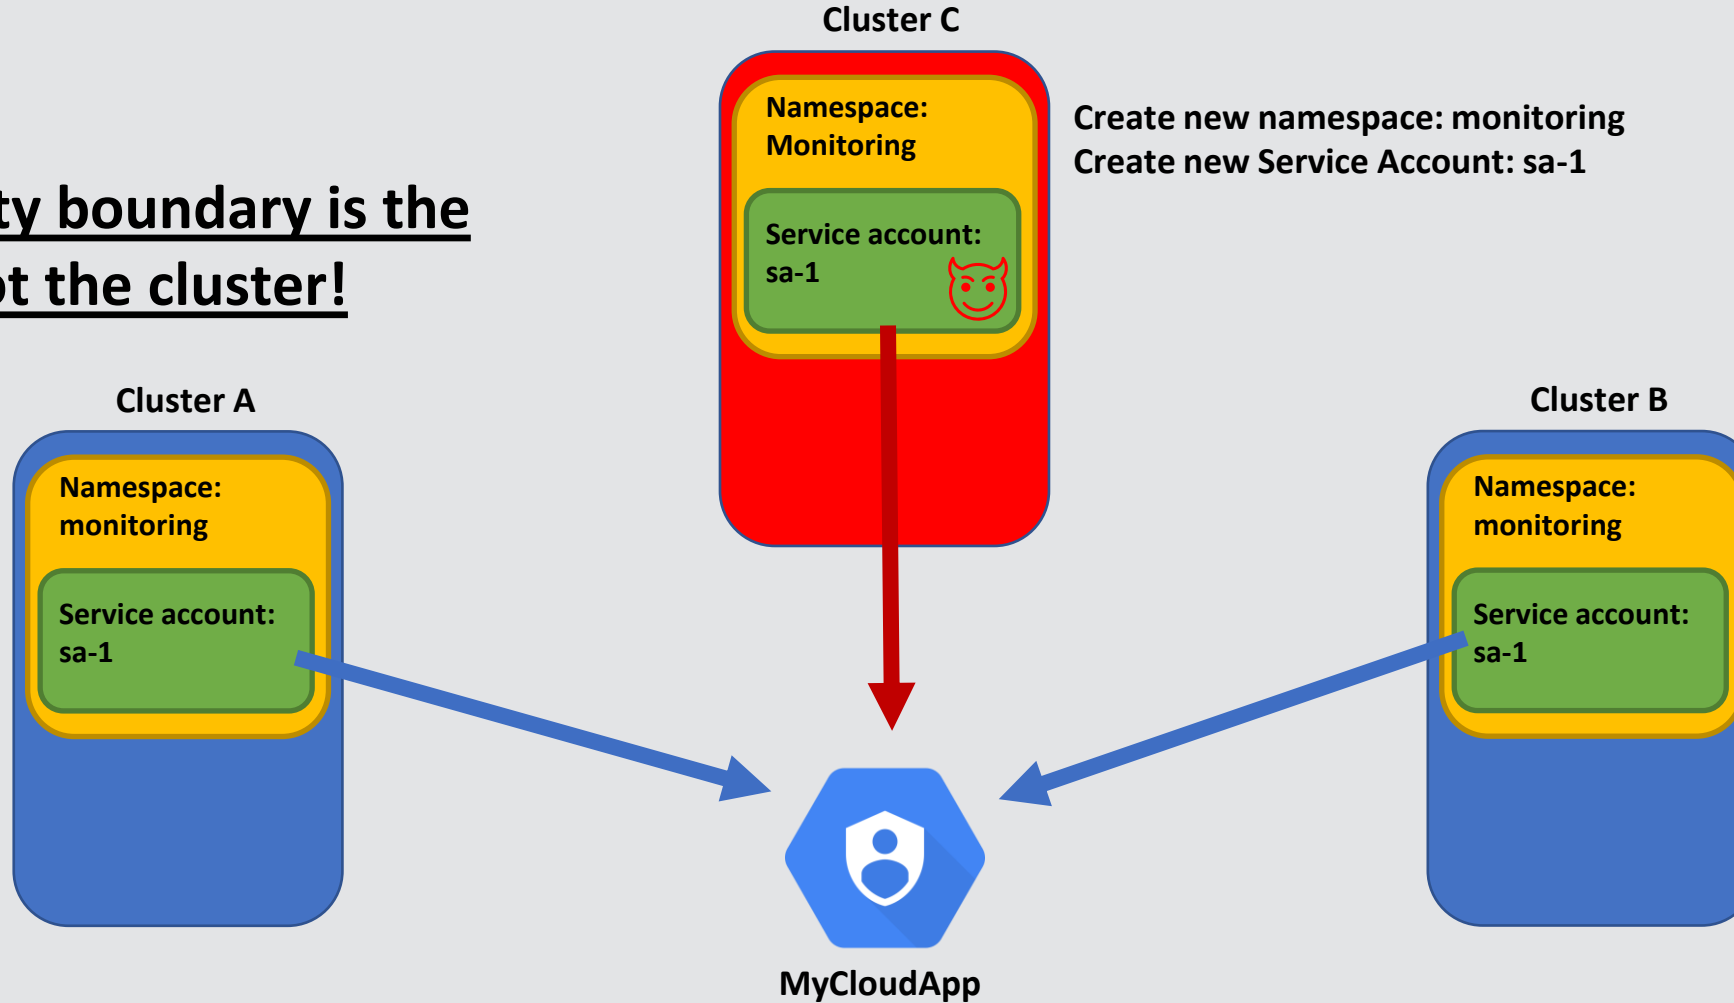


Cluster-to-cloud lateral movement

Using OIDC (identity federation)

GCP

The security boundary is the project, not the cluster!



Cluster-to-cloud lateral movement



What have we seen so far?

- Inner-cluster lateral movement
 - Example: Self-update permissions that lead to cluster-takeover
 - Additional permissions can lead to cluster takeover (partial list)
- Cluster-to-cloud lateral movement:
 - Storing cloud identity credential on the node
 - Direct access to IMDS
 - Indirect access to IMDS
 - Using OIDC (identity federation)

Detections & mitigations



Detections & mitigations



Detections & mitigations



Kubernetes control plane

Monitor suspicious activity in the cluster using K8s audit log (kube-audit). Examples:

1. Deployment of abnormal images
2. Pods with suspicious configurations (sensitive volume mounts, privileged etc.)
3. Reconnaissance activity (for example: SelfSubjectRulesReview API call).
4. Sensitive API calls, such as “get secret”



Cloud provider control plane

Monitor suspicious activity of cloud identities used by K8s workloads\nodes. Examples:

1. Abnormal behavior of cloud identities. Usually, the cloud identities used by the workloads have a consistent behavior.
2. Suspicious access to sensitive cloud services (e.g. storage, secret store etc.)

Detections & mitigations



Kubernetes control plane

Monitor suspicious activity in the cluster using K8s audit log (kube-audit). Examples:

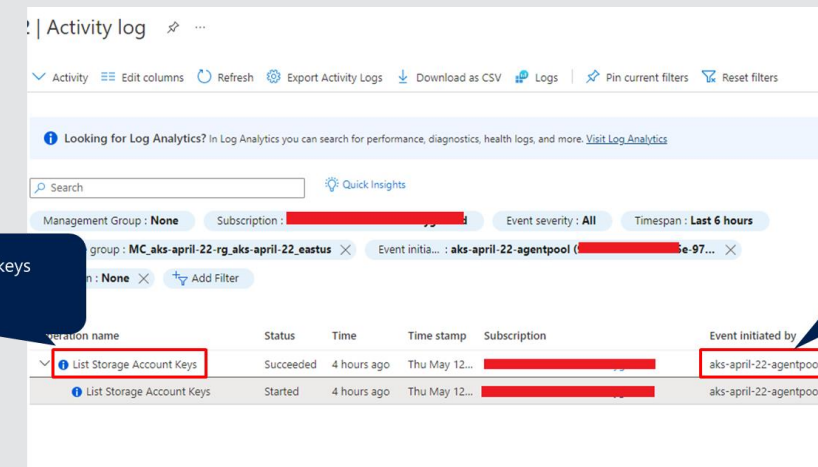
1. Deployment of abnormal images
2. Pods with suspicious configurations (sensitive volume mounts, privileged etc.)
3. Reconnaissance activity (for example: SelfSubjectRulesReview API call).
4. Sensitive API calls, such as “get secret”



Cloud provider control plane

Monitor suspicious activity of cloud identities used by K8s workloads/nodes. Examples:

1. Abnormal behavior of cloud identities. Usually, the cloud identities used by the workloads have a consistent behavior.
2. Suspicious access to sensitive cloud services (e.g. storage, secret store etc.)



Activity log

Activity | Edit columns | Refresh | Export Activity Logs | Download as CSV | Logs | Pin current filters | Reset filters

Looking for Log Analytics? In Log Analytics you can search for performance, diagnostics, health logs, and more. [Visit Log Analytics](#)

Search

Quick Insights

Management Group: None | Subscription: MC_aks-april-22-rg_aks-april-22_eastus | Event severity: All | Timespan: Last 6 hours

Event initiated by: aks-april-22-agentpool

Operation name	Status	Time	Time stamp	Subscription	Event initiated by
List Storage Account Keys	Succeeded	4 hours ago	Thu May 12...	MC_aks-april-22-rg_aks-april-22_eastus	aks-april-22-agentpool
List Storage Account Keys	Started	4 hours ago	Thu May 12...	MC_aks-april-22-rg_aks-april-22_eastus	aks-april-22-agentpool

List storage keys

Managed identity used by Kubernetes node

Azure: Activity Log

AWS: CloudTrail

GCP: Cloud Audit Logs

Detections & mitigations



- In December, a new version of the Threat Matrix for Kubernetes was released (v3).
- The new version now includes mitigation techniques.
- <http://aka.ms/KubernetesThreatMatrix>

Detections & mitigations



Tactics

- Initial Access >
- Execution >
- Persistence >
- Privilege Escalation >
- Defense Evasion >
- Credential Access >
- Discovery >
- Lateral Movement >
- Collection >
- Impact >

Tactics

Initial Access	Execution	Persistence	Privilege Escalation	Defense Evasion	Credential Access	Discovery	Lateral Movement	Collection	Impact
Using cloud credentials	Exec into container	Backdoor container	Privileged container	Clear container logs	List K8S secrets	Access Kubernetes API server	Access cloud resources	Images from a private registry	Data destruction
Compromised image In registry	bash/cmd inside container	Writable hostPath mount	Cluster-admin binding	Delete K8S events	Mount service principal	Access Kubelet API	Container service account	Collecting data from pod	Resource hijacking
Kubeconfig file	New container	Kubernetes CronJob	hostPath mount	Pod / container name similarity	Container service account	Network mapping	Cluster internal networking		Denial of service
Application vulnerability	Application exploit (RCE)	Malicious admission controller	Access cloud resources	Connect from proxy server	Application credentials in configuration files	Exposed sensitive interfaces	Application credentials in configuration files		
Exposed sensitive interfaces	SSH server running inside container	Container service account			Access managed identity credentials	Instance Metadata API	Writable hostPath mount		
	Sidecar injection	Static pods			Malicious admission controller		CoreDNS poisoning		
							ARP poisoning and IP spoofing		

⚠ Disclaimer

The purpose of the Threat Matrix for Kubernetes is to educate readers on the potential of Kubernetes-based tactics, techniques, and procedures (TTPs). It is not to teach how to weaponize or specifically abuse them.

Detections & mitigations

Tactics

- Initial Access >
- Execution >
- Persistence >
- Privilege Escalation >
- Defense Evasion >
- Credential Access >
- Discovery >
- Lateral Movement >
- Collection >
- Impact >

Tactics

Initial Access	Execution	Persistence	Privilege Escalation	Defense Evasion	Credential Access	Discovery	Lateral Movement	Collection	Impact
Using cloud credentials	Exec into container	Backdoor container	Privileged container	Clear container logs	List K8S secrets	Access Kubernetes API server	Access cloud resources	Images from a private registry	Data destruction
Compromised image In registry	bash/cmd inside container	Writable hostPath mount	Cluster-admin binding	Delete K8S events	Mount service principal	Access Kubelet API	Container service account	Collecting data from pod	Resource hijacking
Kubeconfig file	New container	Kubernetes CronJob	hostPath mount	Pod / container name similarity	Container service account	Network mapping	Cluster internal networking		Denial of service
Application vulnerability	Application exploit (RCE)	Malicious admission controller	Access cloud resources	Connect from proxy server	Application credentials in configuration files	Exposed sensitive interfaces	Application credentials in configuration files		
Exposed sensitive interfaces	SSH server running inside container	Container service account			Access managed identity credentials	Instance Metadata API	Writable hostPath mount		
	Sidecar injection	Static pods			Malicious admission controller		CoreDNS poisoning		
							ARP poisoning and IP spoofing		

⚠ Disclaimer

The purpose of the Threat Matrix for Kubernetes is to educate readers on the potential of Kubernetes-based tactics, techniques, and procedures (TTPs). It is not to teach how to weaponize or specifically abuse them.

Detections & mitigations

Tactics

- Initial Access >
- Execution >
- Persistence >
- Privilege Escalation >
- Defense Evasion >
- Credential Access >
- Discovery >
- Lateral Movement >
- Collection >
- Impact >

Tactics

Initial Access	Execution	Persistence	Privilege Escalation	Defense Evasion	Credential Access	Discovery	Lateral Movement	Collection	Impact
Using cloud credentials	Exec into container	Backdoor container	Privileged container	Clear container logs	List K8S secrets	Access Kubernetes API server	Access cloud resources	Images from a private registry	Data destruction
Compromised image In registry	bash/cmd inside container	Writable hostPath mount	Cluster-admin binding	Delete K8S events	Mount service principal	Access Kubelet API	Container service account	Collecting data from pod	Resource hijacking
Kubeconfig file	New container	Kubernetes CronJob	hostPath mount	Pod / container name similarity	Container service account	Network mapping	Cluster internal networking		Denial of service
Application vulnerability	Application exploit (RCE)	Malicious admission controller	Access cloud resources	Connect from proxy server	Application credentials in configuration files	Exposed sensitive interfaces	Application credentials in configuration files		
Exposed sensitive interfaces	SSH server running inside container	Container service account			Access managed identity credentials	Instance Metadata API	Writable hostPath mount		
	Sidecar injection	Static pods			Malicious admission controller		CoreDNS poisoning		
							ARP poisoning and IP spoofing		

⚠ Disclaimer

The purpose of the Threat Matrix for Kubernetes is to educate readers on the potential of Kubernetes-based tactics, techniques, and procedures (TTPs). It is not to teach how to weaponize or specifically abuse them.

Detections & mitigations



Tactics

- Initial Access >
- Execution >
- Persistence >
- Privilege Escalation >
- Defense Evasion >
- Credential Access >
- Discovery >
- Lateral Movement >
- [Access cloud resources](#)
- Container service account
- Cluster internal networking
- Application credentials in configuration files
- Writable hostPath mount
- CoreDNS poisoning
- ARP poisoning and IP spoofing
- Collection >
- Impact >

Access cloud resources

If the Kubernetes cluster is deployed in the cloud, in some cases attackers can leverage their access to a single container to get access to other cloud resources outside the cluster. For example, AKS uses several managed identities that are attached to the nodes, for the cluster operation. Similar identities exist also in EKS and GKE (EC2 roles and IAM service accounts, respectively). By default, running pods can retrieve the identities which in some configurations have privileged permissions. Therefore, if attackers gain access to a running pod in the cluster, they can leverage the identities to access external cloud resources.

Also, AKS has an option to authenticate with Azure using a service principal. When this option is enabled, each node stores service principal credentials that are located in /etc/kubernetes/azure.json. AKS uses this service principal to create and manage Azure resources that are needed for the cluster operation. By default, the service principal has contributor permissions in the cluster's Resource Group. Attackers who get access to this service principal file (by hostPath mount, for example) can use its credentials to access or modify the cloud resources.

Mitigations

ID	Mitigation	Description
MS-M9003	Adhere to least-privilege principle	Grant only necessary permission to the cloud identities.
MS-M9018	Restrict the access of pods to IMDS	Restrict the access of pods to IMDS to restrict pods from getting access to cloud identities.
MS-M9019	Allocate specific identities to pods	Use dedicated allocated identities to pods
MS-M9013	Restrict over permissive containers	Block mounting volumes with access to cloud credentials.

Info

ID: MS-TA9020
Tactic: [Privilege Escalation](#), [Lateral Movement](#)
MITRE technique: [T1078.004](#)

Detections & mitigations

Tactics

- Initial Access >
- Execution >
- Persistence >
- Privilege Escalation >
- Defense Evasion >
- Credential Access >
- Discovery >
- Lateral Movement ▾
 - [Access cloud resources](#)
 - Container service account
 - Cluster internal networking
 - Application credentials in configuration files
 - Writable hostPath mount
 - CoreDNS poisoning
 - ARP poisoning and IP spoofing
- Collection >
- Impact >

Access cloud resources

If the Kubernetes cluster is deployed in the cloud, in some cases attackers can leverage their access to a single container to get access to other cloud resources outside the cluster. For example, AKS uses several managed identities that are attached to the nodes, for the cluster operation. Similar identities exist also in EKS and GKE (EC2 roles and IAM service accounts, respectively). By default, running pods can retrieve the identities which in some configurations have privileged permissions. Therefore, if attackers gain access to a running pod in the cluster, they can leverage the identities to access external cloud resources.

Also, AKS has an option to authenticate with Azure using a service principal. When this option is enabled, each node stores service principal credentials that are located in `/etc/kubernetes/azure.json`. AKS uses this service principal to create and manage Azure resources that are needed for the cluster operation. By default, the service principal has contributor permissions in the cluster's Resource Group. Attackers who get access to this service principal file (by `hostPath` mount, for example) can use its credentials to access or modify the cloud resources.

Mitigations

ID	Mitigation	Description
MS-M9003	Adhere to least-privilege principle	Grant only necessary permission to the cloud identities.
MS-M9018	Restrict the access of pods to IMDS	Restrict the access of pods to IMDS to restrict pods from getting access to cloud identities.
MS-M9019	Allocate specific identities to pods	Use dedicated allocated identities to pods
MS-M9013	Restrict over permissive containers	Block mounting volumes with access to cloud credentials.

Info

ID: MS-TA9020
Tactic: [Privilege Escalation](#), [Lateral Movement](#)
MITRE technique: [T1078.004](#)

Detections & mitigations

Tactics

- Initial Access >
- Execution >
- Persistence >
- Privilege Escalation >
- Defense Evasion >
- Credential Access >
- Discovery >
- Lateral Movement ✓
- [Access cloud resources](#)
- Container service account
- Cluster internal networking
- Application credentials in configuration files
- Writable hostPath mount
- CoreDNS poisoning
- ARP poisoning and IP spoofing
- Collection >
- Impact >

Access cloud resources

If the Kubernetes cluster is deployed in the cloud, in some cases attackers can leverage their access to a single container to get access to other cloud resources outside the cluster. For example, AKS uses several managed identities that are attached to the nodes, for the cluster operation. Similar identities exist also in EKS and GKE (EC2 roles and IAM service accounts, respectively). By default, running pods can retrieve the identities which in some configurations have privileged permissions. Therefore, if attackers gain access to a running pod in the cluster, they can leverage the identities to access external cloud resources.

Also, AKS has an option to authenticate with Azure using a service principal. When this option is enabled, each node stores service principal credentials that are located in `/etc/kubernetes/azure.json`. AKS uses this service principal to create and manage Azure resources that are needed for the cluster operation. By default, the service principal has contributor permissions in the cluster's Resource Group. Attackers who get access to this service principal file (by `hostPath` mount, for example) can use its credentials to access or modify the cloud resources.

Mitigations

ID	Mitigation	Description
MS-M9003	Adhere to least-privilege principle	Grant only necessary permission to the cloud identities.
MS-M9018	Restrict the access of pods to IMDS	Restrict the access of pods to IMDS to restrict pods from getting access to cloud identities.
MS-M9019	Allocate specific identities to pods	Use dedicated allocated identities to pods
MS-M9013	Restrict over permissive containers	Block mounting volumes with access to cloud credentials.

Info

ID: MS-TA9020
Tactic: [Privilege Escalation](#), [Lateral Movement](#)
MITRE technique: [T1078.004](#)

Detections & mitigations



Mitigations

- Multi-factor authentication
- Restrict access to the API server using IP firewall
- Adhere to least-privilege principle
- Secure CI/CD environment
- Image assurance policy
- Enable Just In Time access to API server
- Network intrusion prevention
- Limit access to services over network
- Require strong authentication to services
- Restrict exec commands on pods
- Restrict container runtime using LSM
- Remove tools from container images
- Restrict over permissive containers
- Network segmentation
- Avoid running management interface on containers
- Restrict file and directory permissions
- Ensure that pods meet defined Pod Security Standards
- Restricting cloud metadata API access

[Allocate specific identities to pods](#)

Allocate specific identities to pods

When needed, allocate dedicated cloud identity per pod with minimal permissions, instead of inheriting the node's cloud identity. This prevents other pods from accessing cloud identities that are not necessary for their operation. The features that implement this separation are: Azure AD Pod Identity (AKS), Azure AD Workload identity (AKS), IRSA (EKS) and GCP Workload Identity (GCP).

Techniques Addressed by Mitigation

ID	Name	Use
MS-TA9020	Access cloud resources	Use dedicated allocated identities to pods
MS-TA9028	Access Managed Identity credentials	Allocate specific identities to pods.

Info

ID: MS-M9019
MITRE mitigation: -

Detections & mitigations



Tactics

- Initial Access >
- Execution >
- Persistence >
- Privilege Escalation >
- Defense Evasion >
- Credential Access >
- Discovery >
- Lateral Movement >
- Collection >
- Impact >

Tactics

Initial Access	Execution	Persistence	Privilege Escalation	Defense Evasion	Credential Access	Discovery	Lateral Movement	Collection	Impact
Using cloud credentials	Exec into container	Backdoor container	Privileged container	Clear container logs	List K8S secrets	Access Kubernetes API server	Access cloud resources	Images from a private registry	Data destruction
Compromised image In registry	bash/cmd inside container	Writable hostPath mount	Cluster-admin binding	Delete K8S events	Mount service principal	Access Kubelet API	Container service account	Collecting data from pod	Resource hijacking
Kubeconfig file	New container	Kubernetes CronJob	hostPath mount	Pod / container name similarity	Container service account	Network mapping	Cluster internal networking		Denial of service
Application vulnerability	Application exploit (RCE)	Malicious admission controller	Access cloud resources	Connect from proxy server	Application credentials in configuration files	Exposed sensitive interfaces	Application credentials in configuration files		
Exposed sensitive interfaces	SSH server running inside container	Container service account			Access managed identity credentials	Instance Metadata API	Writable hostPath mount		
	Sidecar injection	Static pods			Malicious admission controller		CoreDNS poisoning		
							ARP poisoning and IP spoofing		

⚠ Disclaimer

The purpose of the Threat Matrix for Kubernetes is to educate readers on the potential of Kubernetes-based tactics, techniques, and procedures (TTPs). It is not to teach how to weaponize or specifically abuse them.

Detections & mitigations



Tactics

- Initial Access >
- Execution >
- Persistence >
- Privilege Escalation >
- Defense Evasion >
- Credential Access >
- Discovery >
- Lateral Movement >
- Collection >
- Impact >

Tactics

Initial Access	Execution	Persistence	Privilege Escalation	Defense Evasion	Credential Access	Discovery	Lateral Movement	Collection	Impact
Using cloud credentials	Exec into container	Backdoor container	Privileged container	Clear container logs	List K8S secrets	Access Kubernetes API server	Access cloud resources	Images from a private registry	Data destruction
Compromised image In registry	bash/cmd inside container	Writable hostPath mount	Cluster-admin binding	Delete K8S events	Mount service principal	Access Kubelet API	Container service account	Collecting data from pod	Resource hijacking
Kubeconfig file	New container	Kubernetes CronJob	hostPath mount	Pod / container name similarity	Container service account	Network mapping	Cluster internal networking		Denial of service
Application vulnerability	Application exploit (RCE)	Malicious admission controller	Access cloud resources	Connect from proxy server	Application credentials in configuration files	Exposed sensitive interfaces	Application credentials in configuration files		
Exposed sensitive interfaces	SSH server running inside container	Container service account			Access managed identity credentials	Instance Metadata API	Writable hostPath mount		
	Sidecar injection	Static pods			Malicious admission controller		CoreDNS poisoning		
							ARP poisoning and IP spoofing		

⚠ Disclaimer

The purpose of the Threat Matrix for Kubernetes is to educate readers on the potential of Kubernetes-based tactics, techniques, and procedures (TTPs). It is not to teach how to weaponize or specifically abuse them.

Detections & mitigations



Tactics

Initial Access >

Execution >

Persistence >

Privilege Escalation >

Defense Evasion >

Credential Access >

Discovery >

Lateral Movement ▾

Access cloud resources

[Container service account](#)

Cluster internal networking

Application credentials in configuration files

Writable hostPath mount

CoreDNS poisoning

ARP poisoning and IP spoofing

Collection >

Impact >

Container service account

Service account (SA) represents an application identity in Kubernetes. By default, a Service Account access token is mounted to every created pod in the cluster and containers in the pod can send requests to the Kubernetes API server using the Service Account credentials. Attackers who get access to a pod can access the Service Account token (located in `/var/run/secrets/kubernetes.io/serviceaccount/token`) and perform actions in the cluster, according to the Service Account permissions. If RBAC is not enabled, the Service Account has unlimited permissions in the cluster. If RBAC is enabled, its permissions are determined by the RoleBindings \ ClusterRoleBindings that are associated with it.

An attacker which get access to the Service Account token can also authenticate and access the Kubernetes API server from outside the cluster and maintain access to the cluster.

Mitigations

ID	Mitigation	Description
MS-M9025	Disable Service Account Auto Mount	Disable service account auto mount.
MS-M9003	Adhere to least-privilege principle	Configure the Kubernetes RBAC such that each service account will have the minimal necessary permissions for the application's functionality.

Info

ID: MS-TA9016

Tactic: [Credential Access](#), [Lateral Movement](#), [Persistence](#)

MITRE technique: [T1528](#)

Detections & mitigations



Tactics

Initial Access >

Execution >

Persistence >

Privilege Escalation >

Defense Evasion >

Credential Access >

Discovery >

Lateral Movement ▾

Access cloud resources

[Container service account](#)

Cluster internal networking

Application credentials in configuration files

Writable hostPath mount

CoreDNS poisoning

ARP poisoning and IP spoofing

Collection >

Impact >

Container service account

Service account (SA) represents an application identity in Kubernetes. By default, a Service Account access token is mounted to every created pod in the cluster and containers in the pod can send requests to the Kubernetes API server using the Service Account credentials. Attackers who get access to a pod can access the Service Account token (located in `/var/run/secrets/kubernetes.io/serviceaccount/token`) and perform actions in the cluster, according to the Service Account permissions. If RBAC is not enabled, the Service Account has unlimited permissions in the cluster. If RBAC is enabled, its permissions are determined by the RoleBindings \ ClusterRoleBindings that are associated with it.

An attacker which get access to the Service Account token can also authenticate and access the Kubernetes API server from outside the cluster and maintain access to the cluster.

Info

ID: MS-TA9016

Tactic: [Credential Access](#), [Lateral Movement](#), [Persistence](#)

MITRE technique: [T1528](#)

Mitigations

ID	Mitigation	Description
MS-M9025	Disable Service Account Auto Mount	Disable service account auto mount.
MS-M9003	Adhere to least-privilege principle	Configure the Kubernetes RBAC such that each service account will have the minimal necessary permissions for the application's functionality.

Detections & mitigations



Mitigations

- Multi-factor authentication
- Restrict access to the API server using IP firewall
- Adhere to least-privilege principle
- Secure CI/CD environment
- Image assurance policy >
- Enable Just In Time access to API server
- Network intrusion prevention
- Limit access to services over network
- Require strong authentication to services
- Restrict exec commands on pods
- Restrict container runtime using LSM
- Remove tools from container images
- Restrict over permissive containers
- Network segmentation
- Avoid running management interface on containers
- Restrict file and directory permissions
- Ensure that pods meet defined Pod Security Standards
- Restricting cloud metadata API access
- Allocate specific identities to

Disable service account auto mount

By default, a service account is mounted to every pod. If the application doesn't require access to the Kubernetes API, disable the service account auto-mount by specifying `automountServiceAccountToken: false` in the pod configuration.

Techniques Addressed by Mitigation

ID	Name	Use
MS-TA9016	Container service account	Disable service account auto mount.

Info

ID: MS-M9025
MITRE mitigation: -

Threat Matrix for Kubernetes



1. Threat Matrix for Kubernetes v1 – April 2020.
 - a. MITRE ATT&CK framework became the standard for EDR evaluations
 - b. OS focused (missing coverage of Cloud-ish TTPs)

Threat Matrix for Kubernetes



1. Threat Matrix for Kubernetes v1 – April 2020.
 - a. MITRE ATT&CK framework became the standard for EDR evaluations
 - b. OS focused (missing coverage of Cloud-ish TTPs)
2. Consensually acceptance by the industry
 - a. Customers
 - b. Competitors
 - c. MITRE (+Microsoft's matrix v2) – April 2021

Threat Matrix for Kubernetes



1. Threat Matrix for Kubernetes v1 – April 2020.
 - a. MITRE ATT&CK framework became the standard for EDR evaluations
 - b. OS focused (missing coverage of Cloud-ish TTPs)
2. Consensually acceptance by the industry
 - a. Customers
 - b. Competitors
 - c. MITRE (+Microsoft's matrix v2) – April 2021
3. v3 – December 2022
 - a. New attack TTPs
 - b. Introducing Mitigation

Key takeaways



1. Implement a holistic strategy for K8s security by considering both the cluster and the cloud levels.
2. Identities are a key aspect of K8s security: monitor their activity using auditing tools.
3. Adhere to the least-privilege principle.
4. Use mitigation measures to prevent potential attacks.

Thank You!



CLOUDNATIVE
SECURITYCON

NORTH AMERICA 2023

Yossi Weizman

 yossi-weizman

Ram Pliskin

 rampliskin

