



# CLOUDNATIVE **SECURITYCON**

**NORTH AMERICA 2023**





CLOUDNATIVE  
**SECURITYCON**

NORTH AMERICA 2023

# Network Security at Scale

*Bernard Van De Walle, Splunk*

*Mitch Connors, Aviatrix*



# Meet Your Speakers



## **Bernard Van De Walle**

Principal software engineer, Splunk  
K8s, Istio, Envoy operations at scale  
Previously at Cruise



## **Mitch Connors**


Principal software engineer, Aviatrix  
Istio Member since 2018  
Istio TOC

splunk® > cloud™

- **Splunk Cloud:** Cloud Native splunk
- Scale: ~ 35 K8s clusters
  - Distributed across all regions
  - AWS and GCP
- Cloud agnostic, K8s, Istio, Envoy



# AGENDA



CLOUDNATIVE  
**SECURITYCON**

NORTH AMERICA 2023

- L3/L4: Cloud Provider
  - VPC
  - Network Load Balancer
- L3/L4: Kubernetes
  - Services
  - Network Policies
- L7: Istio/Envoy
  - AuthN/AuthZ



# Cloud Provider L3/L4



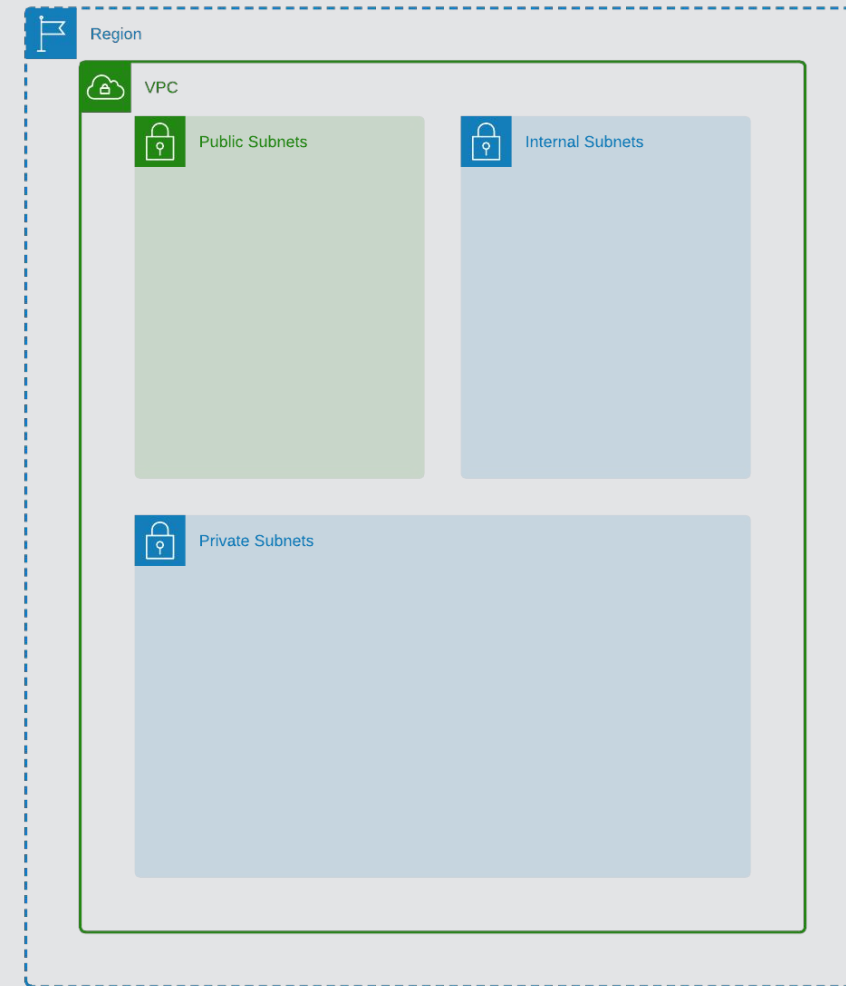
CLOUDNATIVE  
**SECURITYCON**

NORTH AMERICA 2023



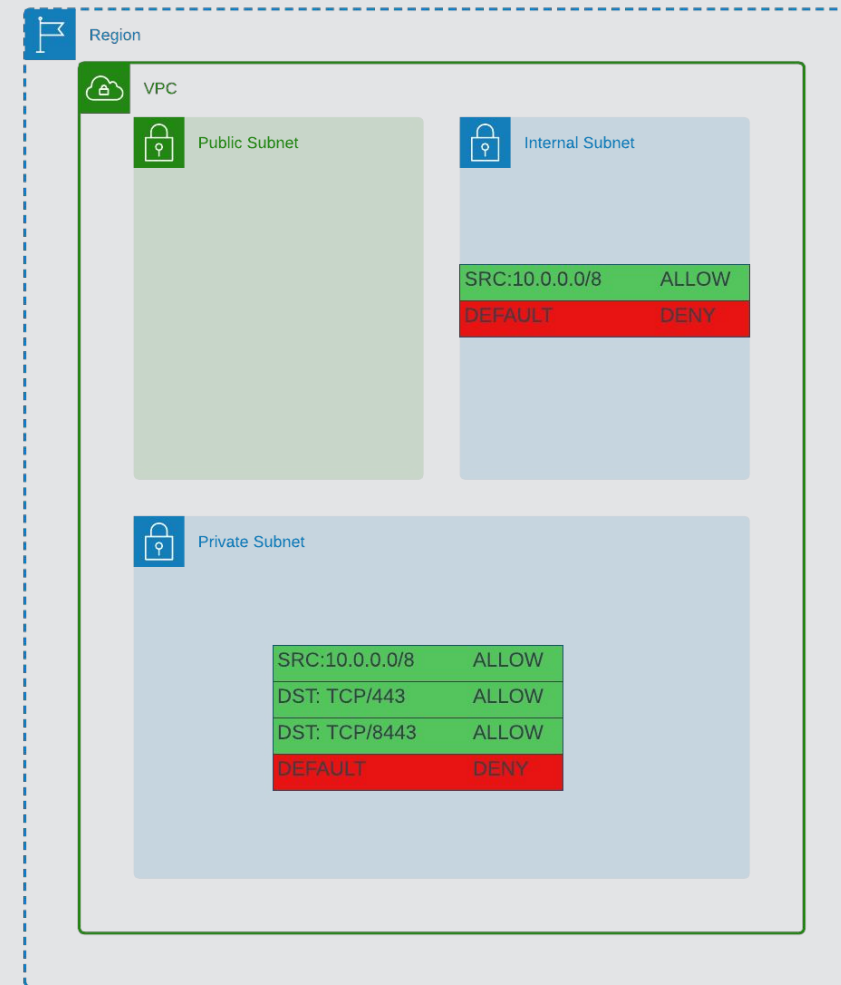
# Standard VPC

- Isolated standard VPC
- 3 set of subnets:
  - **Private** for Kubernetes workloads
  - **Public** for Internet connectivity
  - **Internal** for Splunk connectivity



# Network ACLs

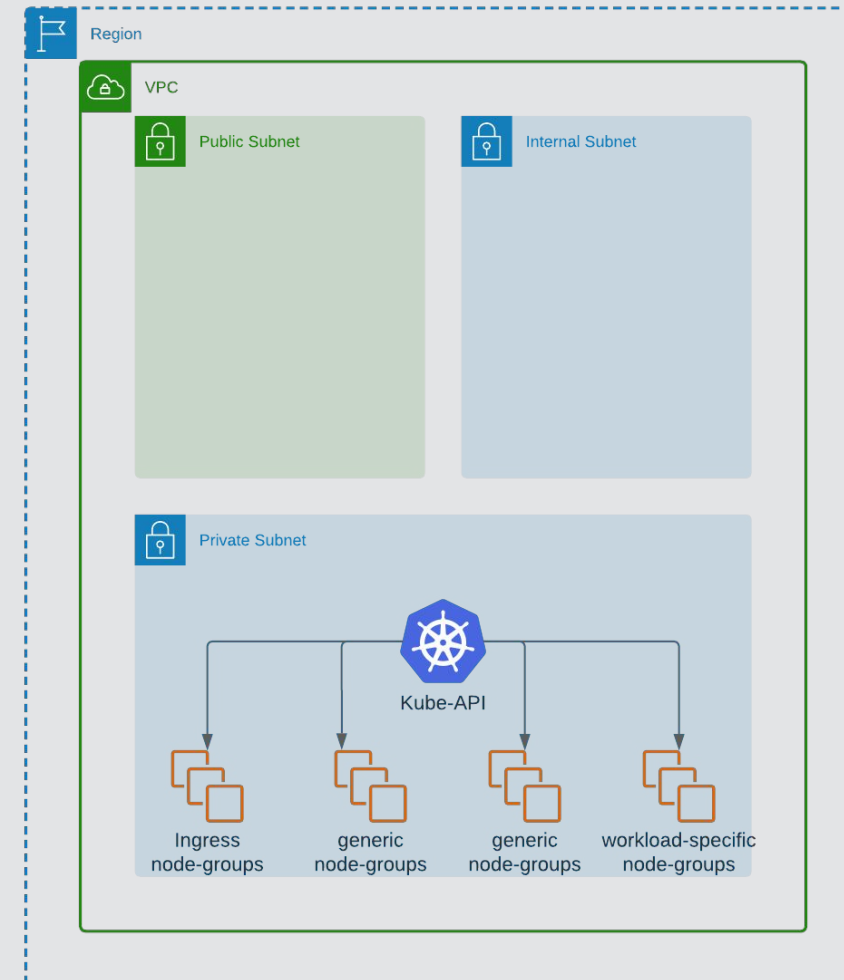
- Stateless
- Applied by subnet
- Basic L3/L4 capabilities
- Provides a catch-all last resort set of rules
- Example:
  - Internal subnet SRC 10.x/8





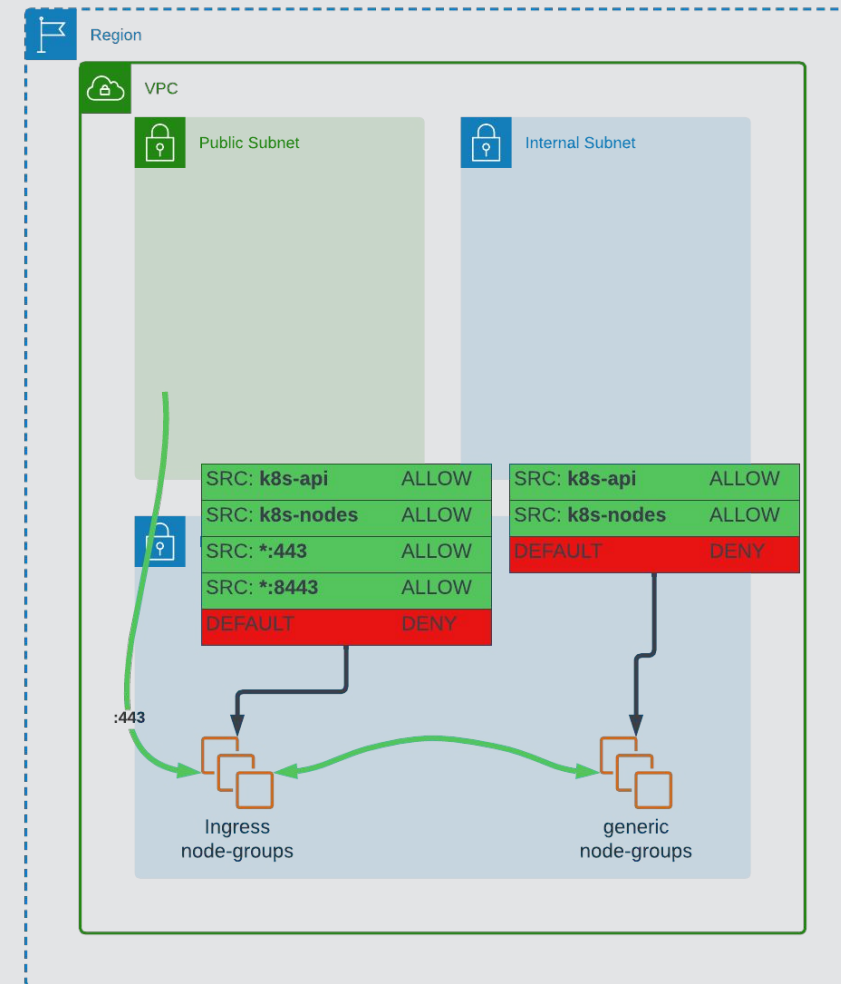
# Kubernetes nodes

- Cluster nodes deployed on private IP space
- Multiple node-groups:
  - Ingress (Gateways)
  - Generic
  - Workload-specific



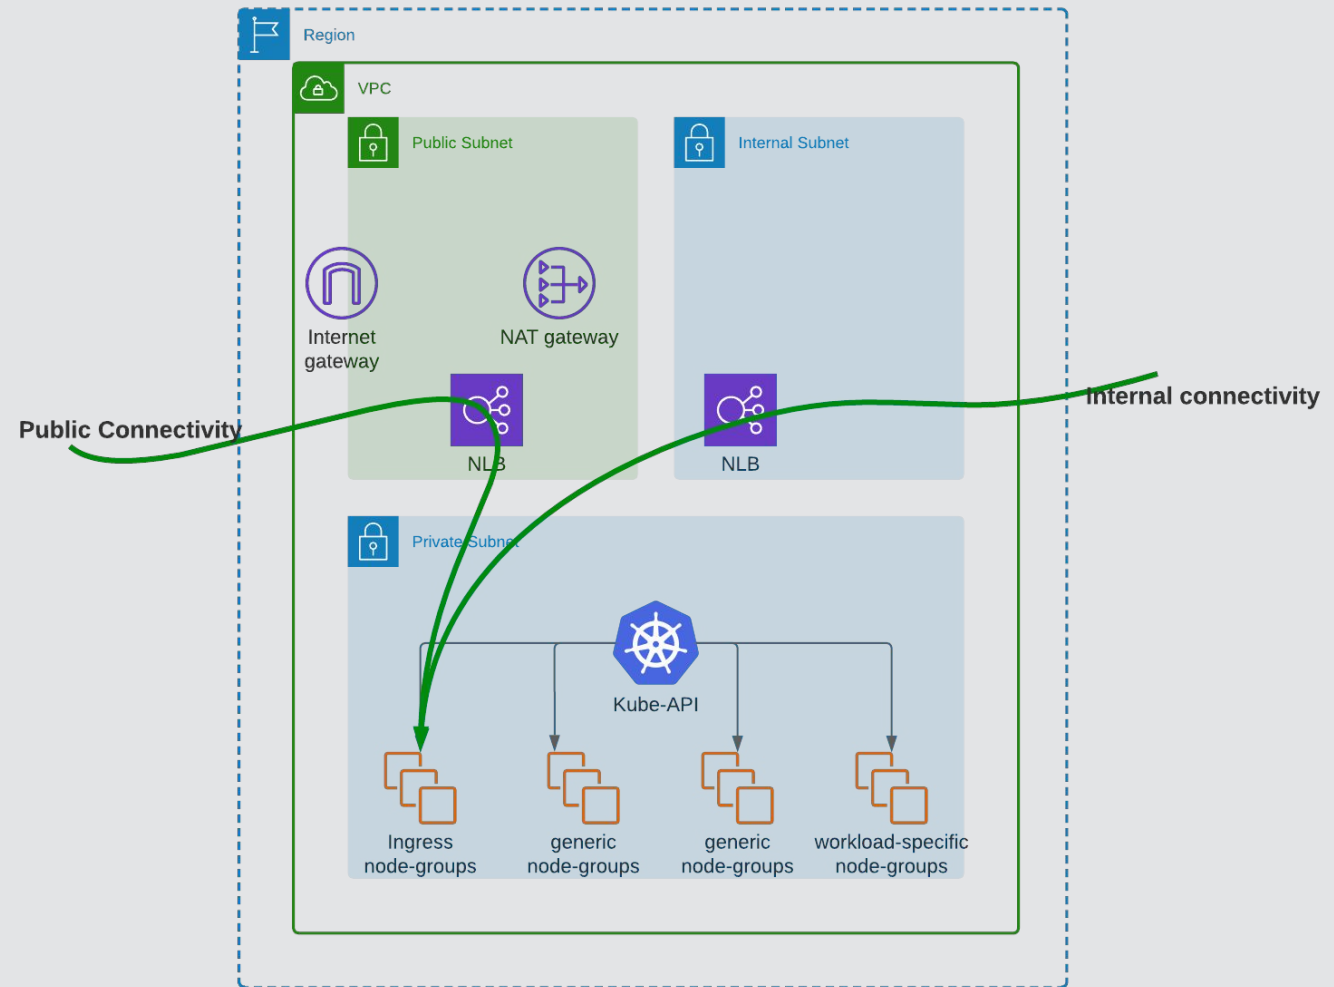
# Security Groups

- Stateful
- Applied per instance
- Allow for fine-grained traffic across specific instances
- Example:
  - Ingress node-group -> Generic node-group



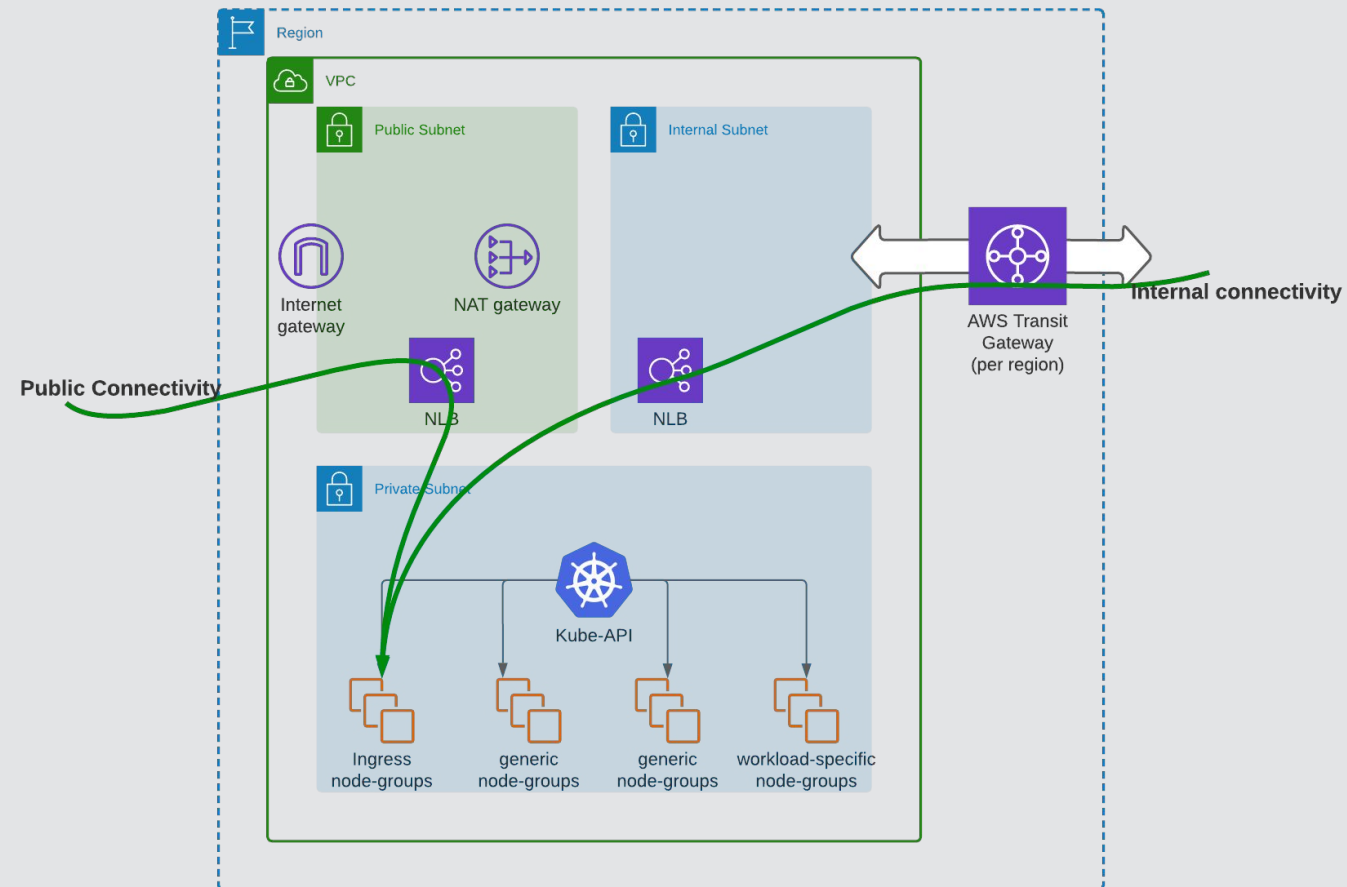
# Ingress connectivity

- Connectivity through NLBs
  - Public NLBs
  - Internal NLBs
- Connectivity through Ingress Gateways (Envoy/Istio)



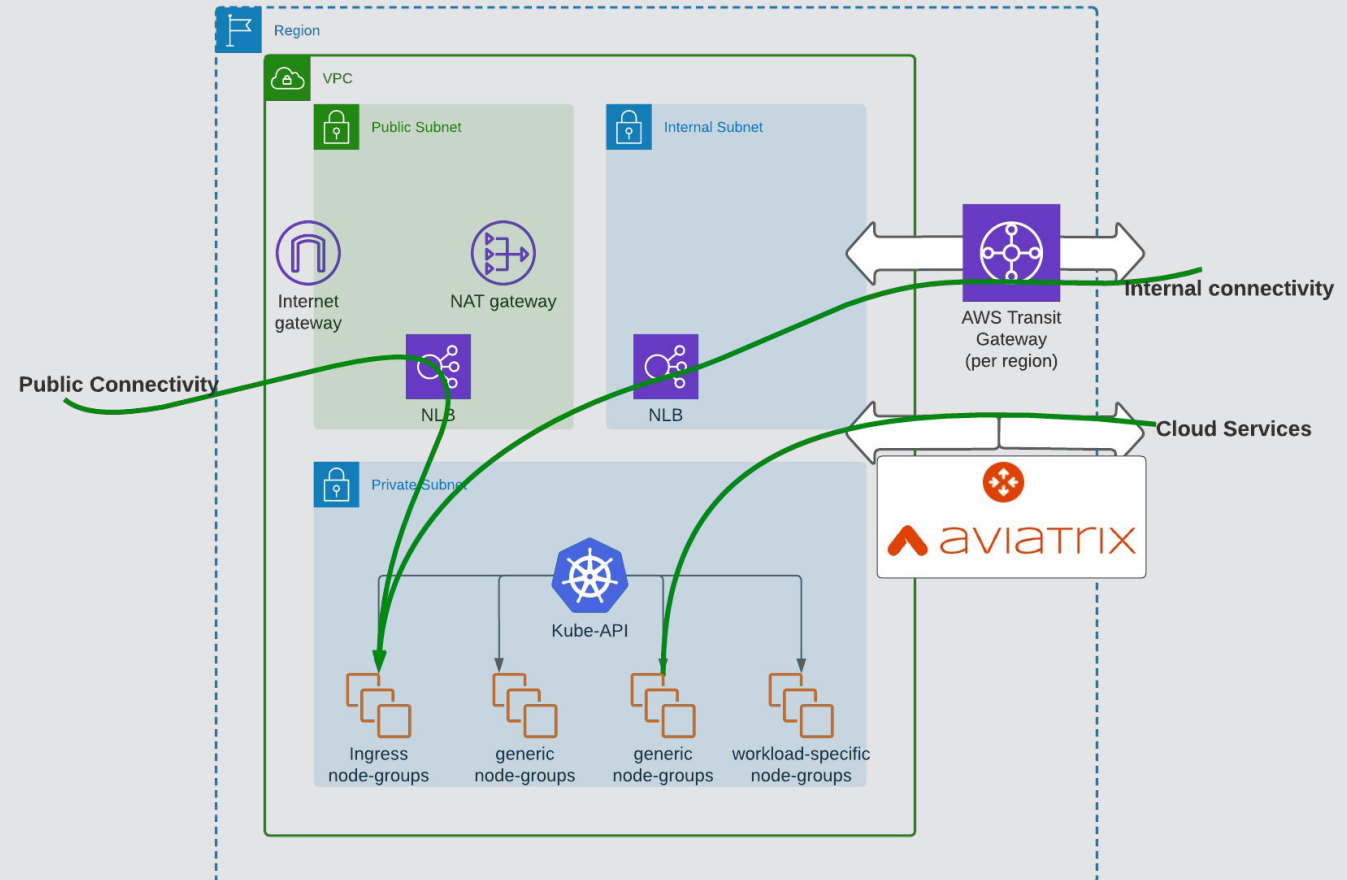
# Internal connectivity (transit gateway)

- Internal connectivity through AWS transit gateway
  - Internal SRC IP advertised only
  - Connected to Splunk firewall



# Cloud connectivity (Aviatrix)

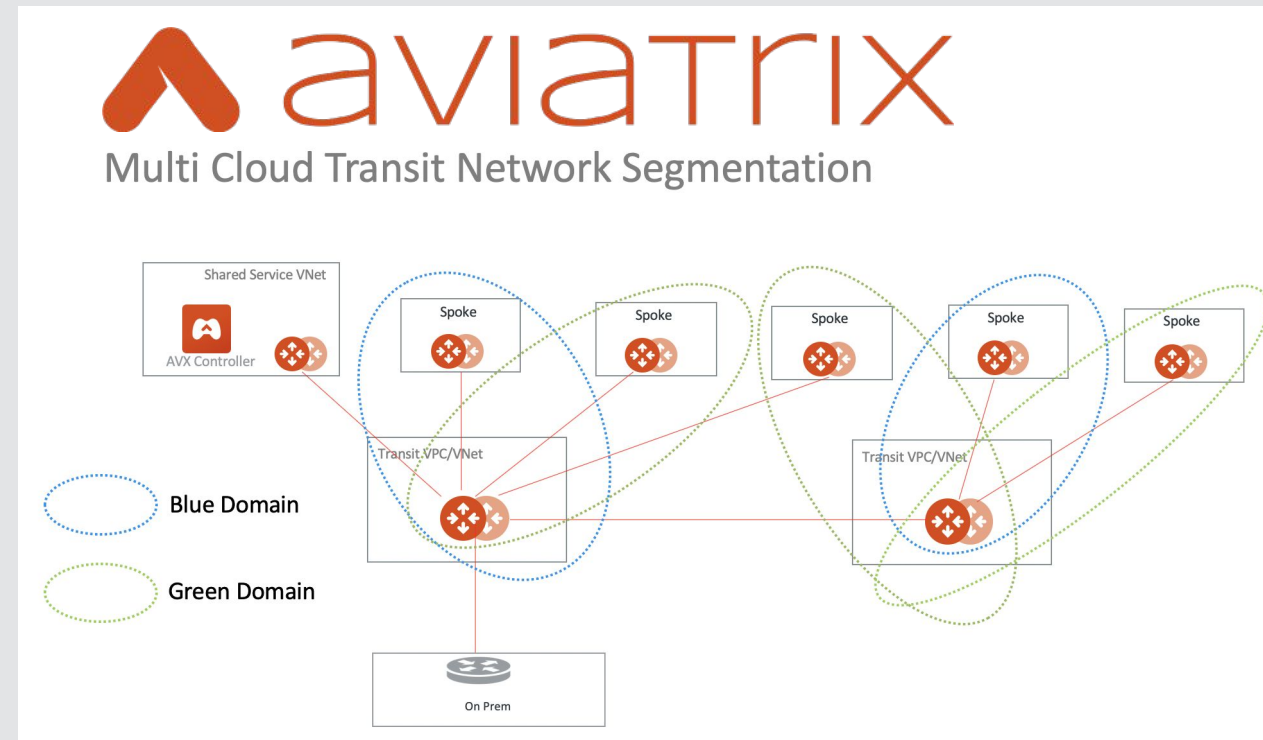
- Cloud services connectivity through Aviatrix
  - Flat network
  - Supports overlapping IPs





# Aviatrix Network Domains

- By Default All connected domains are routable in flat network across clouds, regions, on-prem
- Network Domains limit routes
- Prod can talk to Shared, Dev can talk to Shared, Dev cannot talk to Prod





CLOUDNATIVE  
**SECURITYCON**

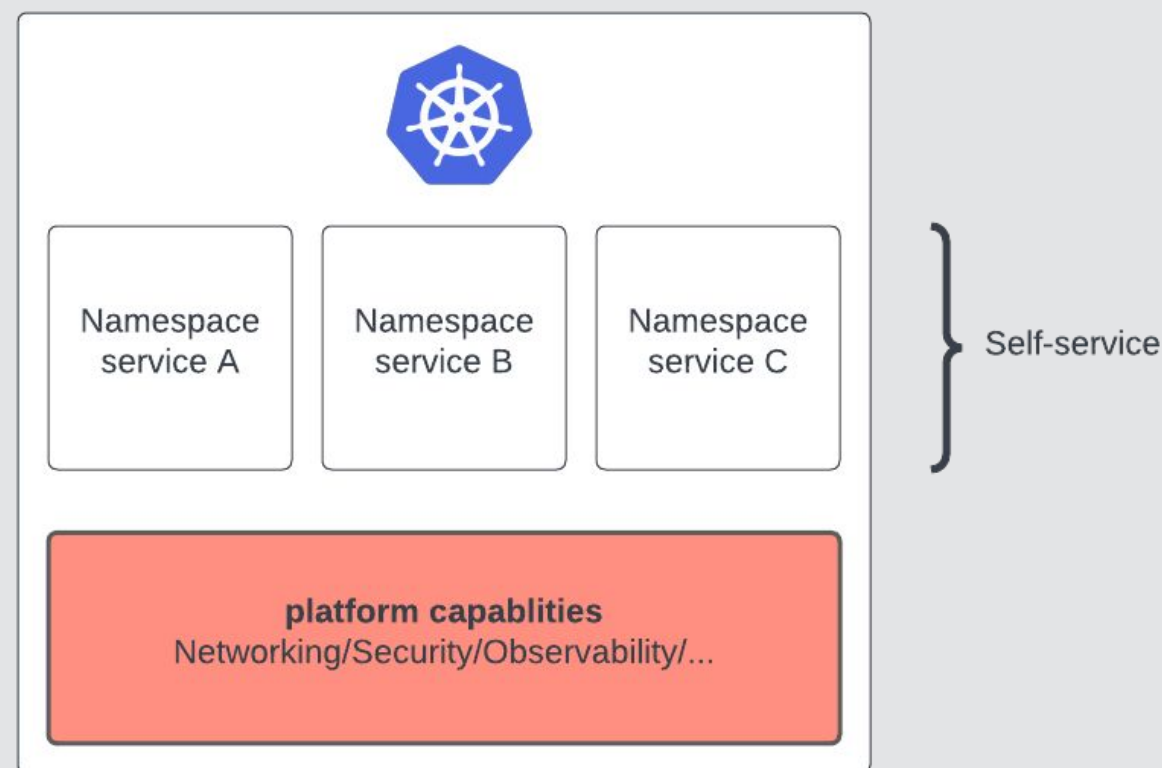
NORTH AMERICA 2023

# Kubernetes L3/L4



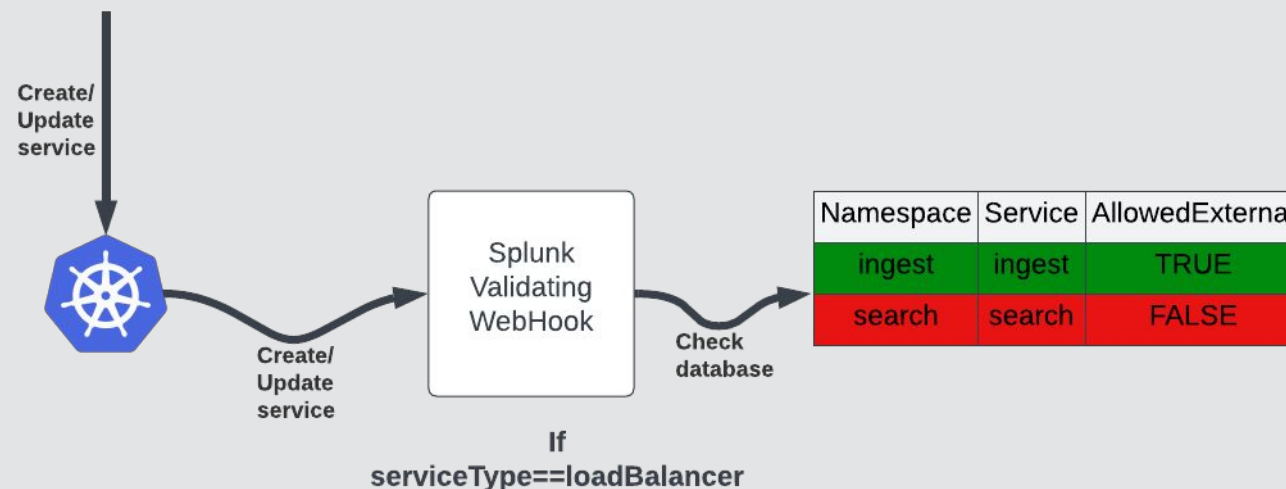
# Kubernetes deployment

- Self-Service platform
- External connectivity **ONLY** through NLBs
- Pod to pod connectivity and Network policies through Calico



# ValidatingWebhook

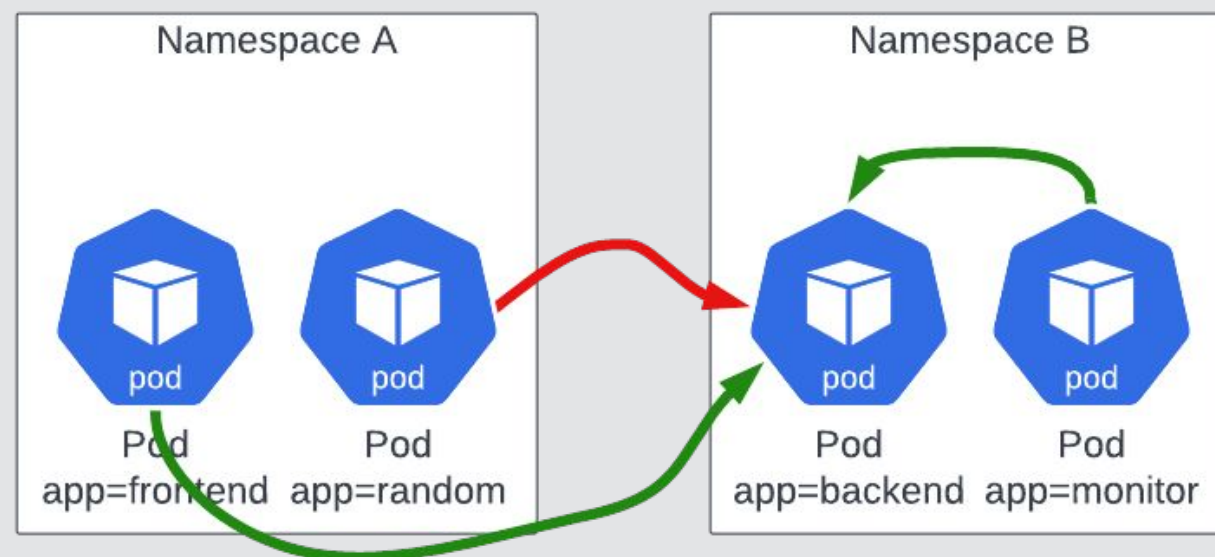
- Splunk validating webhook
- Denies service type load-balancer
- Plenty of open implementations
  - OPA
  - K-Rail (<https://github.com/cruise-automation/k-rail>)



# NetworkPolicies on K8S

- K8S-centric, L3/L4, stateful
- Implemented by your CNI plugin (Cilium, Calico,...)

```
1  apiVersion: networking.k8s.io/v1
2  kind: NetworkPolicy
3  spec:
4    podSelector:
5      matchLabels:
6        app: backend
7    ingress:
8      - from:
9          - podSelector:
10              matchLabels:
11                app: monitor
12          - from:
13              namespaceSelector:
14                matchLabels:
15                  name: frontend
16              podSelector:
17                matchLabels:
18                  app: frontend
```





# Istio L7



CLOUDNATIVE  
**SECURITYCON**

NORTH AMERICA 2023



# Istio Service Mesh

- Application Layer Networking
  - All POST requests from svc A route to this subset of svc B
  - What is the success latency from svc A to svc B?
  - svc A may send GET requests to svc B only at path /foo/\*/bar
- How do we do it?



## Traffic Routing



## Observability



## Security

# Istio Service Mesh



# In-transit encryption

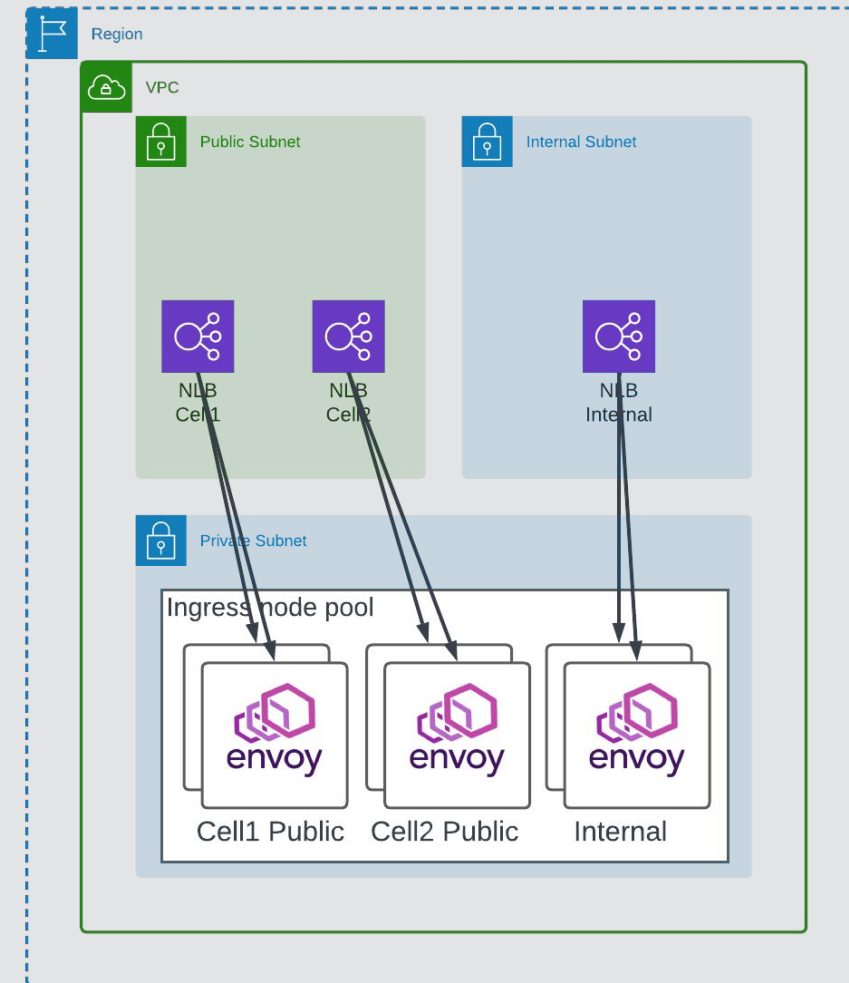
- Run the mesh in **permissive** mode
- Monitor the **passthrough cluster** and alert teams not using mTLS
- Alternative to **strict** mode

Metrics ▾

```
sum(federate:istio_requests_total:sum_rate2m{
  reporter="source",
  k8s_cluster="$cluster",
  destination_service_name="PassthroughCluster"})
by (destination_service)
```

# Gateway provisioning

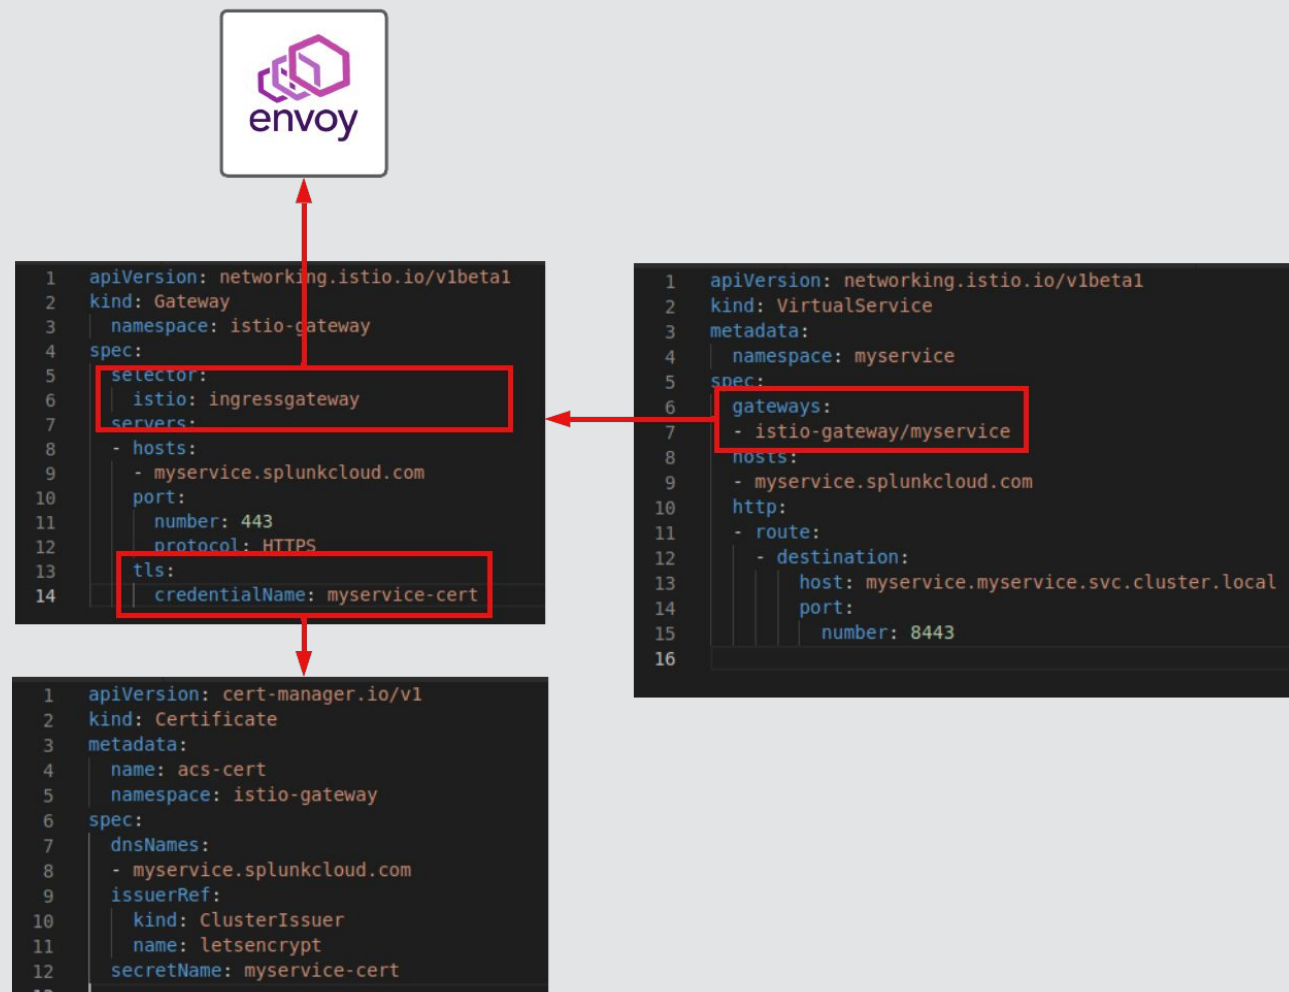
- Ingress only through NLB/gateways
  - For all types of traffic (HTTP/TCP)
- NLB/Gateways per
  - Ingress source (public/internal)
  - Workload types/Blast radius
- Gateway provisioned in **istio-gateway** namespace





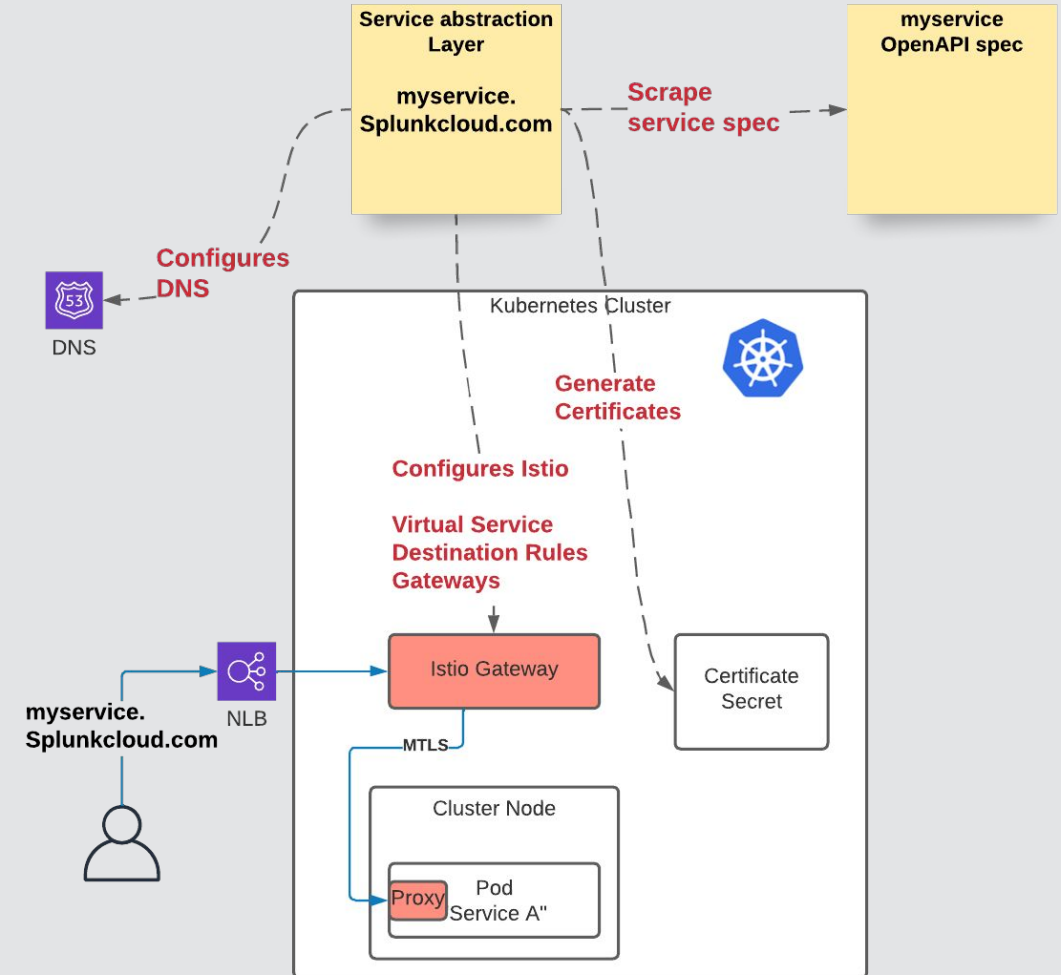
# Ingress setup

- Gateway CRD
  - Istio-gateway namespace
- VirtualService CRD
  - Workload namespace
- Certificate through Let'sEncrypt
- Gated through ValidationWebhook



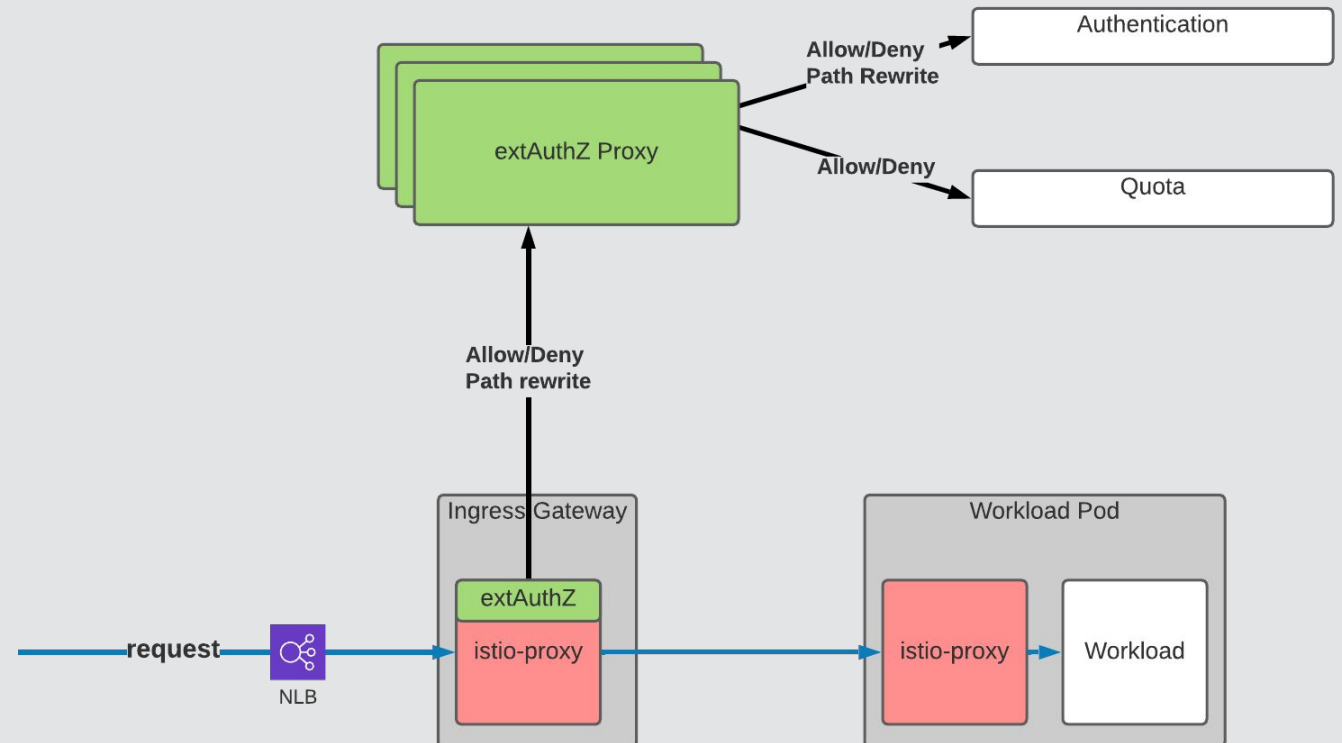
# Service abstraction layer

- “Golden path” abstraction layer for 80% of the use cases
- A single abstraction layer for:
  - **VirtualServices**, **DestinationRules**, **Gateways** and **ServiceEntry** CRD
  - Certificate management
  - DNS management
- OpenAPI spec per service
- Abstraction Layer controller scrapes those openAPI specs



# Layer7 Authentication

- Gateways authenticate requests through envoy ExtAuthZ
- extAuthZ-proxy allows plugins by adding them inline
- Bogus requests are blocked on the gateway
- More info: [External Authorization — envoy 1.26.0-dev-7cc893 documentation](#)



# IP/HTTP AllowList with Istio

- AllowLists on L4-L7
- Mix and match IP and HTTP concepts

```
1  apiVersion: security.istio.io/v1beta1
2  kind: AuthorizationPolicy
3  spec:
4    action: ALLOW
5    rules:
6    - from:
7      - source:
8        ipBlocks:
9        - 1.2.3.4
10     to:
11     - operation:
12       methods:
13       - GET
14       paths:
15       - /myservice/api/v1/*
16   selector:
17     matchLabels:
18     istio: ingressgateway-default
```

# JWT Auth with Istio

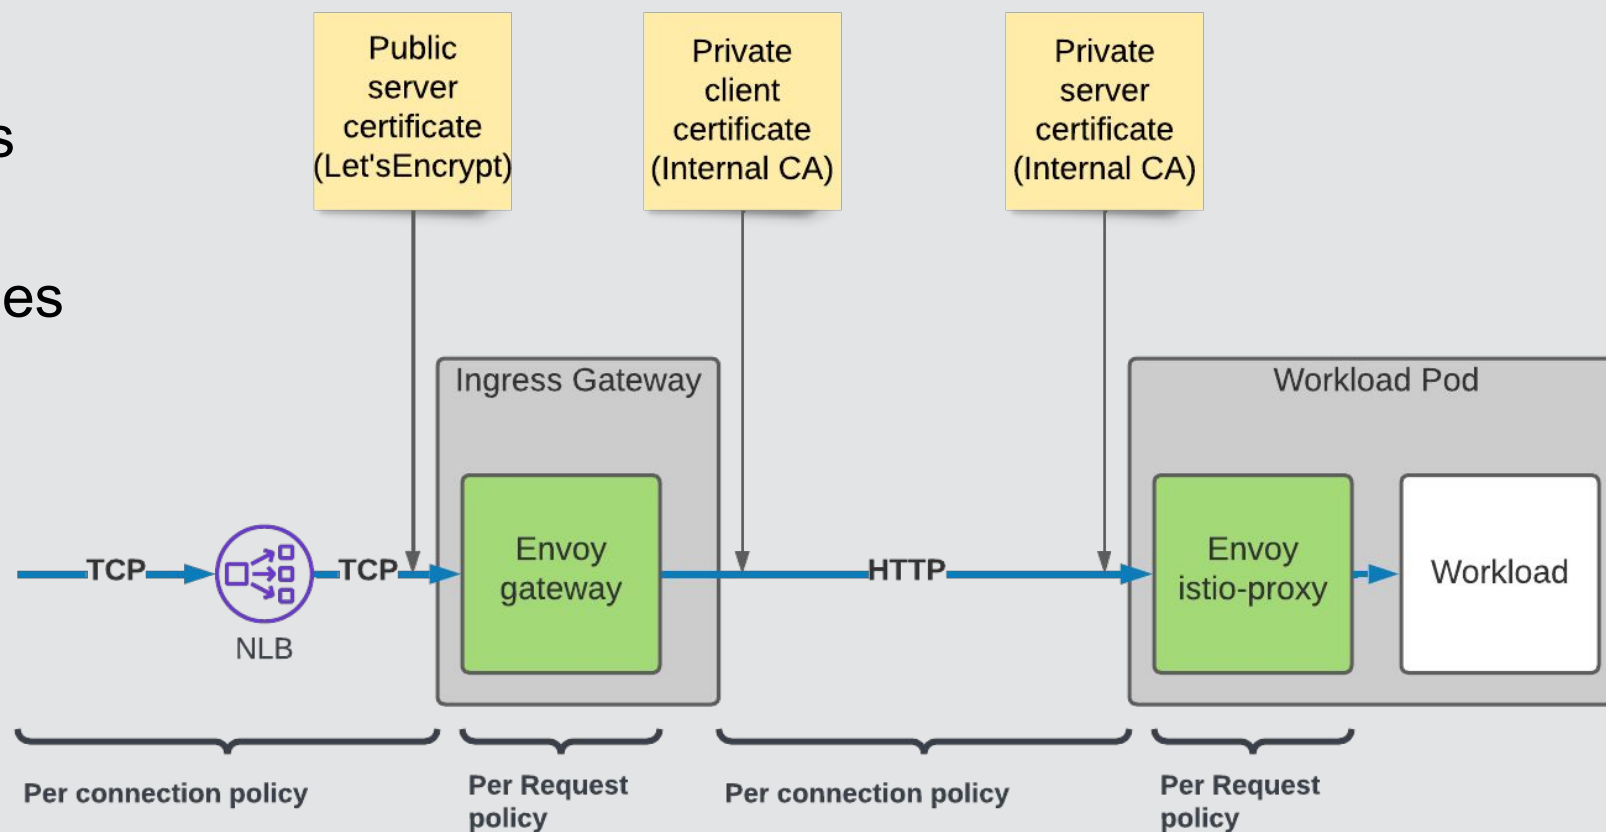
- JWT validation in Istio

```
1  apiVersion: security.istio.io/v1beta1
2  kind: RequestAuthentication
3  metadata:
4    namespace: myservice
5  spec:
6    jwtRules:
7      - issuer: vault.splunkcloud
8        jwks: '{ "keys": [ {"kty": "RSA", "e": "AQAB", "use": "sig", "alg": "RS256", "n": "..."} ] }'
9      selector:
10        matchLabels:
11          app.kubernetes.io/name: myservice
```



# Life of an ingress request

- VPC/K8s apply policies **per connection**
- Istio/Envoy apply policies **per request**





CLOUDNATIVE  
**SECURITYCON**

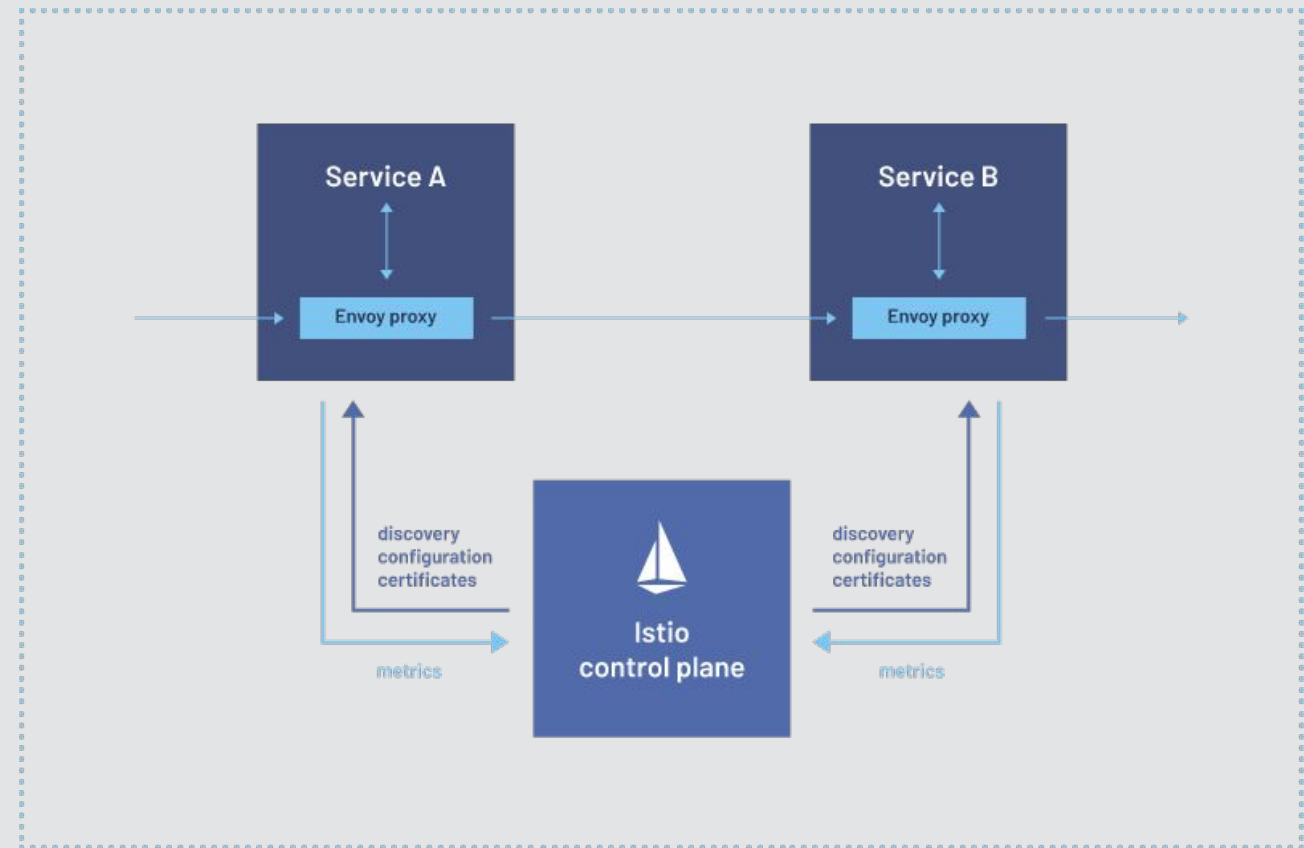
NORTH AMERICA 2023

# Pain Points



# Running One Proxy Per Instance

- Every instance of every application gets a sidecar instance
- Pros: Envoy can control all traffic
- Cons: So. Many. Envoy.
  - Vertical Scaling becomes extremely expensive.



# Managing the Magic

- Pods ≠ Deployment Spec
  - Injection modifies Pods, not Deployments
- Pods ≠ Deployment + Injection
  - Injection only occurs at Pod creation time
- Which version of Envoy?
  - ͇͇(ツ)͇͇

Sidecars can be automatically added to applicable Kubernetes pods using a [mutating webhook admission controller](#) provided by Istio.

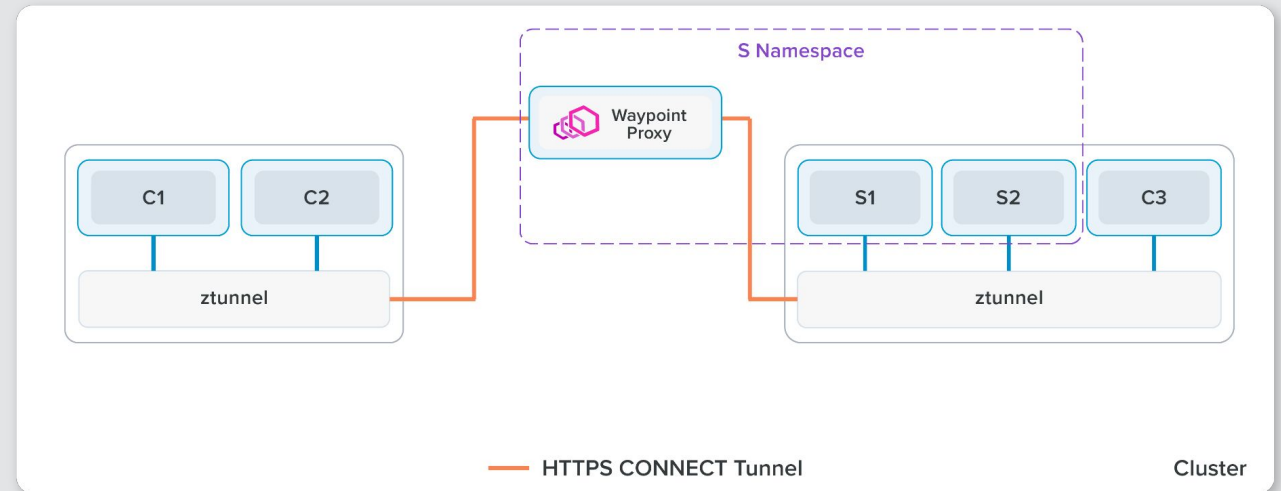
When you set the `istio-injection=enabled` label on a namespace and the injection webhook is enabled, any new pods that are created in that namespace will automatically have a sidecar added to them.

-istio.io

# Ambient Reduces User Pain

- One L4 Proxy per Node
- One scalable L7 Proxy per Service Account + Gateway
- All proxies are managed through Deployments/Daemonsets
- Pods = Deployment Spec
- For more info:

[istio.io/latest/blog/2022/introducing-ambient-mesh/](https://istio.io/latest/blog/2022/introducing-ambient-mesh/)





CLOUDNATIVE  
**SECURITYCON**





NORTH AMERICA 2023

# Conclusion






# Defense at Every Layer

		Identity	Policy	Observability
	VPC Network ACLs	IP/Ports/VM	Network ACLs SecurityGroups	VPC Flow Logs
	Aviatrix 	Network Domains	Network Domains	Copilot Flow Logs
	K8S 	IP/Ports/Pods	NetworkPolicy	Node Telemetry
Strong Identity 	Istio 	ServiceAccount Request Headers	AuthN/AuthZ AuthorizationPolicies RequestPolicies	Istio Telemetry

- **Self-service platforms are hard**
  - Safeguards to avoid users shooting themselves in the foot
  - Provide a Golden Path to avoid configuration errors
- **Defense in depth: Add redundant security at all layers**
- **Observability is key**
  - Help debug
  - Detect misconfiguration



# Thank You!



Please scan the QR Code above  
to leave feedback on this session