

# Package Transparency for WebAssembly Registries

Kyle Brown, SingleStore

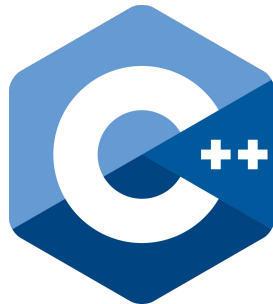
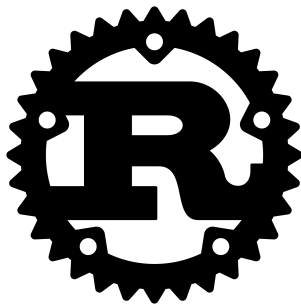
# Overview

- Introduce **WebAssembly (Wasm)**
- What is a **Package Registry**?
- Applying **Certificate Transparency** to **Package Registries**
- **Package Transparency & warg**
- **Package Transparency vs. Various Attacks**

# WebAssembly (Wasm)

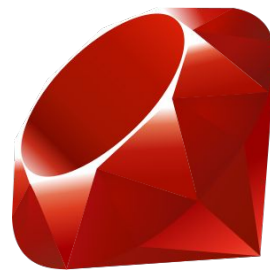
WebAssembly (Wasm) is...

A platform-agnostic “**compile target**” or...  
something you can compile programs to.



Well-supported

Work is in progress for many more languages



Swift



Kotlin



# History



Wasm was created as a web technology



It became a W3C Standard in 2019

supported by major browsers since 2017



Wasm **isn't** *just* for  
the web



# Wasm has really valuable properties

- **Portability**
- **Speed**
  - Low startup latency
  - Near-native performance
- **Security**
  - Capability safety
  - Sandboxing & memory isolation

People using **Wasm** outside the browser



Database Extensibility



Distributed Apps



Edge Computing



Serverless





**BYTECODE**  
**ALLIANCE**

Wasm applications  
and libraries will be  
**composed, shared,**  
and **deployed**

Wasm needs a native  
package registry

# We need a registry as secure as Wasm itself



You wouldn't seal a vault...



with a Cheeto

What is a Package Registry?

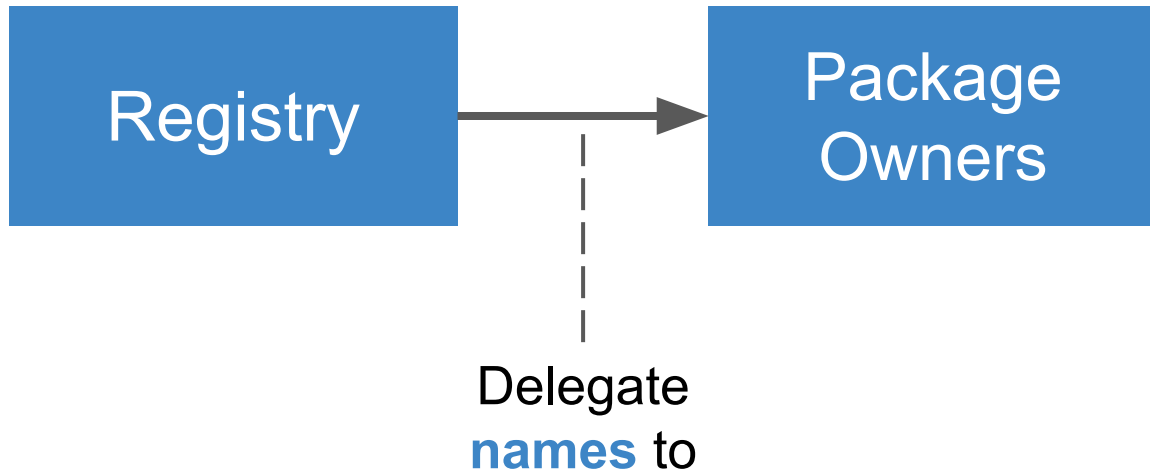
# Package Registries' Role



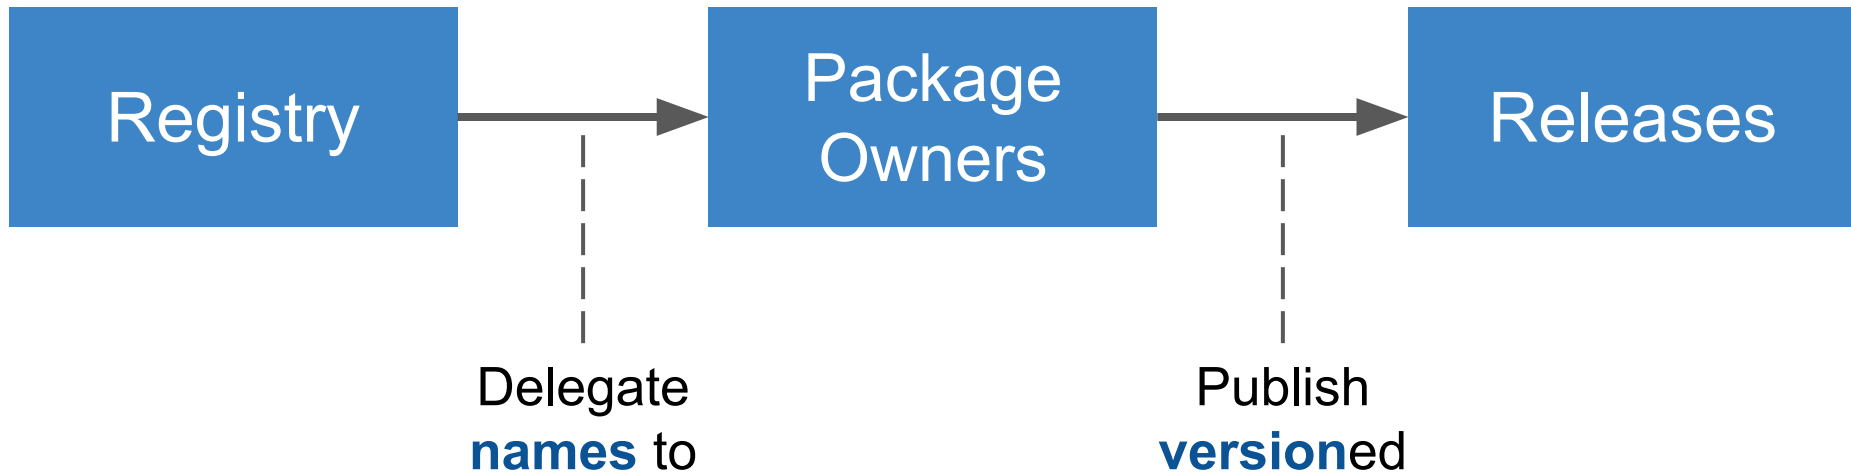
Registry



# Package Registries' Role



# Package Registries' Role



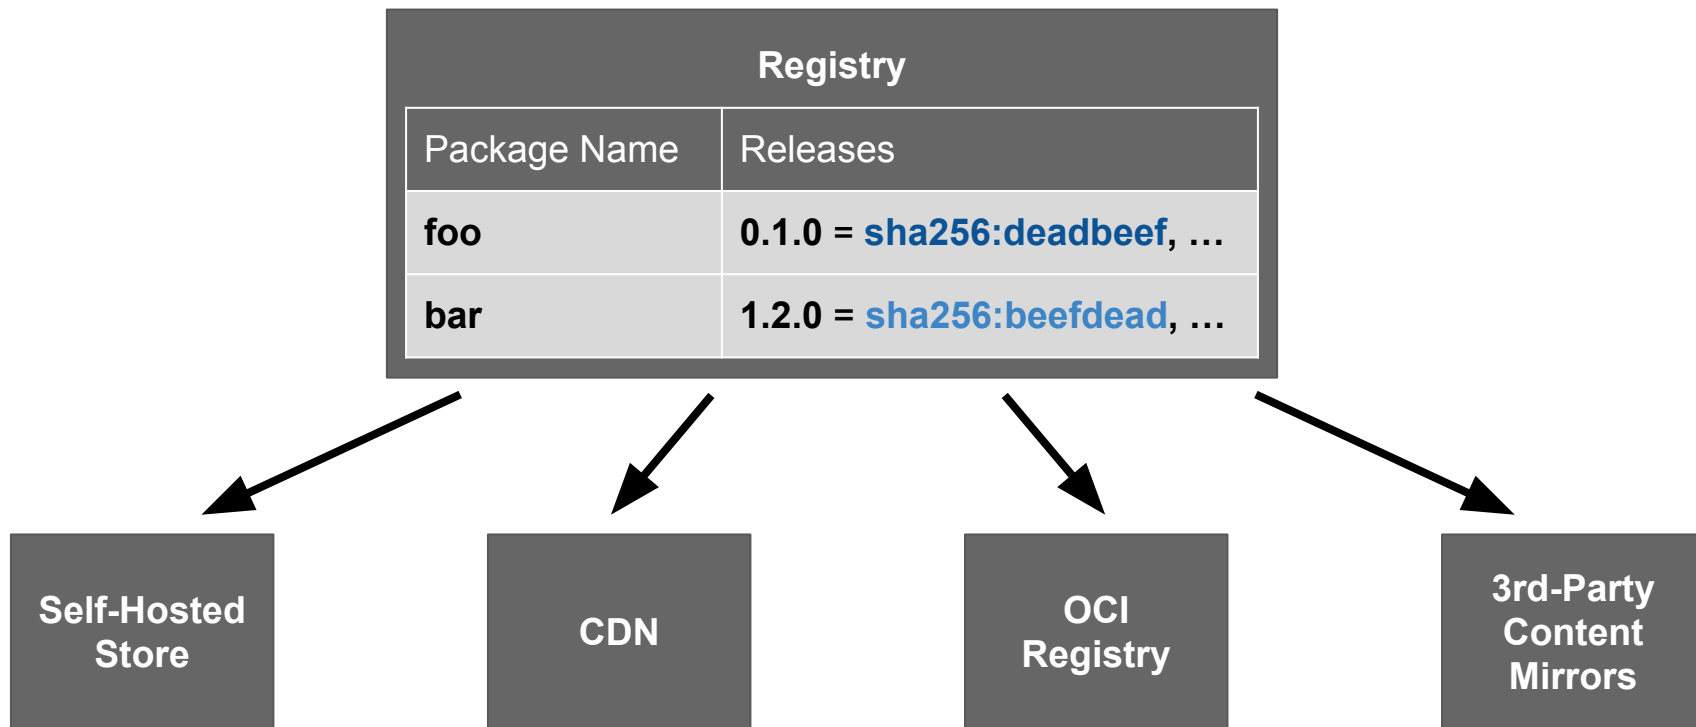
# A Registry is an Index

(name, version) → package digest

# Registries must control metadata

Registry	
Package Name	Releases
foo	0.1.0 = sha256:deadbeef, ...
bar	1.2.0 = sha256:beefdead, ...

# Registries may delegate content hosting



How can we apply  
Certificate Transparency to  
Package Registries?

People are able to  
detect when CAs  
**misissue** certificates

Clients *should* be able to  
detect when registries  
accept **invalid**  
package updates



# Package Transparency

“Package transparency is

**Lann Martin, Fermyon**

“Package transparency is publishing  
cryptographically-verifiable commitments

**Lann Martin, Fermyon**

“Package transparency is publishing  
cryptographically-verifiable commitments  
to the state of a package repository

**Lann Martin, Fermyon**

“Package transparency is publishing cryptographically-verifiable commitments to the state of a package repository to allow auditing of the actions of package authors and the registry itself over time.”

**Lann Martin, Fermyon**

# Components of Package Transparency

1. Publicly Available Package Registry State
2. Cryptographically-Verifiable Commitments
3. Auditing Package Authors and the Registry

**warg** is a protocol for  
Package Transparency

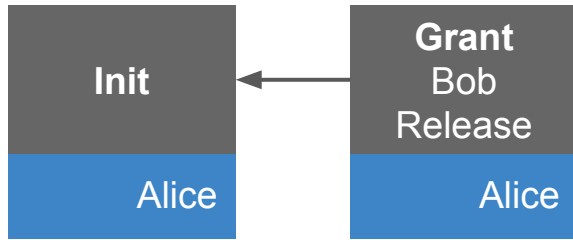
Publicly Available  
Registry State



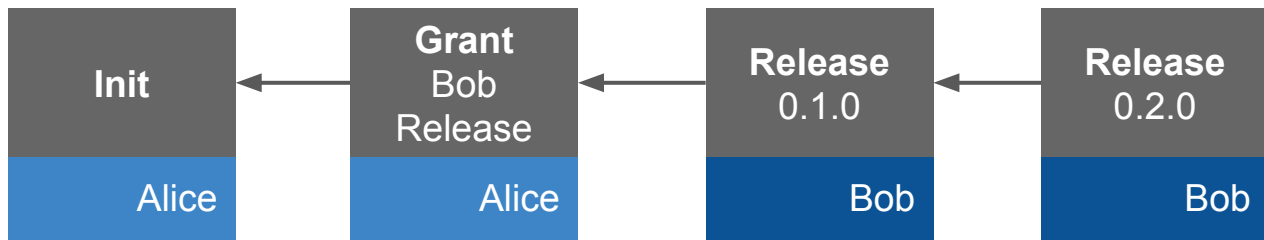
# Append-only Package Logs



# Append-only Package Logs



# Append-only Package Logs



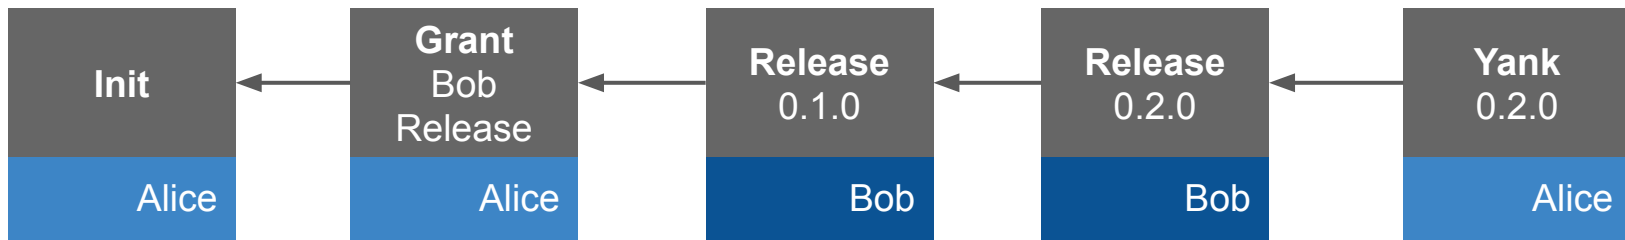
Releases

- 0.1.0

Releases

- 0.1.0
- 0.2.0

# Append-only Package Logs



Releases

- 0.1.0

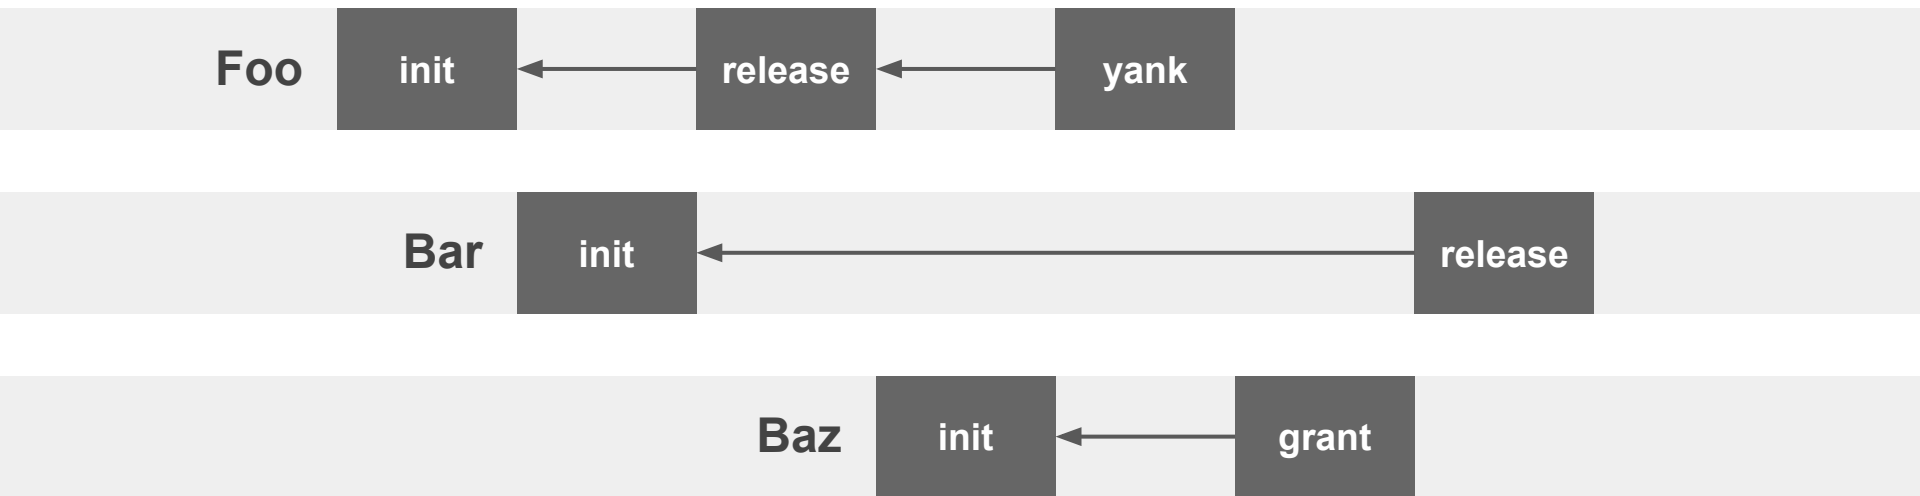
Releases

- 0.1.0
- 0.2.0

Releases

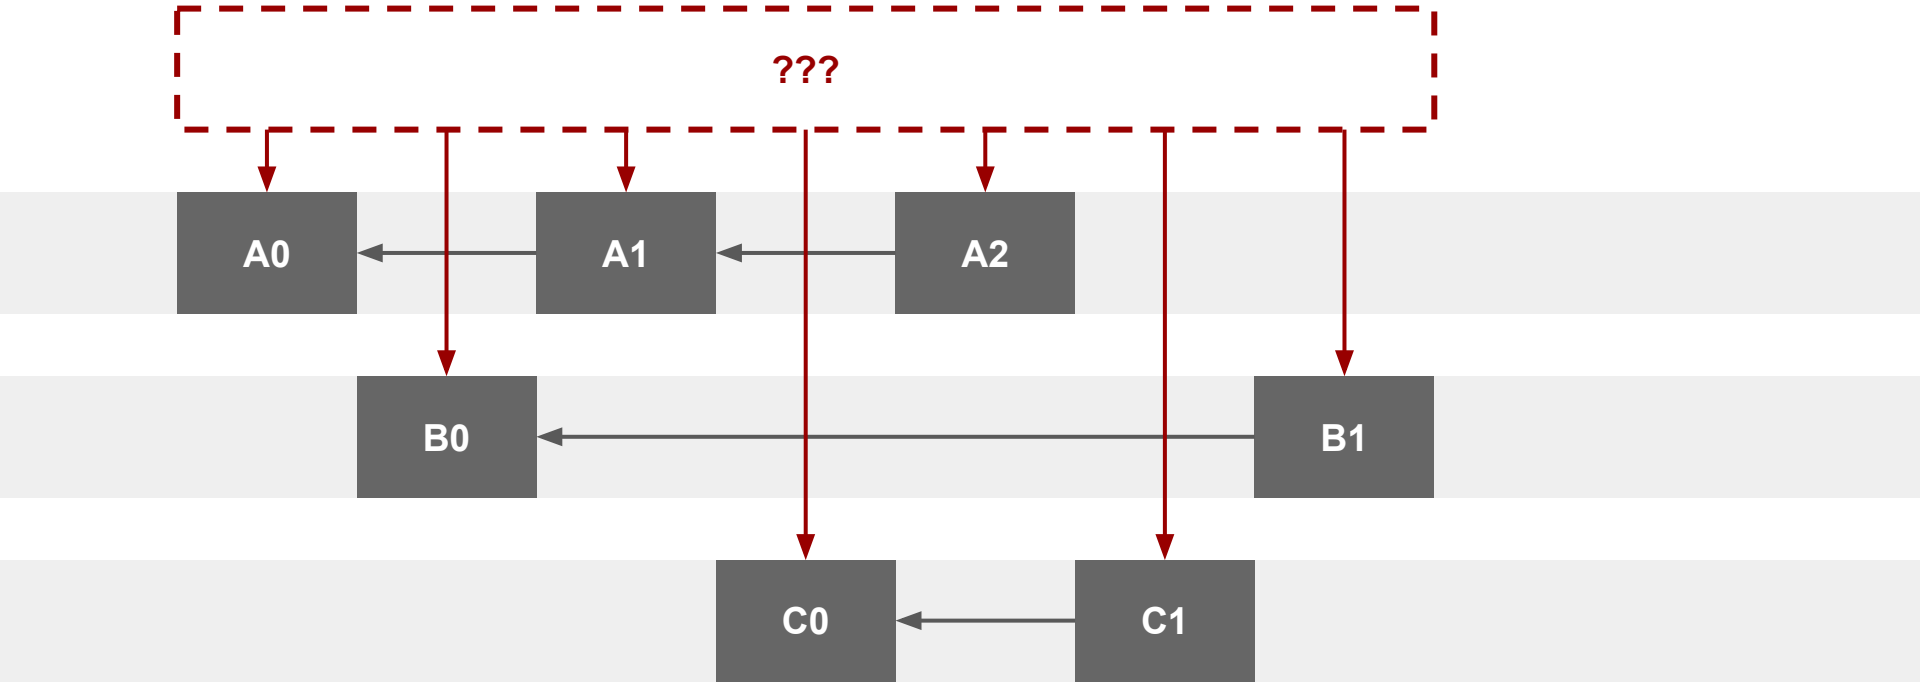
- 0.1.0
- ~~0.2.0~~

# Package Logs



# Cryptographically-Verifiable Commitments

What does the registry claim has happened?



# Verifiable Log

## (Merkle Tree)



# What does a Verifiable Log do?

- A **Verifiable Log** is a total ordering of records
- Each **Verifiable Log** is described by a unique hash
- You can cryptographically check if a record is in the log

# Verifiable Log (Merkle Tree)

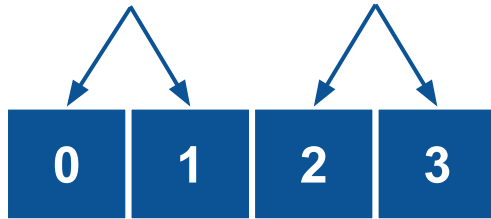


# Verifiable Log (Merkle Tree)



**Leaf Hash**  
`hash ( 0b00 || leaf )`

# Verifiable Log (Merkle Tree)



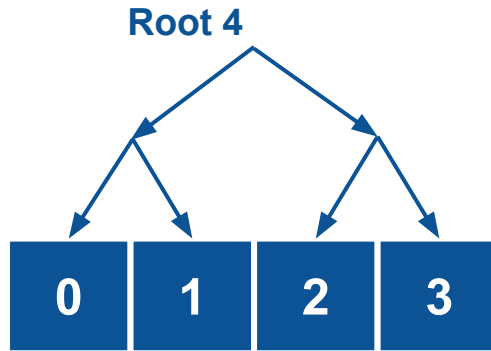
## Branch Hash

`hash ( 0b01 || left || right )`

## Leaf Hash

`hash ( 0b00 || leaf )`

# Verifiable Log (Merkle Tree)



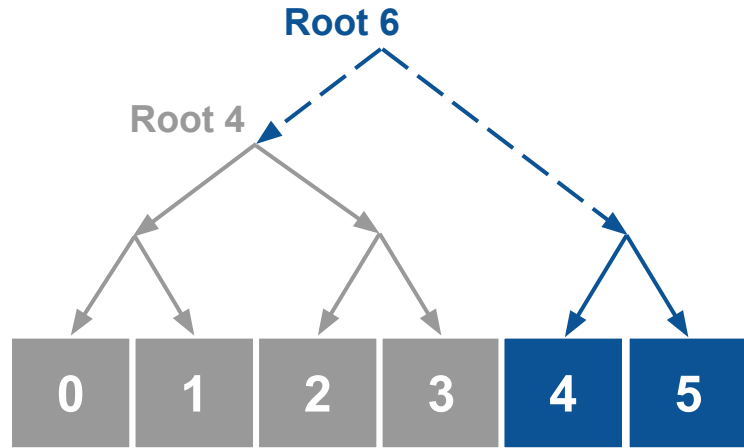
## Branch Hash

`hash ( 0b01 || left || right )`

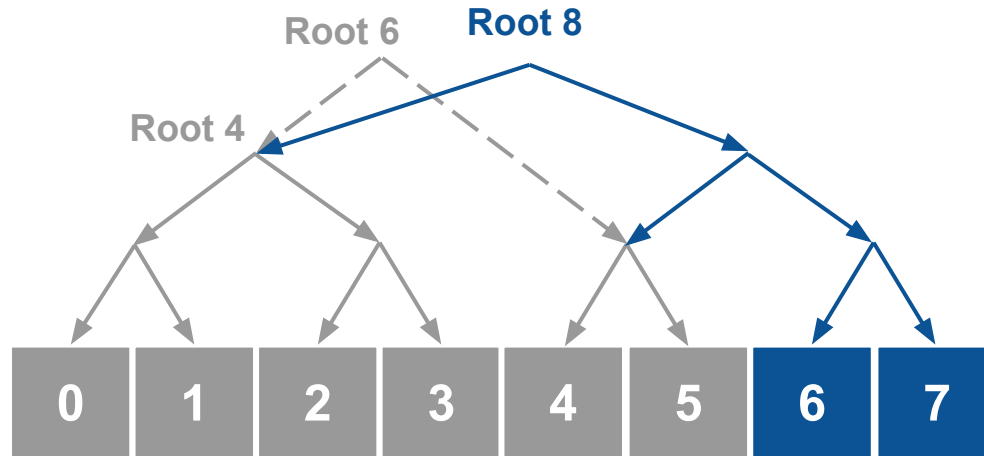
## Leaf Hash

`hash ( 0b00 || leaf )`

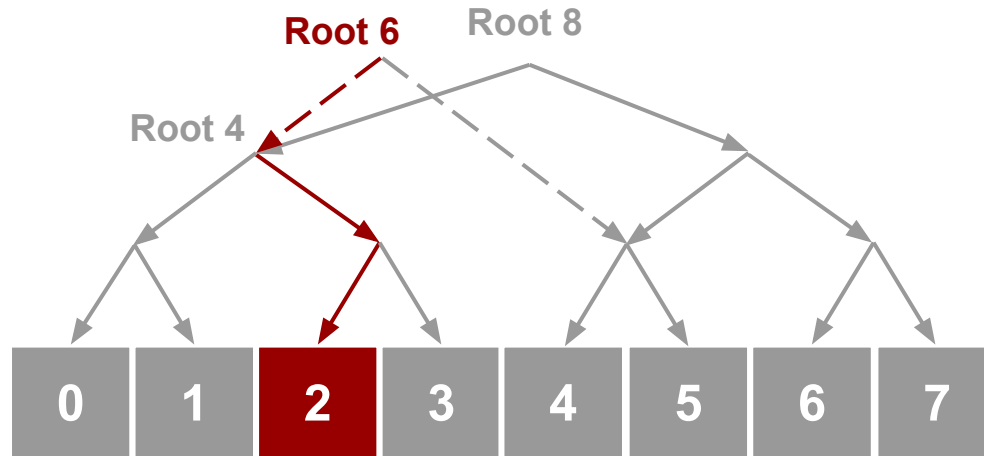
# Verifiable Log (Merkle Tree)



# Verifiable Log (Merkle Tree)

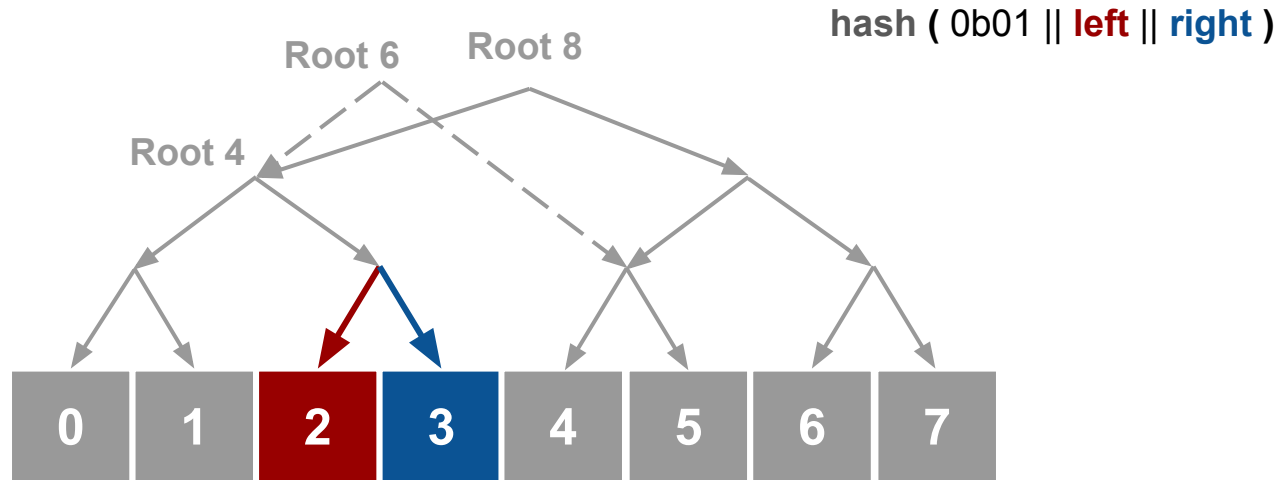


# Verifiable Log (Merkle Tree) - Inclusion Proofs

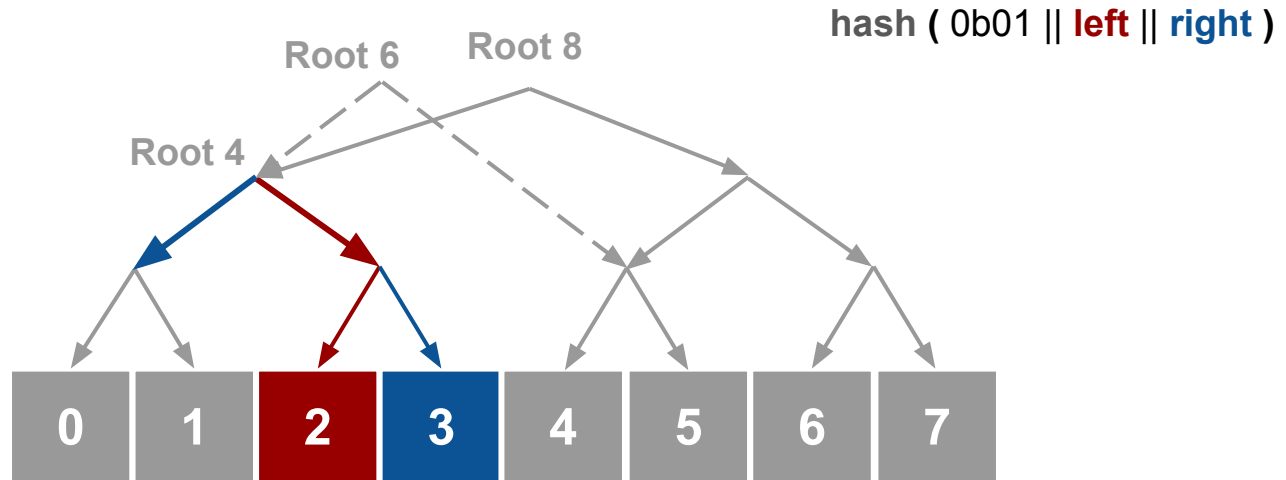




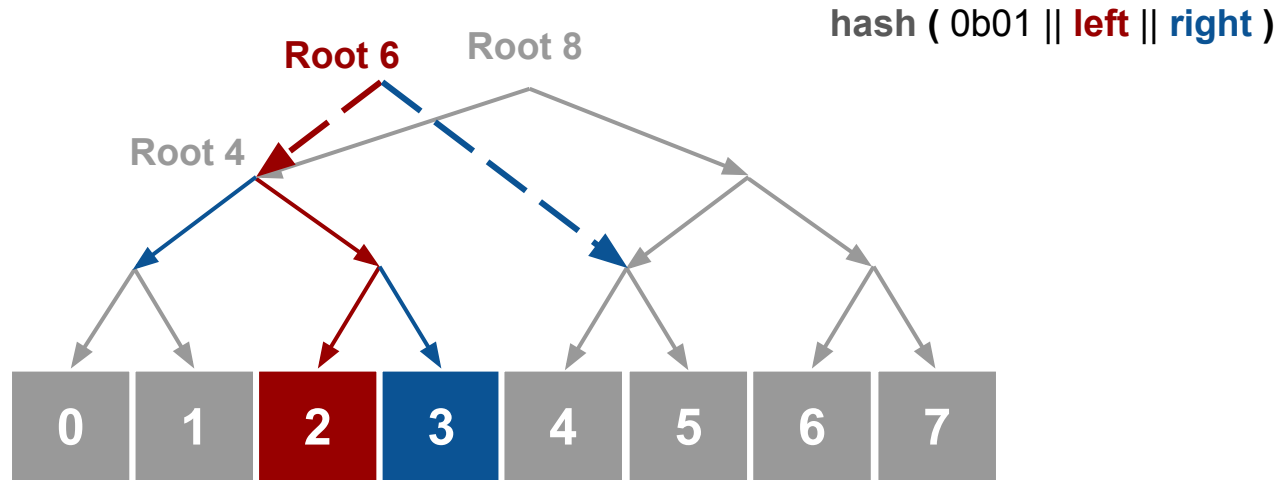
# Verifiable Log (Merkle Tree) - Inclusion Proofs



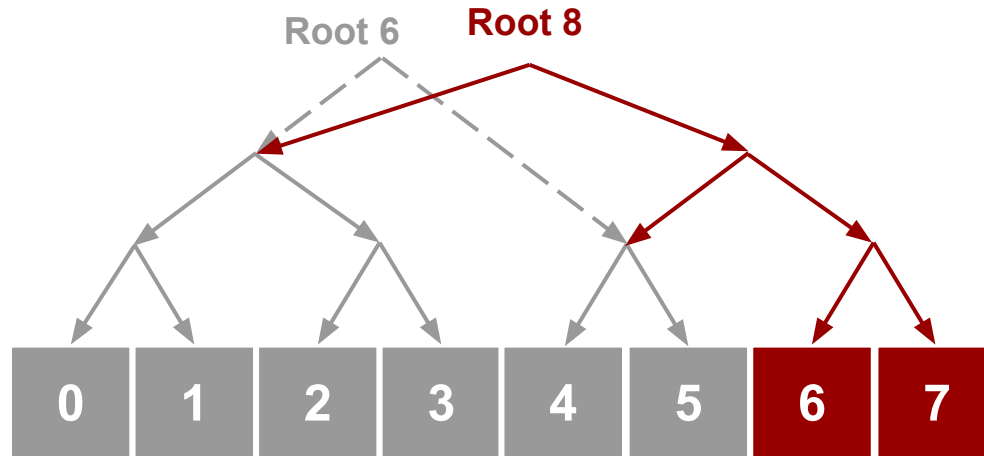
# Verifiable Log (Merkle Tree) - Inclusion Proofs



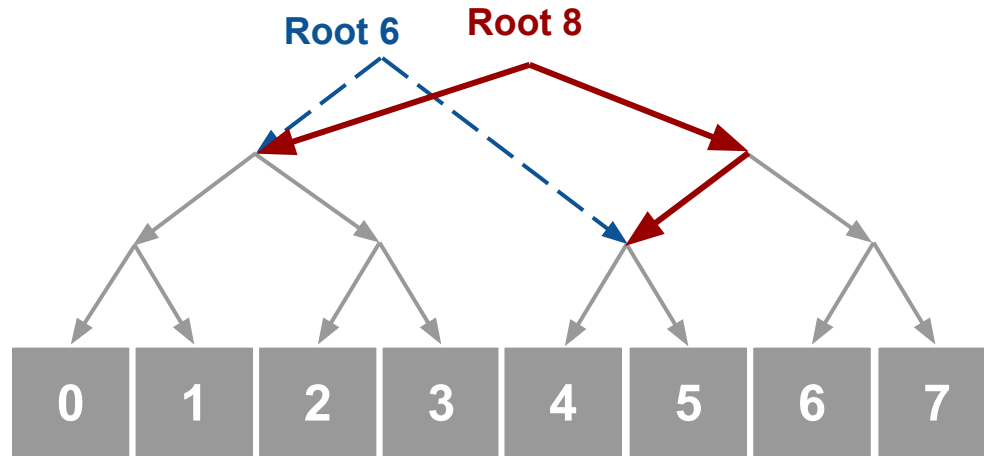
# Verifiable Log (Merkle Tree) - Inclusion Proofs



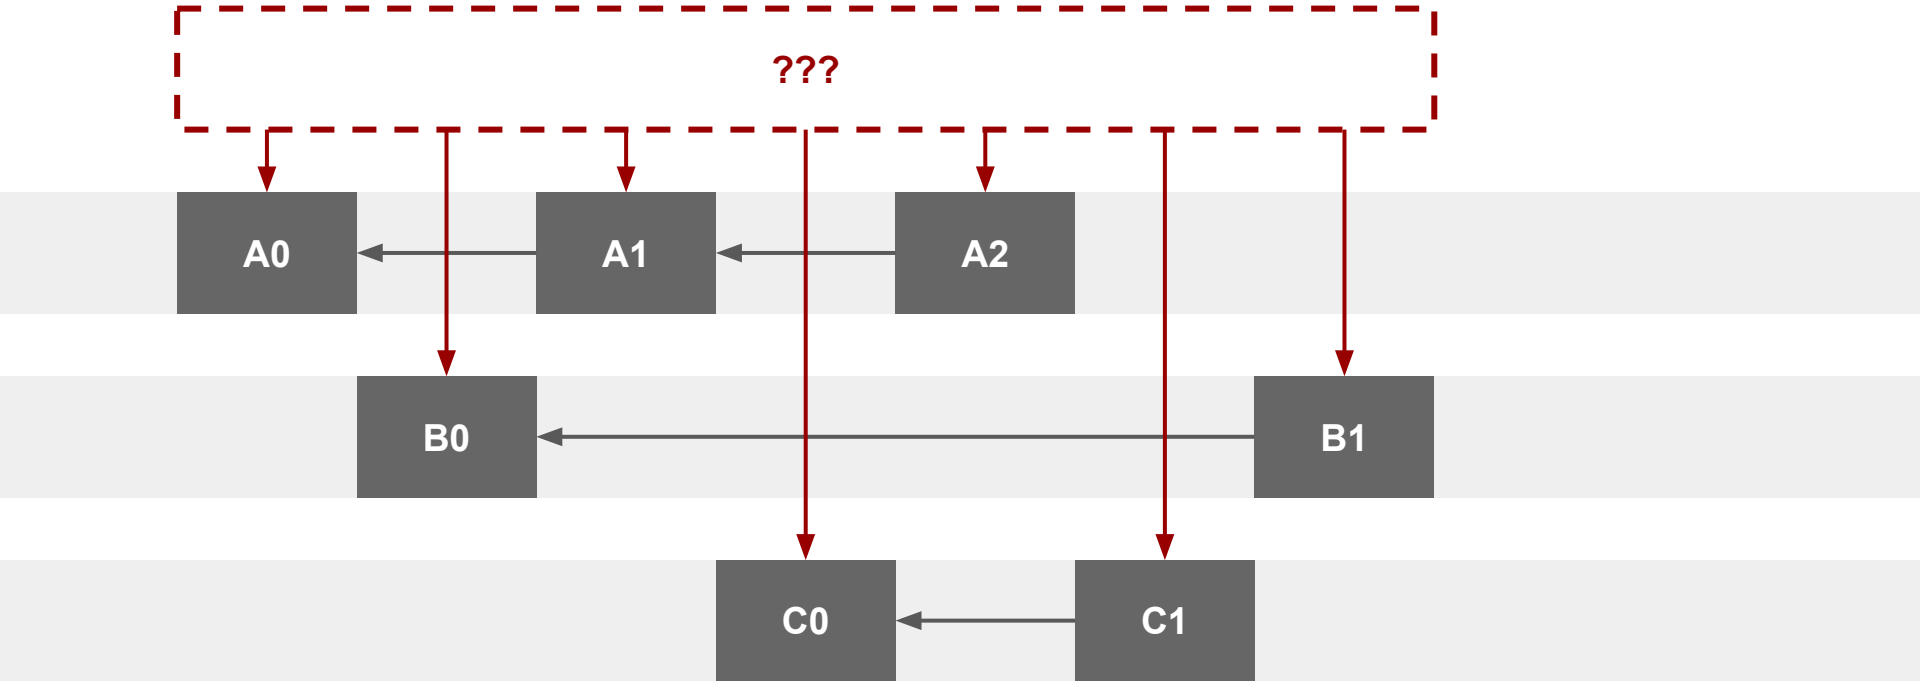
# Verifiable Log (Merkle Tree) - Consistency Proofs



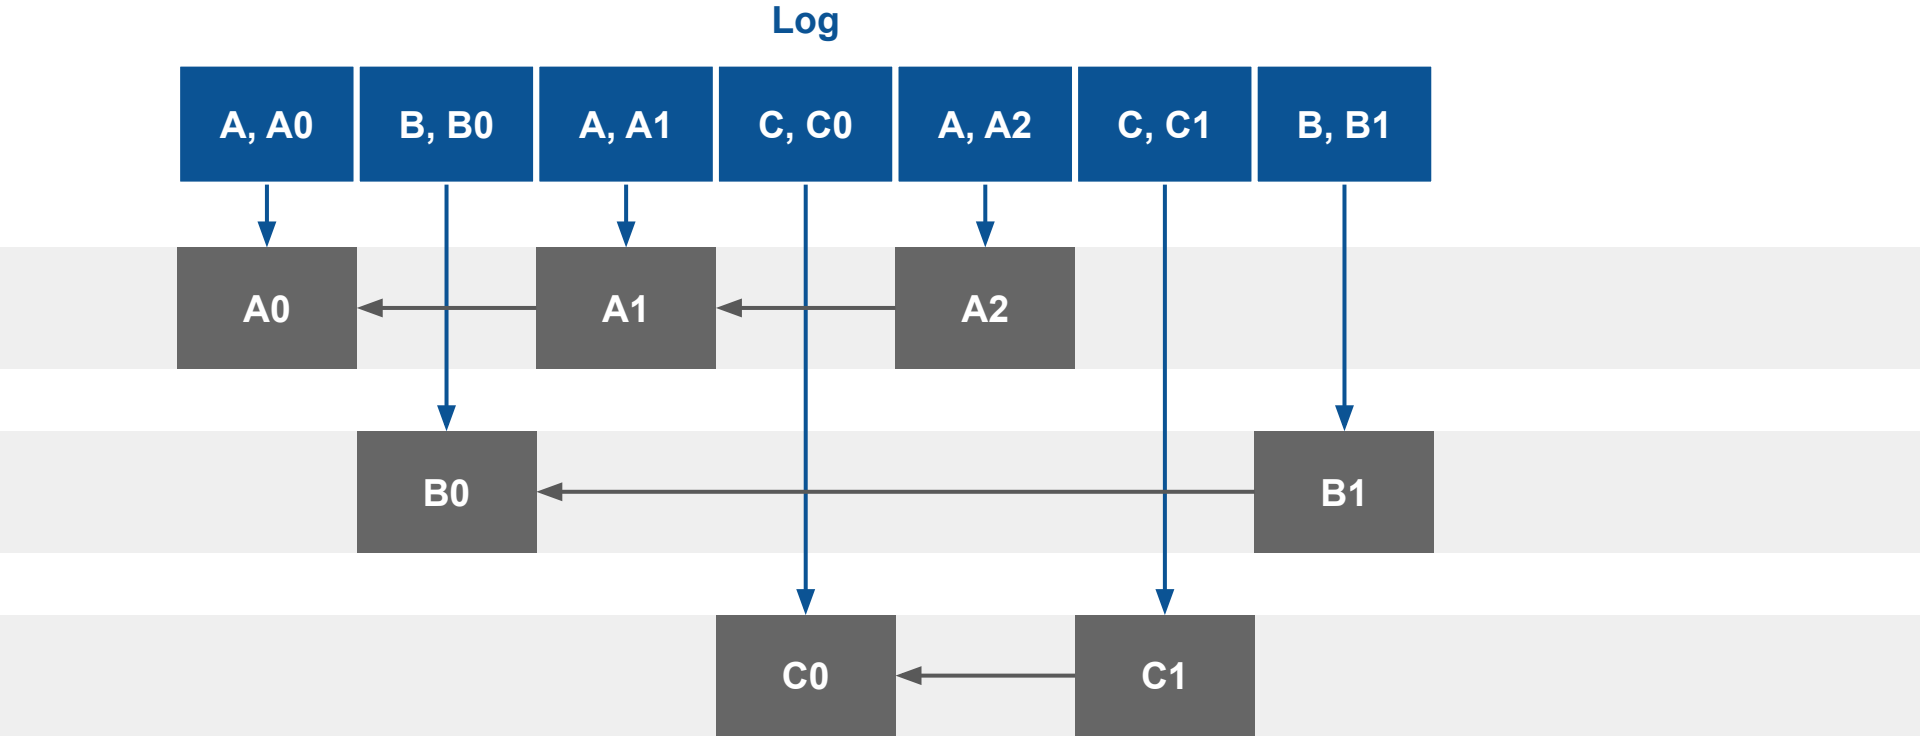
# Verifiable Log (Merkle Tree) - Consistency Proofs



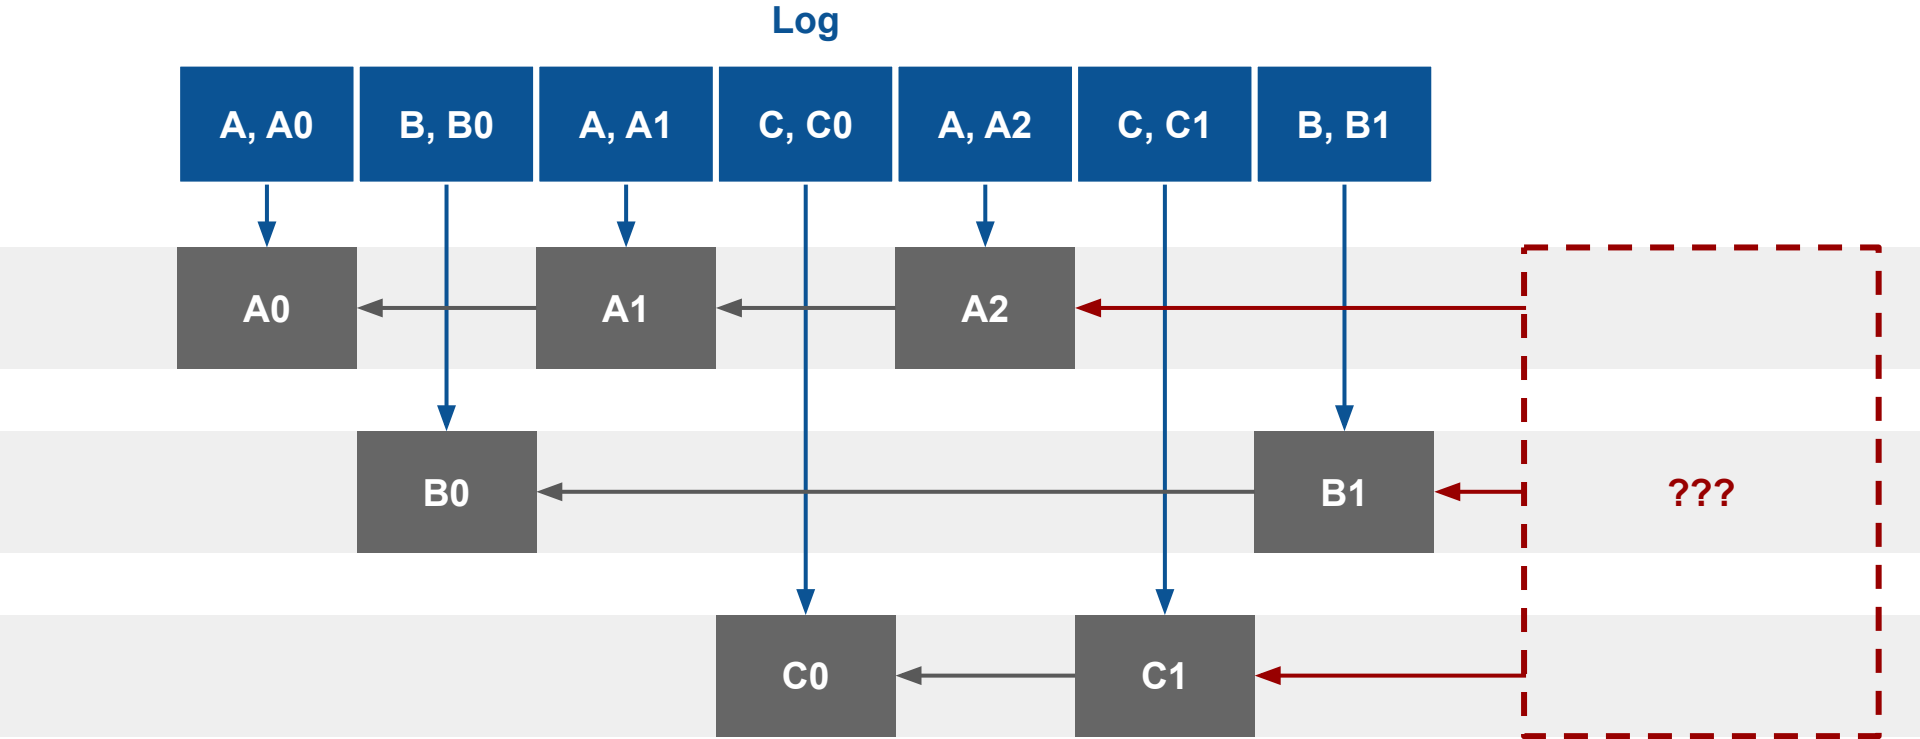
So, what does the registry claim has happened?



# The Package Records in the Verifiable Log



# How do clients know what the latest record is?





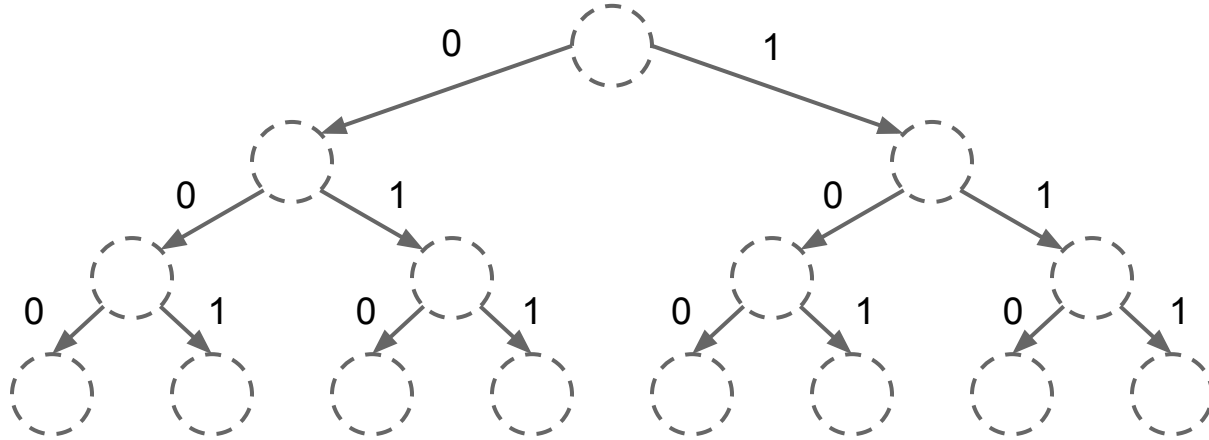
# Verifiable Map

## (Sparse Merkle Tree)

# What does a Verifiable Map do?

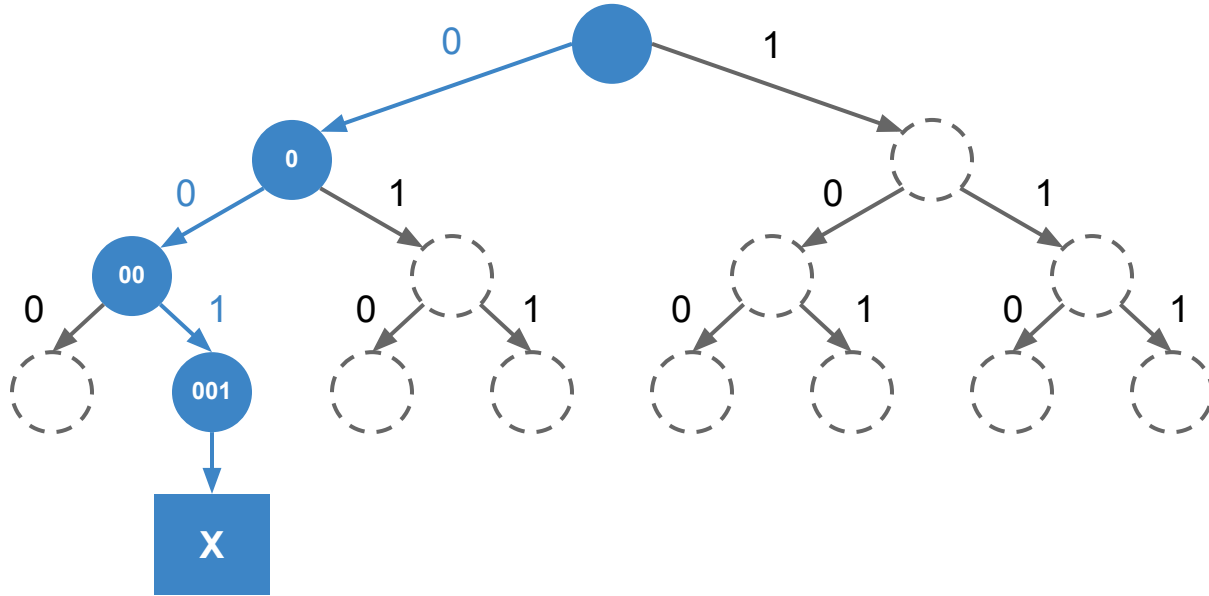
- A **Verifiable Map** is a key-value mapping
- Each **Verifiable Map** is described by a unique hash
- You can cryptographically check if a value is associated with a key

# Verifiable Map (Sparse Merkle Tree)



# Verifiable Map (Sparse Merkle Tree)

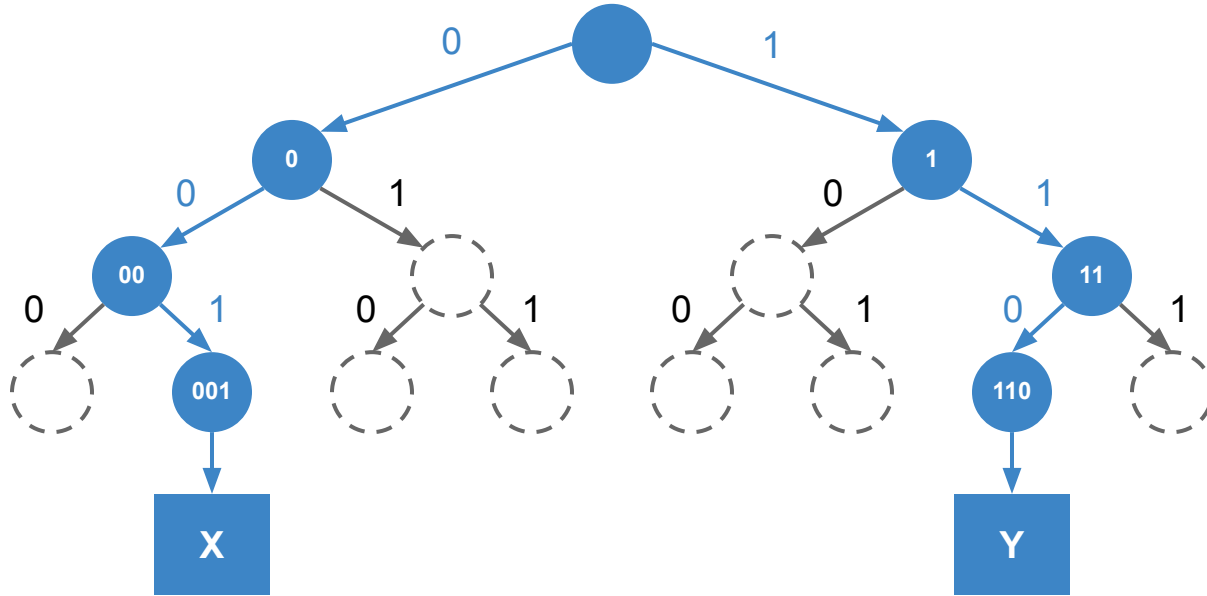
001 → X



# Verifiable Map (Sparse Merkle Tree)

001 → X

110 → Y




# Verifiable Map (Sparse Merkle Tree)

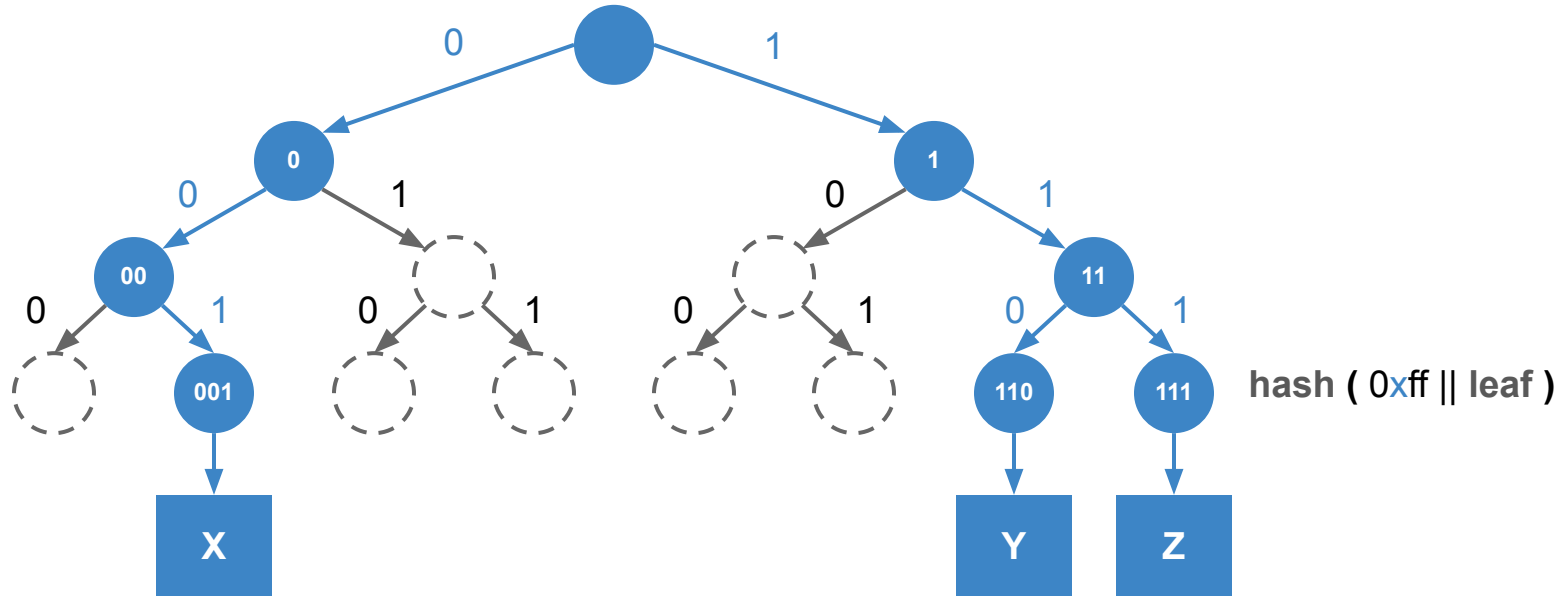
001 → X

110 → Y

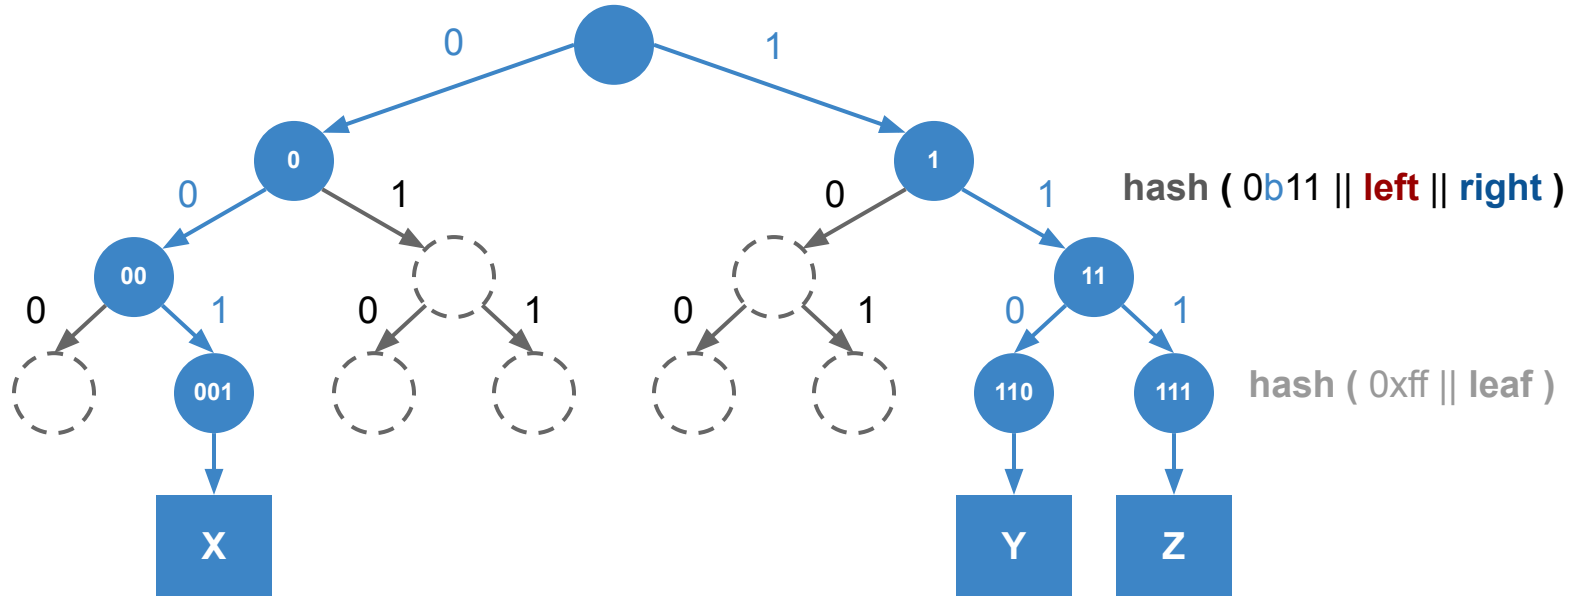
111 → Z



# Verifiable Map (Sparse Merkle Tree)

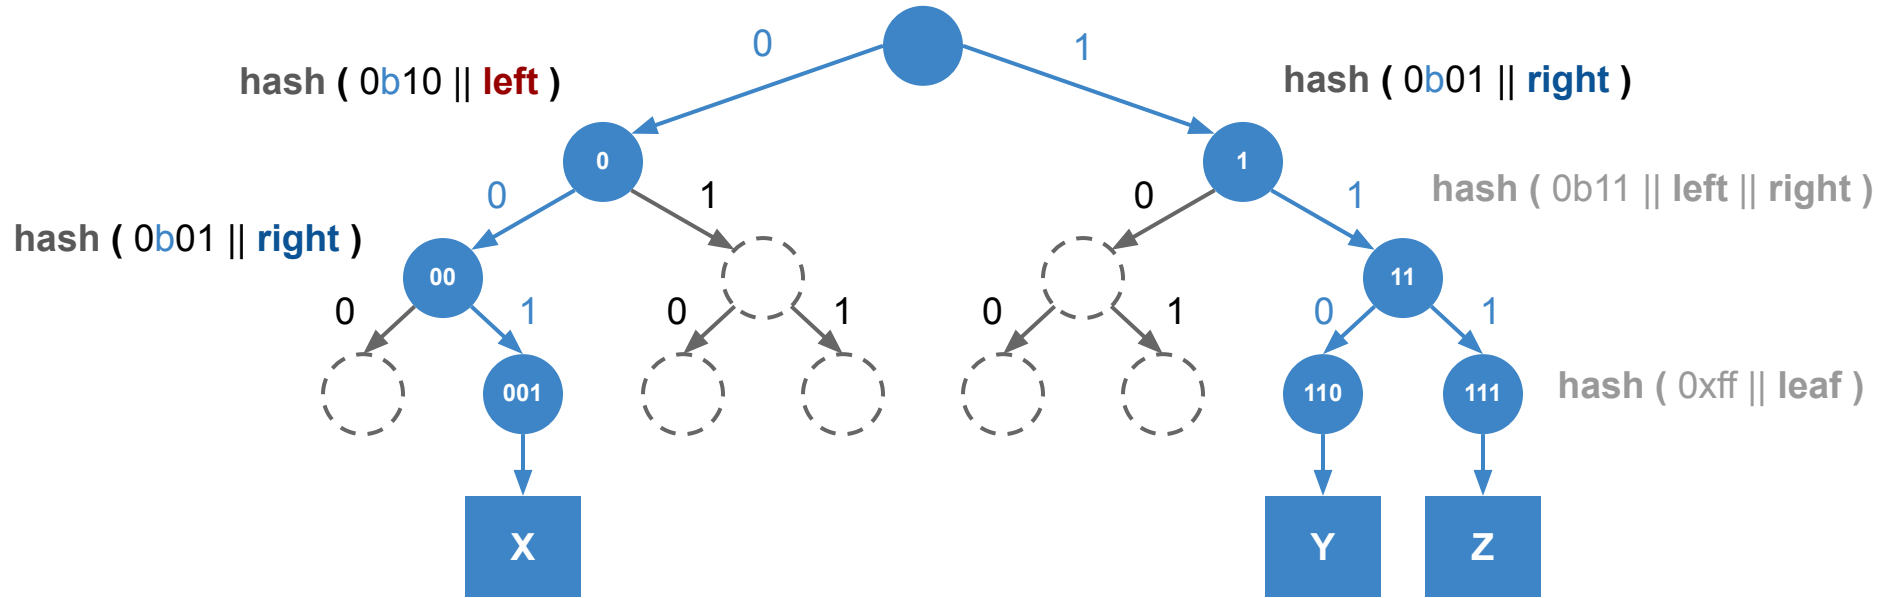


# Verifiable Map (Sparse Merkle Tree)

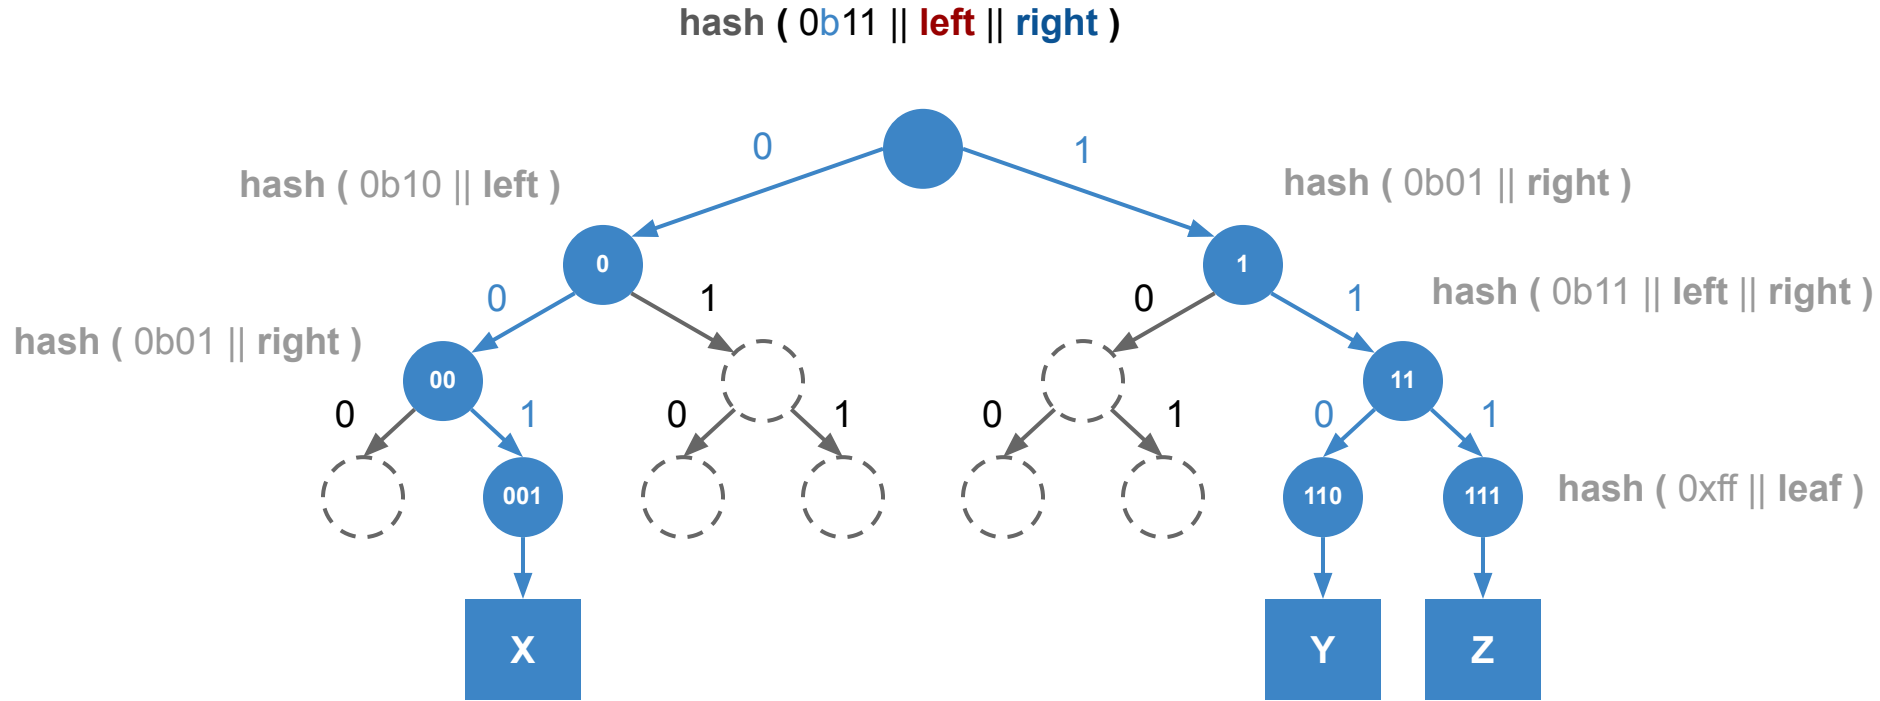




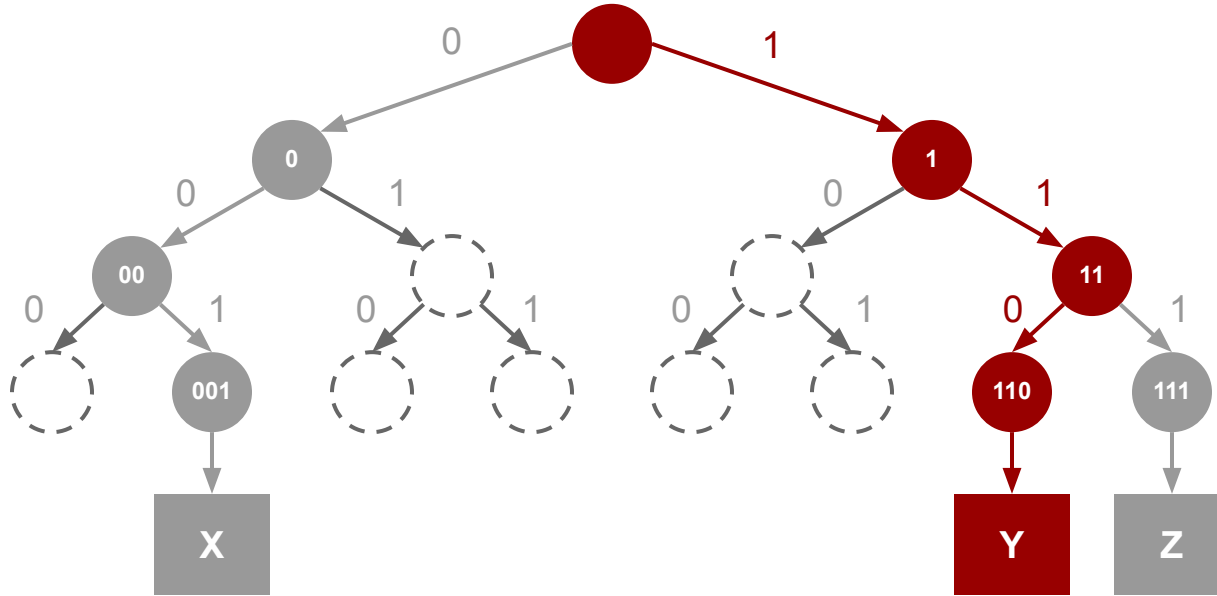
# Verifiable Map (Sparse Merkle Tree)



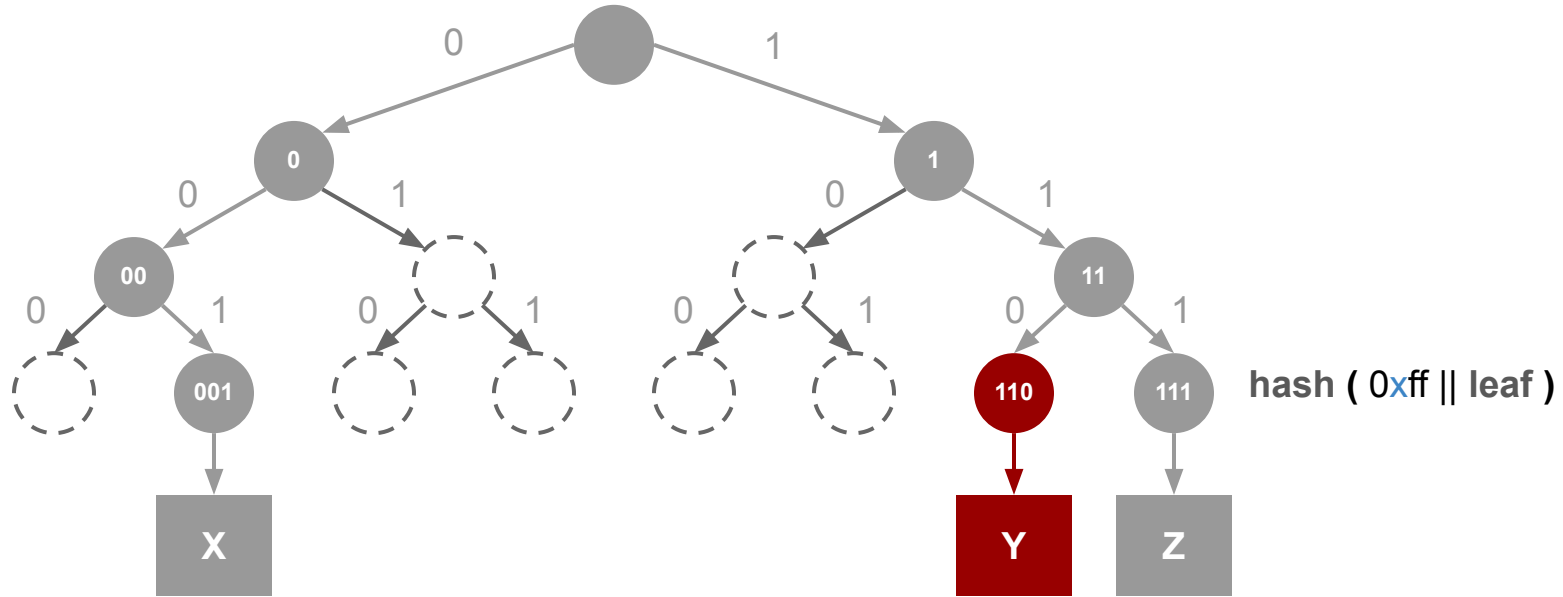
# Verifiable Map (Sparse Merkle Tree)



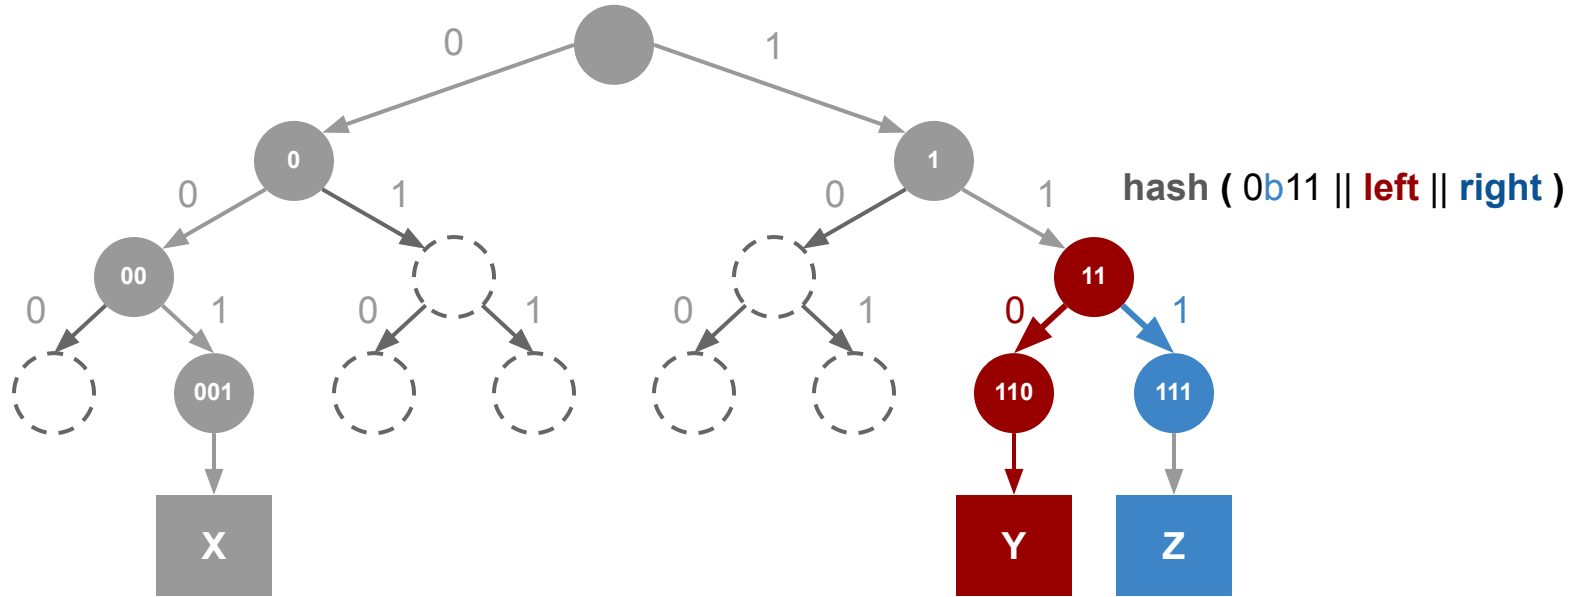
# Verifiable Map (Sparse Merkle Tree) - Inclusion Proof



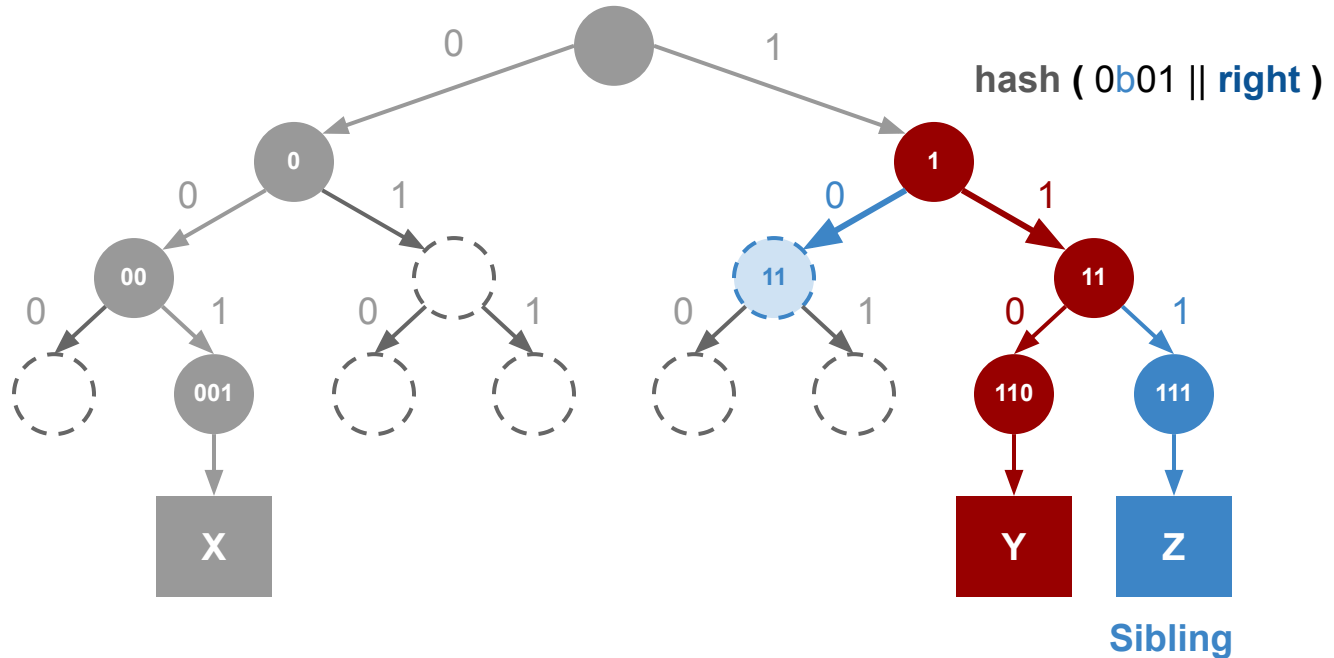
# Verifiable Map (Sparse Merkle Tree) - Inclusion Proof



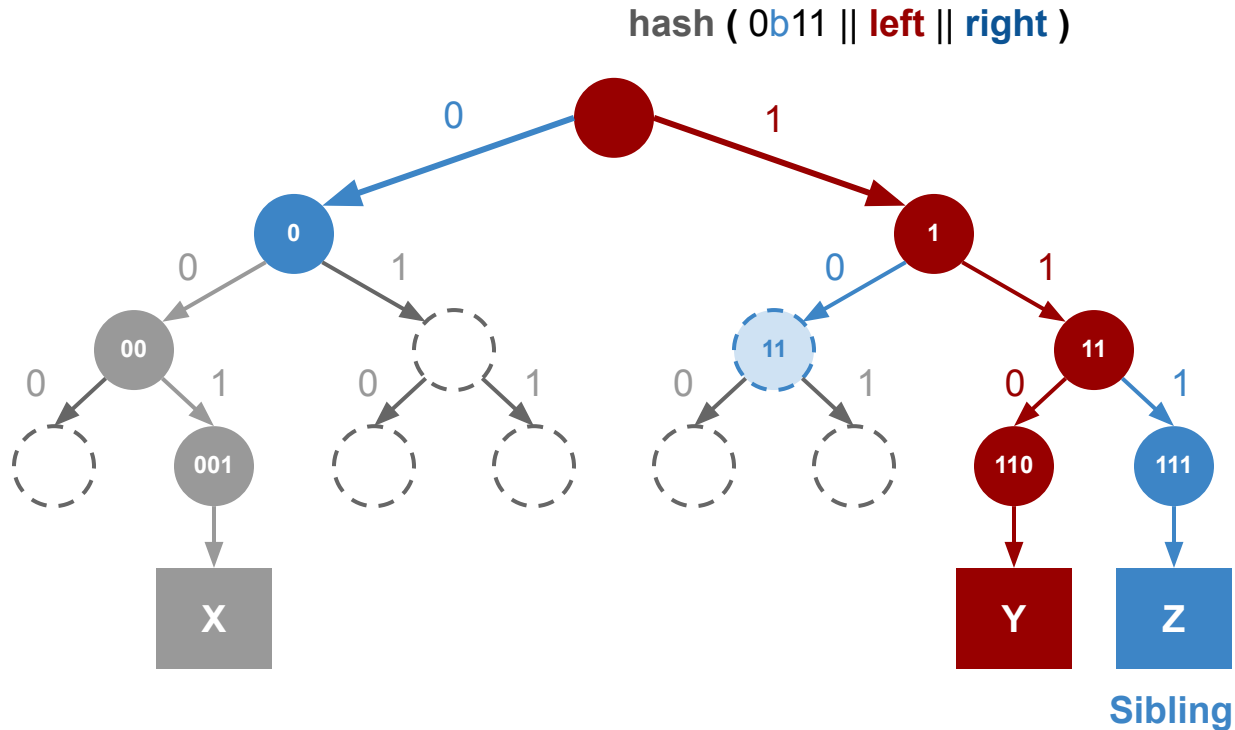
# Verifiable Map (Sparse Merkle Tree) - Inclusion Proof



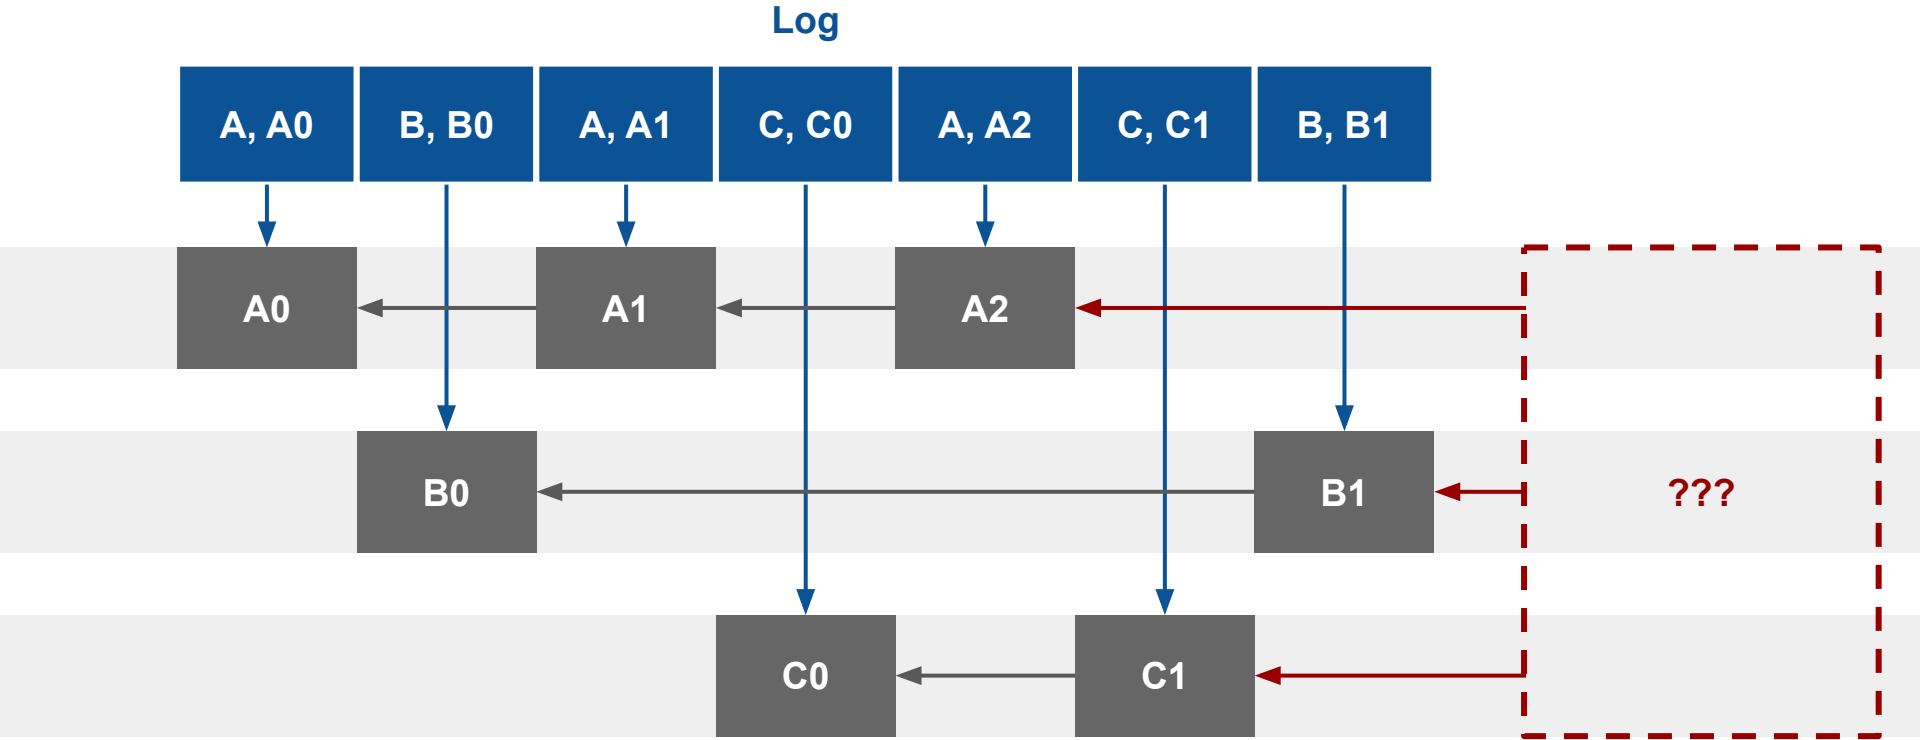
# Verifiable Map (Sparse Merkle Tree) - Inclusion Proof



# Verifiable Map (Sparse Merkle Tree) - Inclusion Proof

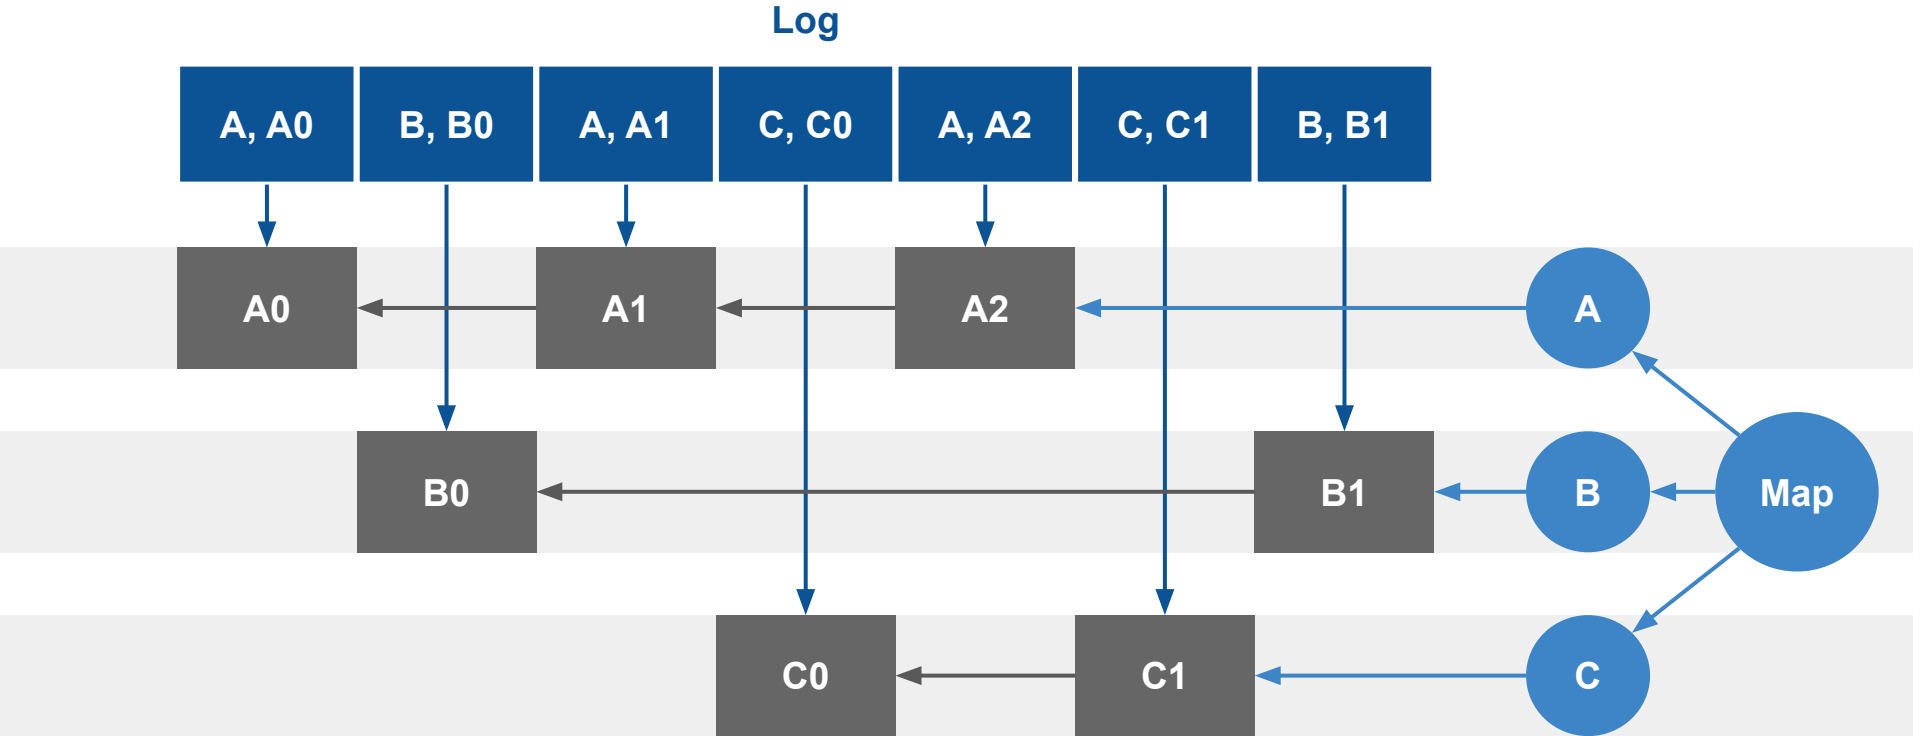


So, how do clients know what the latest record is?

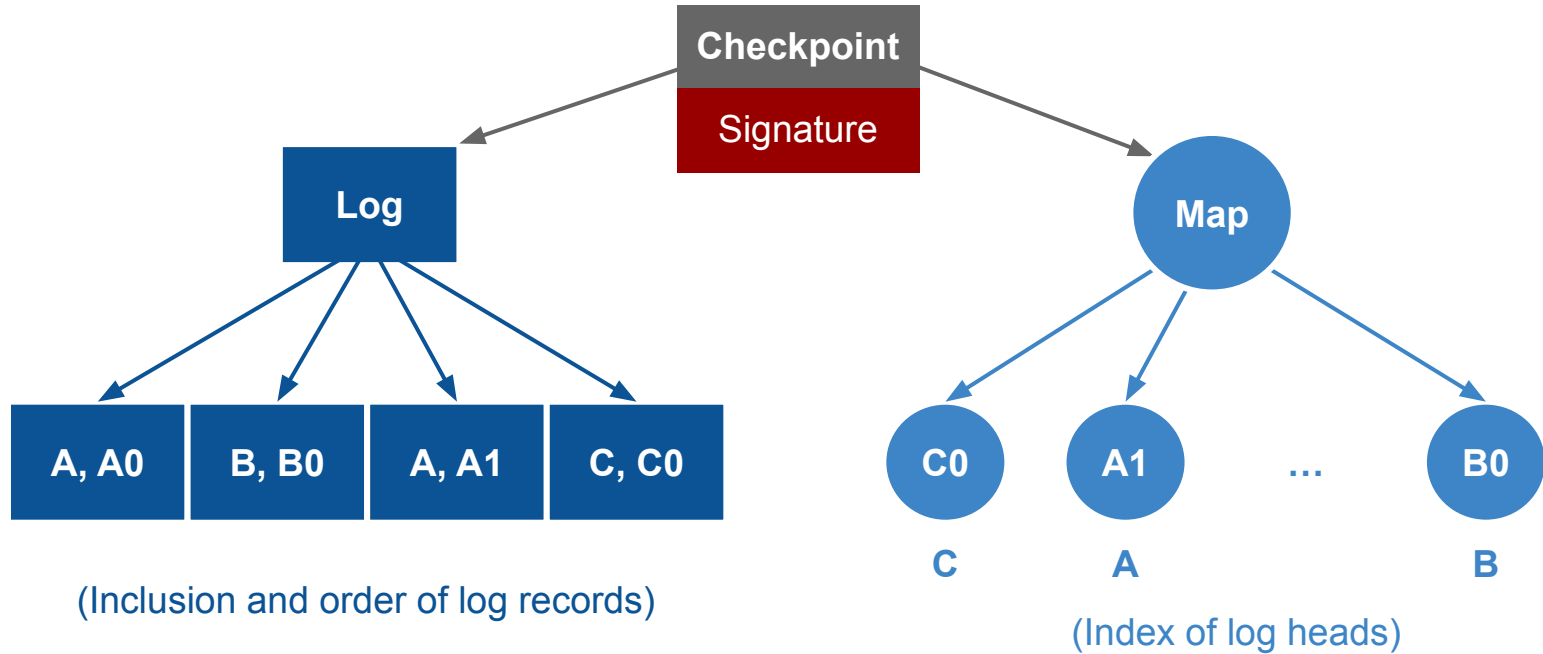




# Use a verifiable map of package log heads



# Overall Checkpoint / Commitment



# Auditing & Verification

# Clients

They are

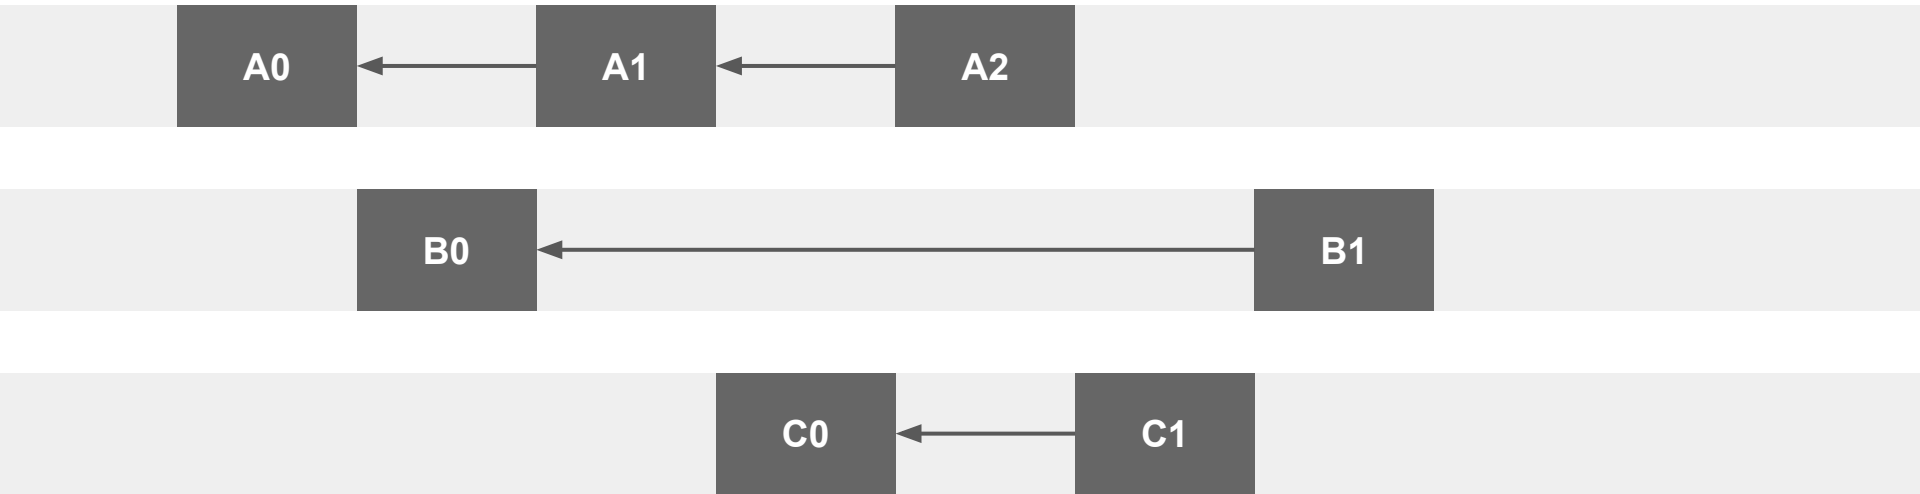
- resource constrained
- only interested in some of the state

So, clients verify

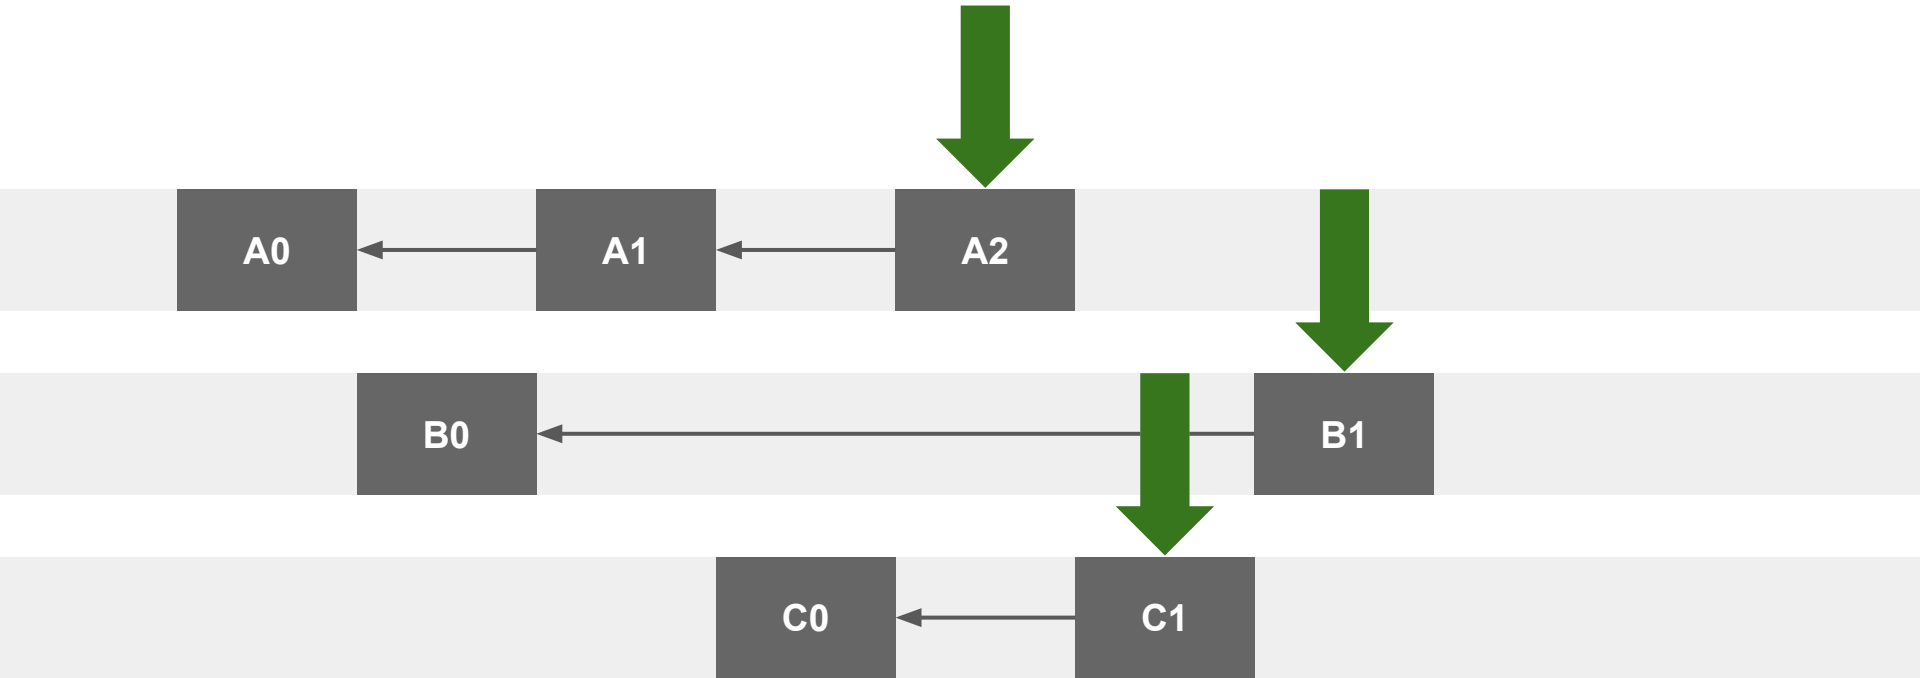
- relevant package state
- that the registry committed to that state
- the commitment is valid / correct

Clients know nothing at the start

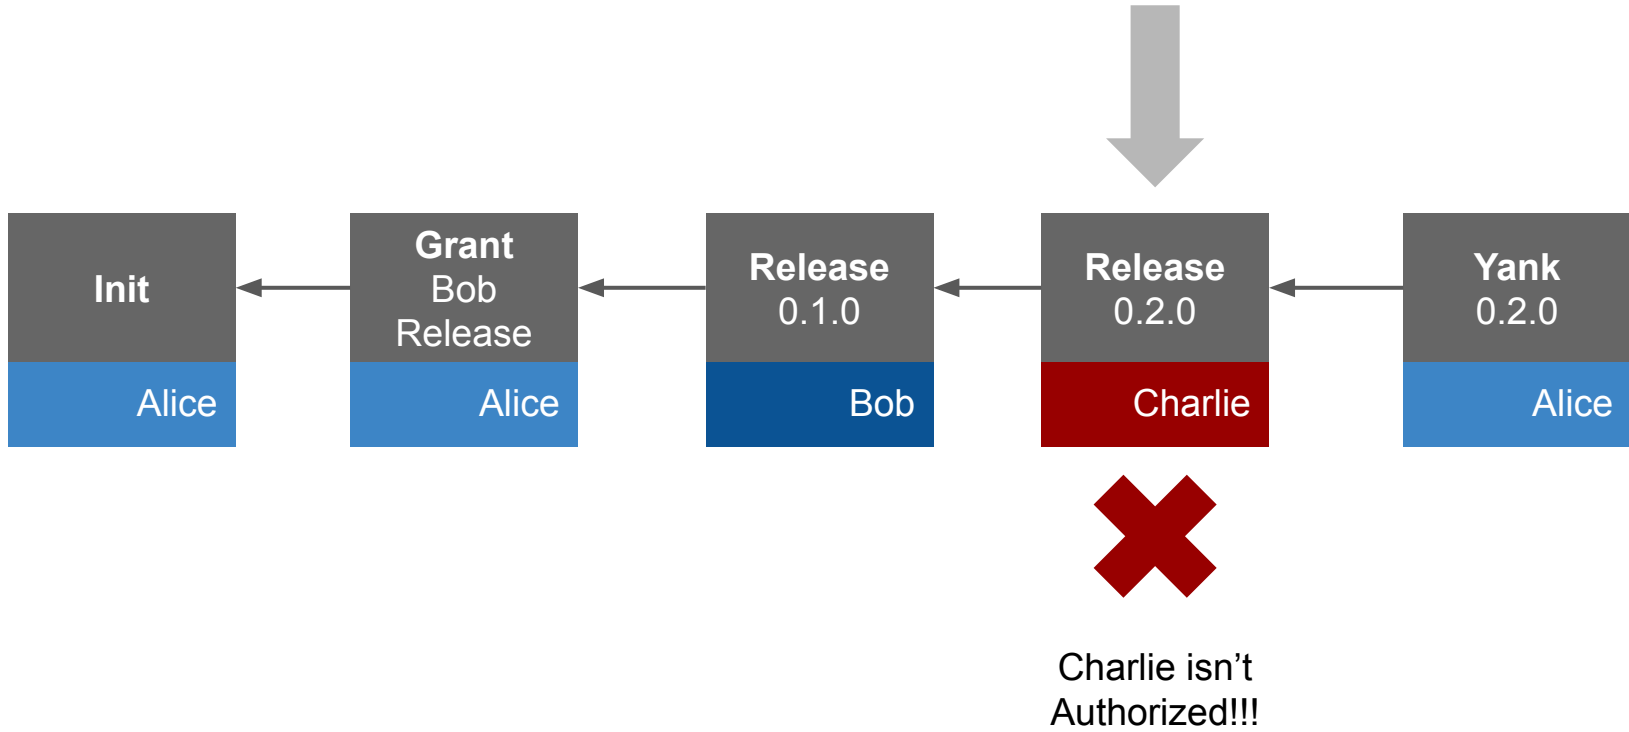
# Clients **download** the Package Logs



# Clients **verify** Package Logs



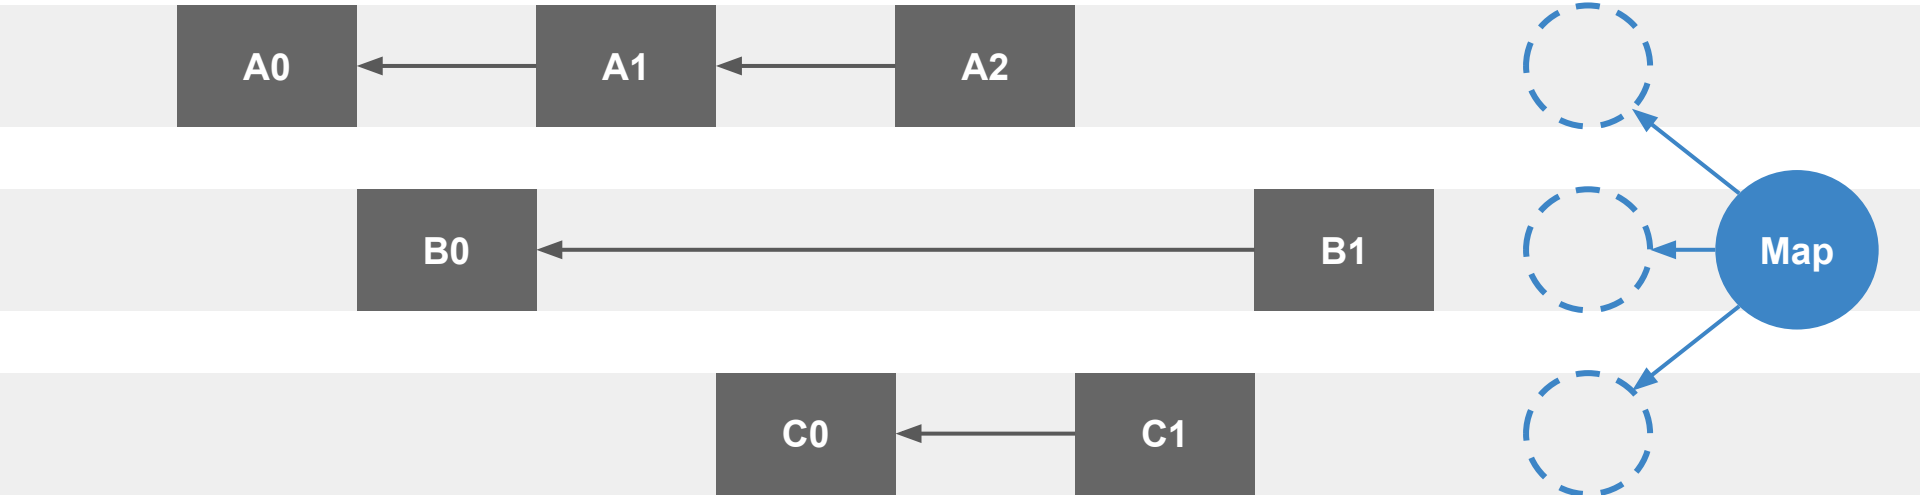
# Clients **verify** Package Logs



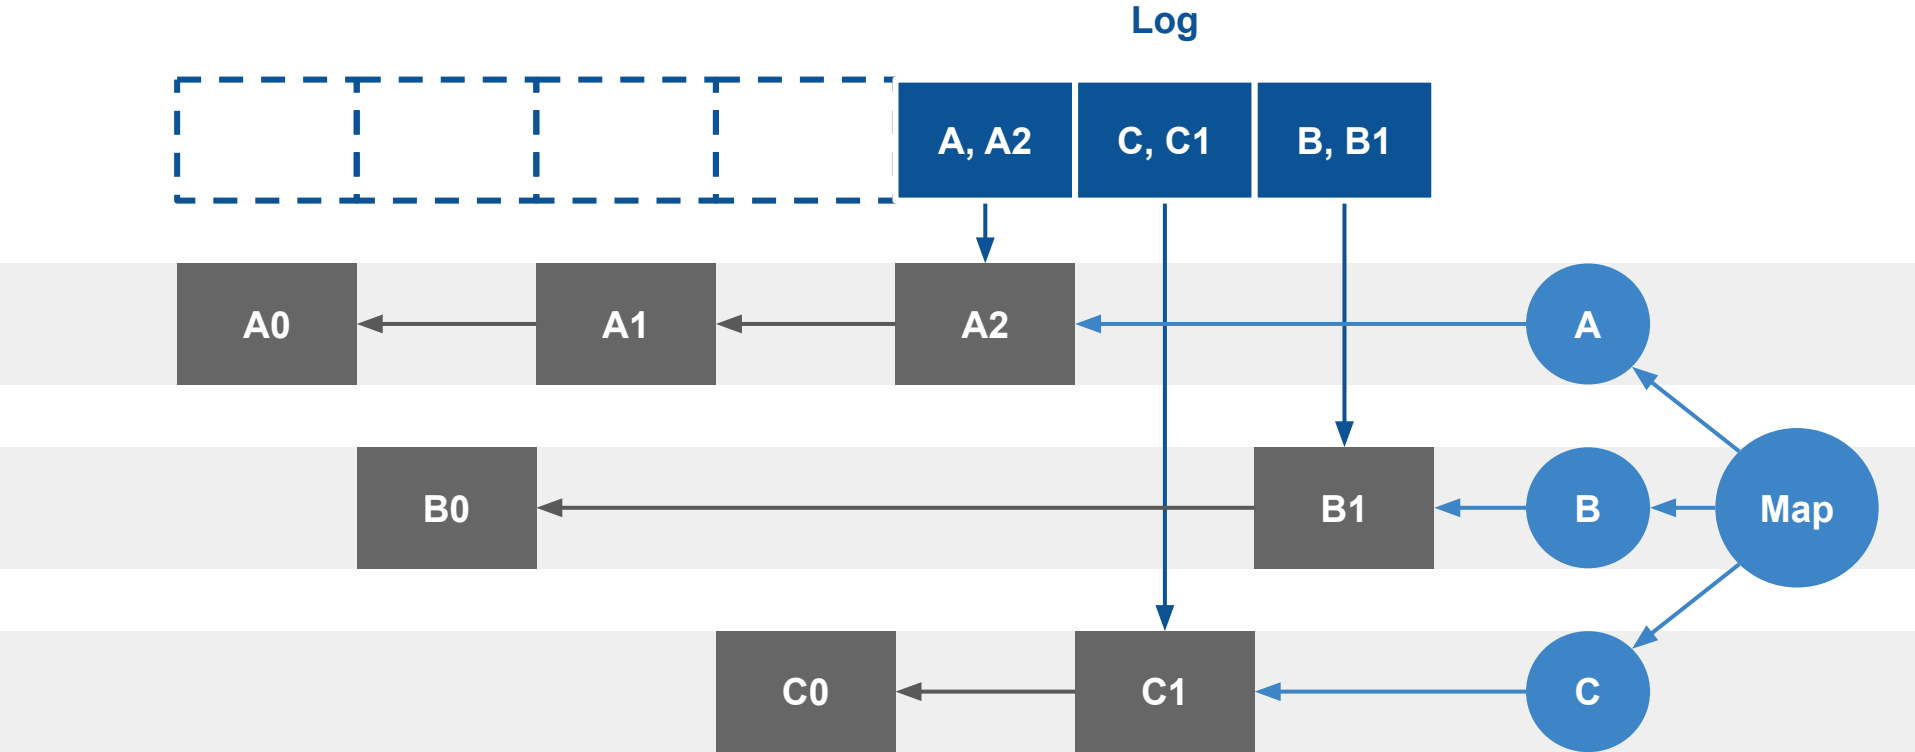


Clients **get** checkpoints for the map and log

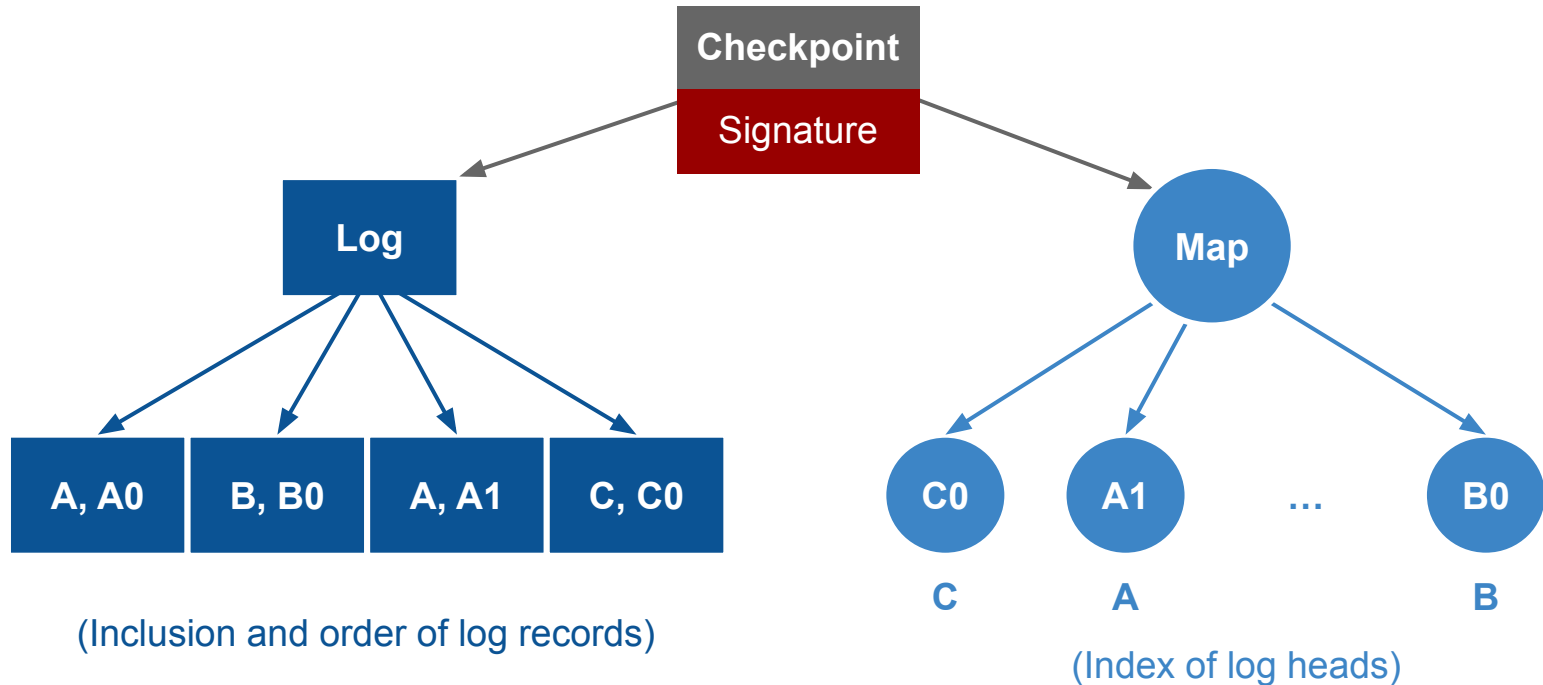
Log



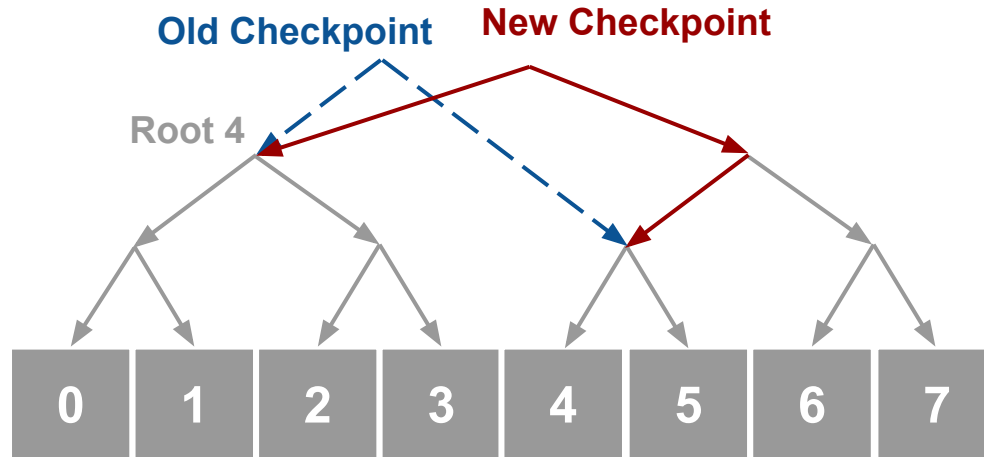
Clients **get** & **verify** proofs for map and log inclusion



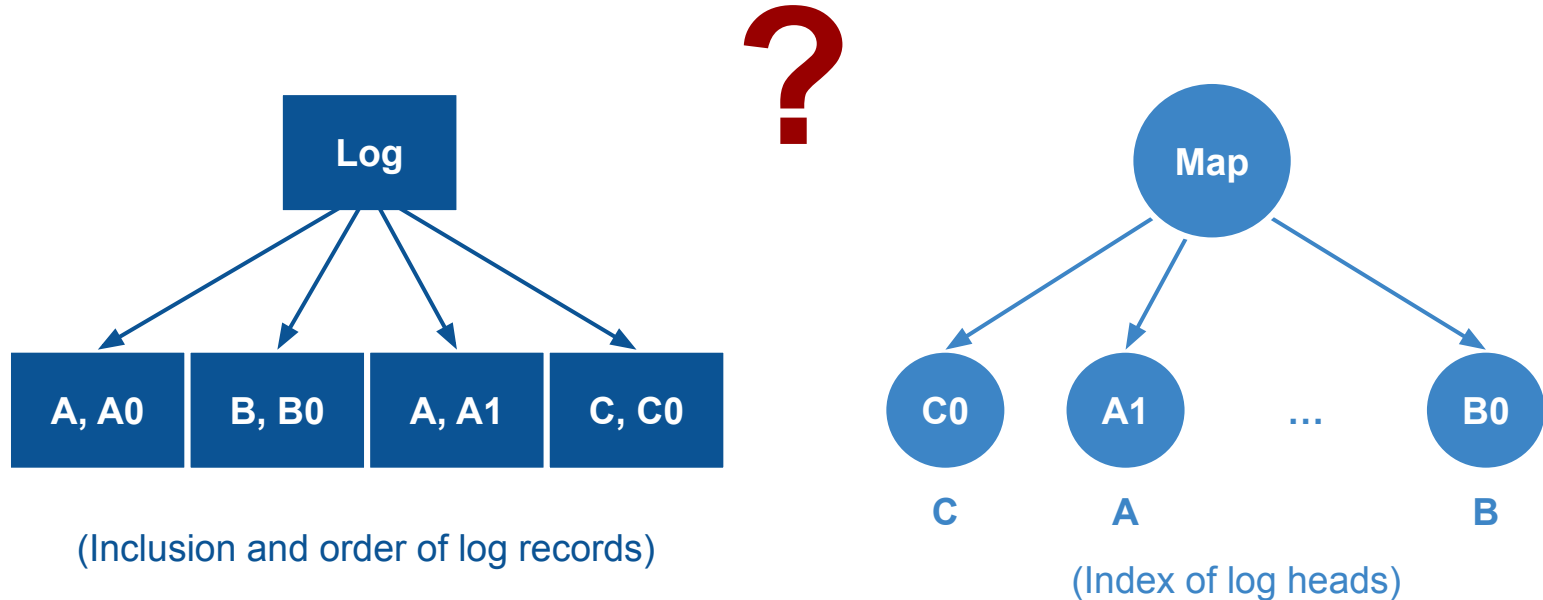
# Client **validate** Checkpoint Signature



# Client Validation - Log Checkpoint Consistency



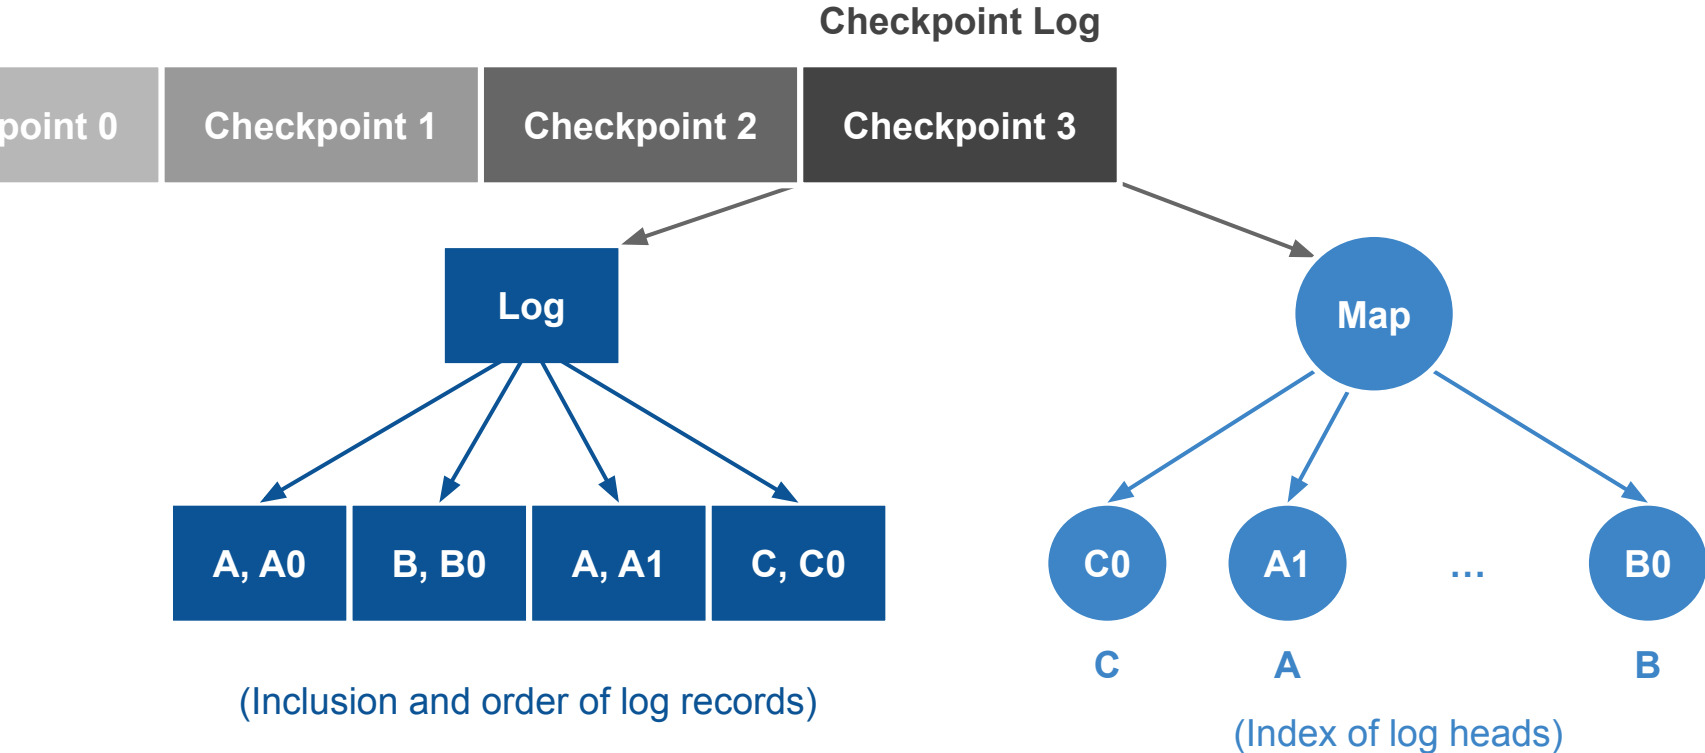
# Client can't validate consistency of the **Log** & **Map**



Who are clients  
gonna call?

# Monitors

# Monitors Audit Log/Map Consistency Over Time





# Components of Package Transparency (Revisited)

## 1. Publicly Available Package Registry State

- A collection of package logs

## 2. Cryptographically-Verifiable Commitments

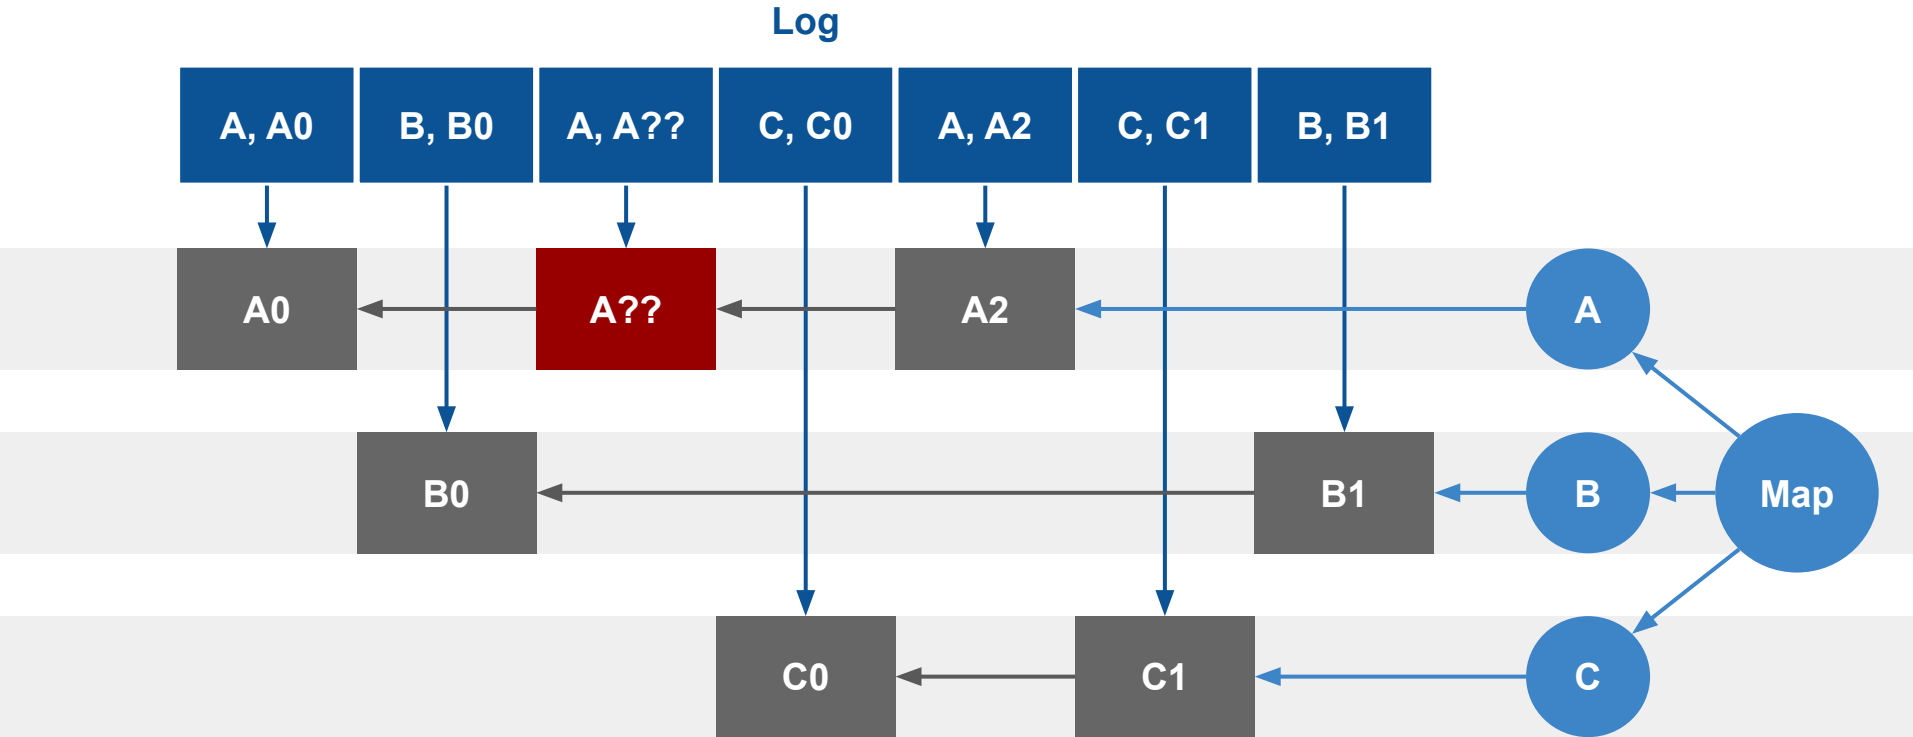
- Signed log + map checkpoints

## 3. Auditing Package Authors and the Registry

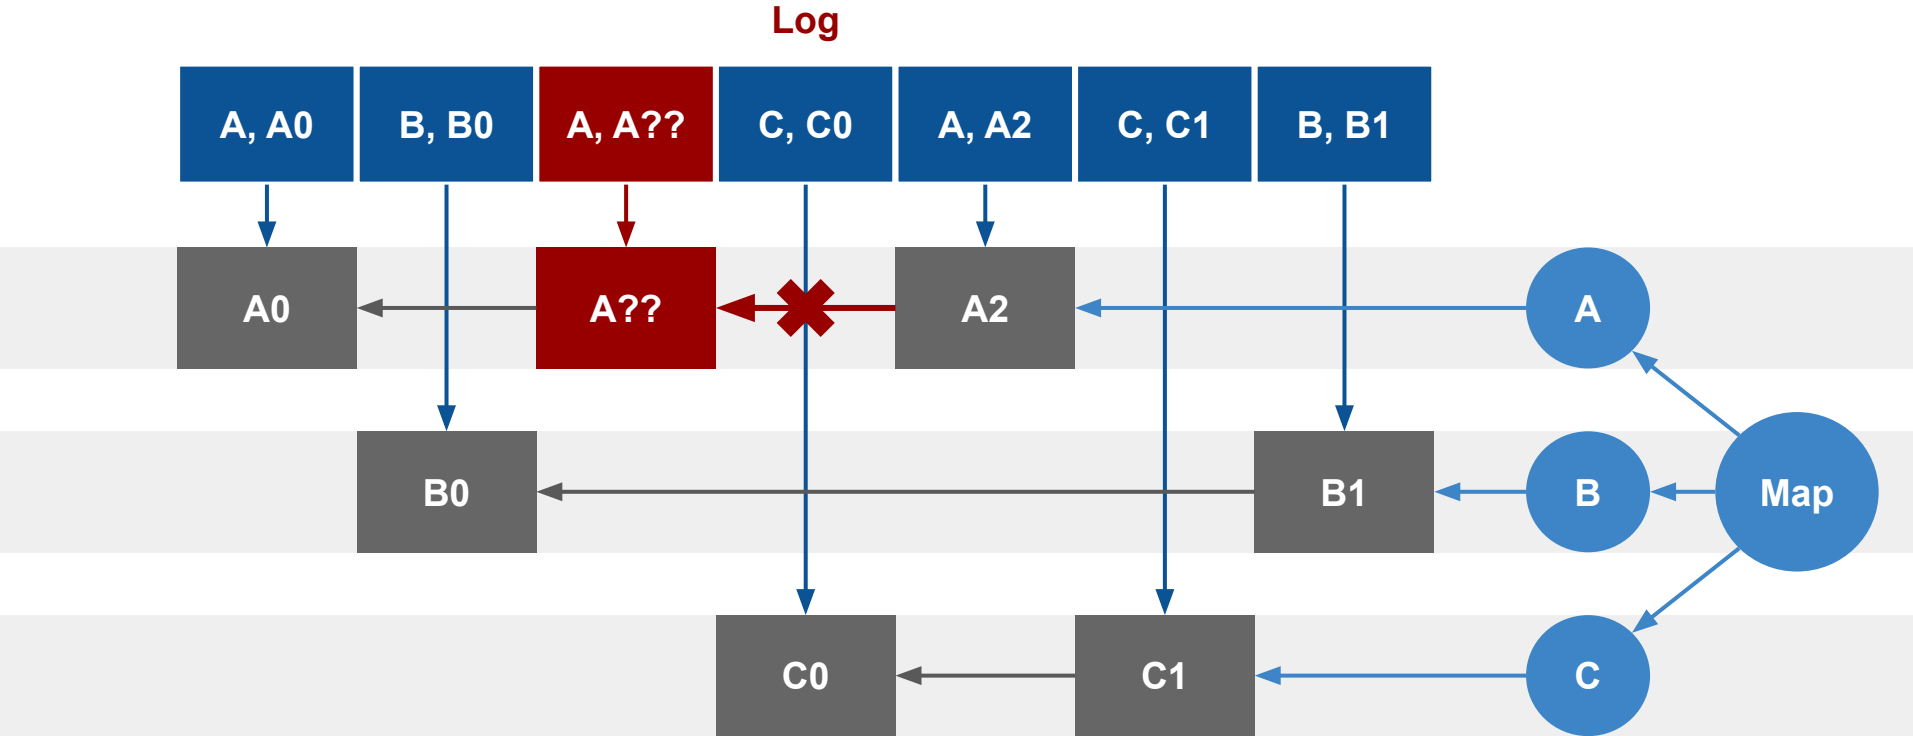
- Client and monitor validation

# Package Transparency vs. Various Attacks

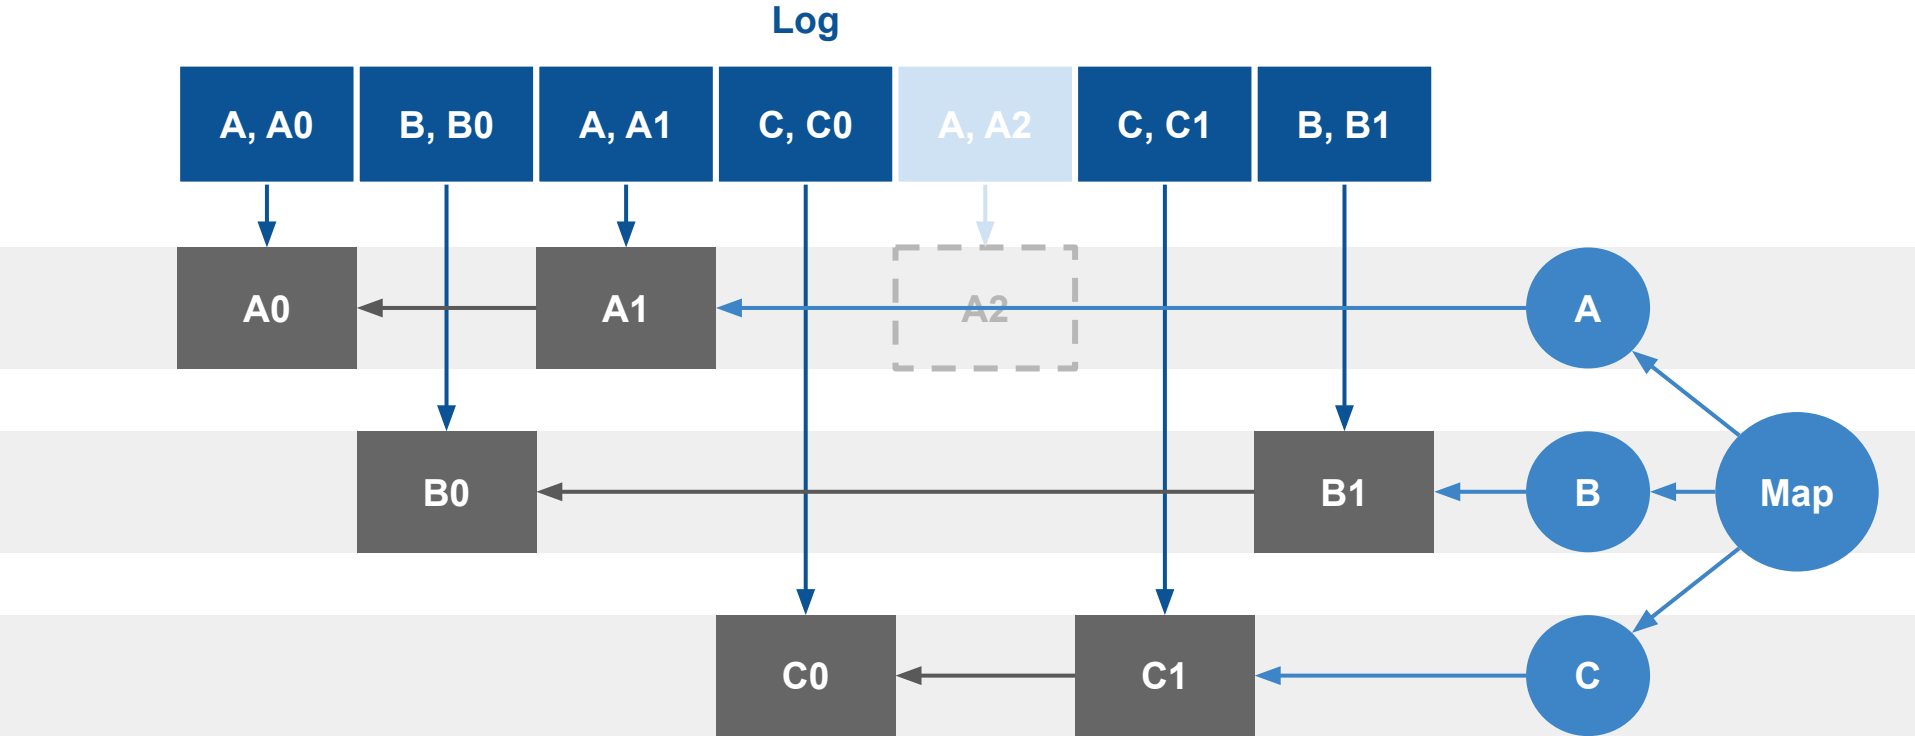
# What if an attacker changes a record?



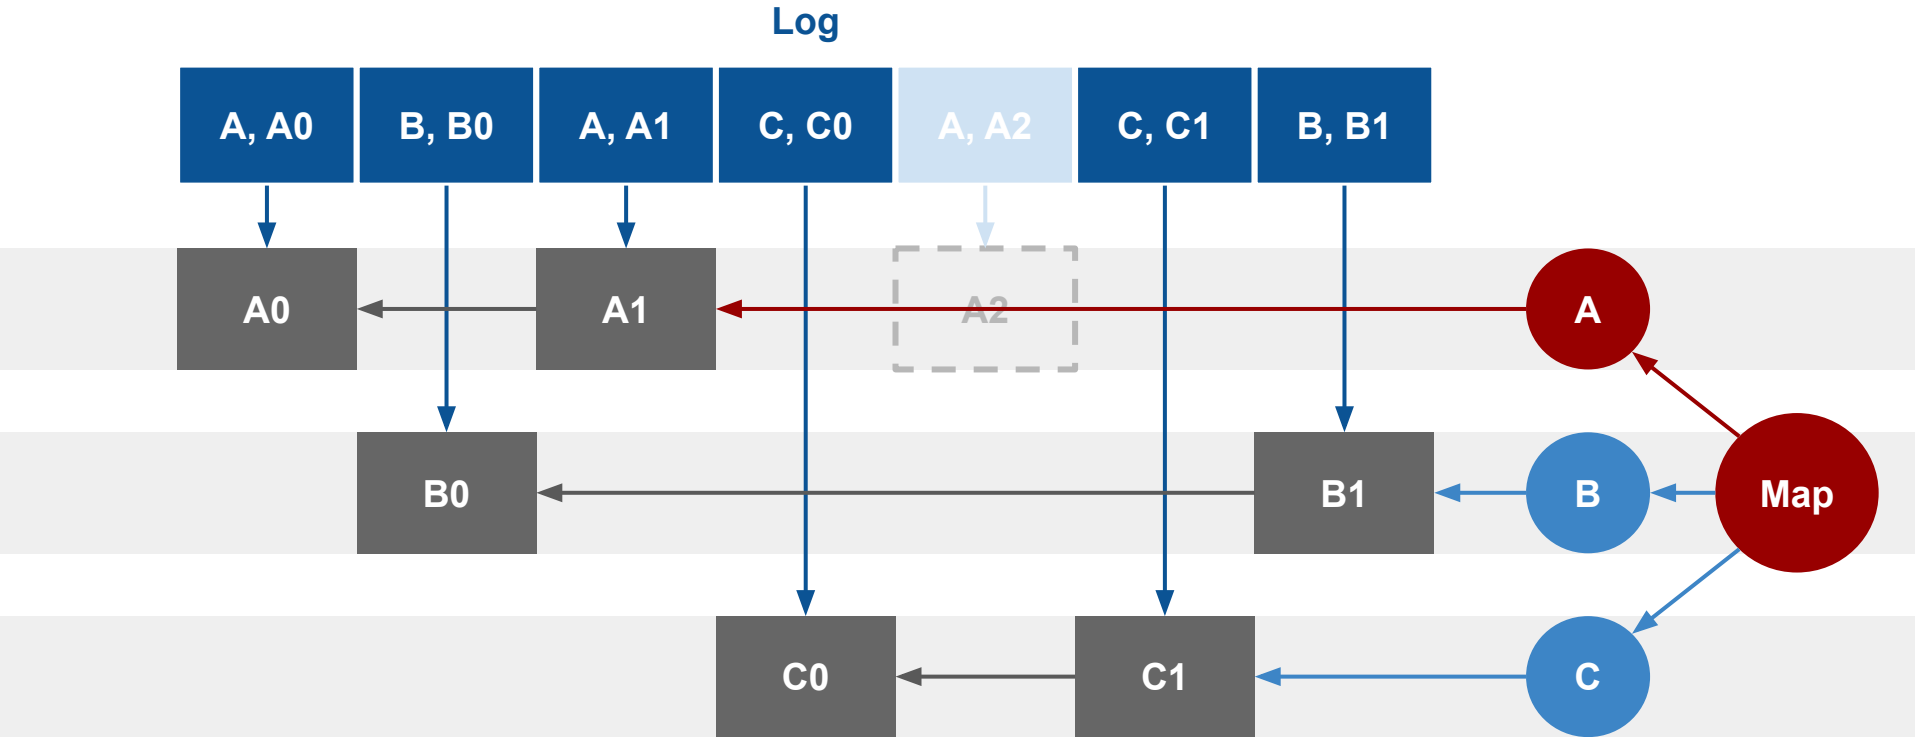
Internal hash linking will fail! (and log inclusion too)



What if an attacker hides new record(s)?



The map checkpoint will no longer match!



# Summary

- **WebAssembly (Wasm)** is a promising way to make *portable and secure* software
- **Package Registries** are *indexes of content*, not necessarily content stores or providers
- **Package Transparency** builds on the the ideas of *Certificate Transparency* to offer registries valuable properties
- **Package Transparency** provides *defenses* against a variety of *attacks*

# Special Thanks to

- Folks who reviewed / advised this presentation,
  - [Lann Martin](#), Fermyon
  - [Luke Wagner](#), Fastly
- Other contributors to the **warg** project,
  - [Bailey Hayes](#), Cosmonic
  - [Peter Huene](#), Fastly
- and finally **SingleStore**
  - for supporting **warg** by enabling me to contribute to this project